
Implementation and Performance Analyses of a Highly Efficient Algorithm for Pressure-Velocity Coupling

Masterarbeit | Fabian Gabel, Studienbereich CE

Betreuer: Dipl.-Ing. U. Falk | Prof. Dr. rer. nat. M. Schäfer



TECHNISCHE
UNIVERSITÄT
DARMSTADT



fnb Fachgebiet
Numerische Berechnungsverfahren
im Maschinenbau





Thesis Statement pursuant to § 22 paragraph 7 of APB TU Darmstadt

I herewith formally declare that I have written the submitted thesis independently. I did not use any outside support except for the quoted literature and other sources mentioned in the paper. I clearly marked and separately listed all of the literature and all of the other sources which I employed when producing this academic work, either literally or in content. This thesis has not been handed in or published before in the same or similar form.
In the submitted thesis the written copies and the electronic version are identical in content.

Date:

Signature:



Abstract

One crucial decision in designing algorithms to numerically solve the incompressible Navier-Stokes equations is how to treat the inherent coupling between the dependent variables velocity and pressure. Due to advances regarding processor and memory hardware technology in the past decades, fully coupled solution algorithms, which denominate algorithms that simultaneously solve a system of partial differential equations for all involved variables, have become attractive in the field of CFD (*Computational Fluid Dynamics*). The objective of the present thesis is to implement and analyse a pressure based fully coupled solution algorithm based on a finite-volume discretization of the Navier-Stokes equations including additional coupling to the temperature equation. The method uses a cell-centered co-located variable arrangement and is able to handle non-orthogonal, three-dimensional, locally refined block-structured grids with hanging nodes. The implementation has been parallelized using the PETSc framework. After the verification using a manufactured solution, the single processor and parallel performance of the implementation are analysed on a parallel benchmark to analyse strong and weak scalability, a benchmark involving a duct flow with complex geometry and the heated cavity benchmark for buoyancy driven flows. The performance analysis always consists of a comparison involving an implementation of a segregated solution algorithm that was developed using the same environment. Furthermore three correctors for non-orthogonality of numerical grids are analysed with respect on their effect on the convergence of the coupled solution algorithm, showing that the *over-relaxed approach* is the most robust for different degrees of non-orthogonality of the grid. In all tests the fully coupled algorithm did outperform the segregated algorithm, provided that the number of involved unknowns was high enough. The solution of the involved linear systems was carried out using black-box solvers from the PETSc framework. It is assumed that major improvements regarding the parallel scalability of the implemented algorithm can be achieved through the use of a special algebraic multigrid methods as preconditioner to the PETSc Krylov subspace solvers.

Contents

1	Introduction	1
1.1	State of Research and Knowledge	1
1.2	Structure of the Thesis	2
2	Fundamentals of Continuum Physics for Thermo-Hydrodynamical Problems	4
2.1	Conservation of Mass – Continuity Equation	4
2.2	Conservation of Momentum – Cauchy-Equations	4
2.3	Closing the System of Equations – Newtonian Fluids	5
2.4	Conservation of Scalar Quantities	5
2.5	Necessary Simplification of Equations	5
2.5.1	Incompressible Flows and Hydrostatic Pressure	5
2.5.2	Variation of Fluid Properties – The Boussinesq Approximation	6
2.6	Final Form of the Set of Equations	7
3	Finite-Volume Methods for Incompressible Flows – Theoretical Basics	8
3.1	Numerical Grid	8
3.2	Approximation of Integrals and Derivatives	9
3.3	Treatment of Non-Orthogonality of Grid Cells	10
3.3.1	Minimum Correction Approach	10
3.3.2	Orthogonal Correction Approach	10
3.3.3	Over-Relaxed Approach	10
3.3.4	Deferred Non-Orthogonal Correction	10
3.4	Numerical Solution of Non-Linear Systems – Picard Iteration	11
4	Implicit Finite-Volume Method for Incompressible Flows – Segregated Approach	12
4.1	Discretization of the Mass Balance	12
4.2	A Pressure-Weighted Interpolation Method for Velocities	12
4.3	Implicit Pressure Correction and the SIMPLE-Algorithm	15
4.4	Discretization of the Mass Fluxes and the Pressure Correction Equation	17
4.5	Discretization of the Momentum Balance	18
4.5.1	Linearization and Discretization of the Convective Term	18
4.5.2	Discretization of the Diffusive Term	19
4.5.3	Discretization of the Source Terms	20
4.6	Discretization of the Temperature Equation	21
4.7	Boundary Conditions	21
4.7.1	Dirichlet Boundary Conditions	21
4.7.2	Treatment of Wall Boundaries	22
4.7.3	Treatment of Block Boundaries	24
4.8	Treatment of the Singularity of the Pressure Correction Equation with Neumann Boundaries	24
4.9	Structure of the Assembled Linear Systems	25
5	Implicit Finite-Volume Method for Incompressible Flows – Fully Coupled Approach	28
5.1	The Fully Coupled Algorithm – Pressure-Velocity Coupling Reconsidered	28
5.2	Coupling to the Temperature Equation	29
5.2.1	Decoupled Approach – Explicit Velocity-to-Temperature Coupling	29
5.2.2	Implicit Velocity-to-Temperature Coupling	30
5.2.3	Temperature-to-Velocity/Pressure Coupling – Newton-Raphson Linearization	30
5.3	Boundary Conditions on Domain and Block Boundaries	31
5.4	Assembly of Linear Systems – Final Form of Equations	32
6	CAFFA Framework	36
6.1	PETSc Framework	36
6.2	Grid Generation and Preprocessing	36
6.3	Implementation of CAFFA	37
6.3.1	The Message-Passing Model	37
6.3.2	Convergence Control	37

6.3.3	Indexing of Variables and Treatment of Boundary Values	38
6.3.4	Domain Decomposition and the Exchange of Ghost Values	38
7	Verification of the developed CAFFA Framework	40
7.1	The Method of Manufactured Solutions for Navier-Stokes Equations	40
7.2	Manufactured Solution for the Navier-Stokes Equations and the Temperature Equation	40
7.3	Measurement of Error and Calculation of Order	41
7.4	Influence of the Under-Relaxation Factor for the Velocities	42
8	Comparison of Solver Concepts	45
8.1	Parallel Performance	45
8.1.1	Employed Hardware and Software – The Lichtenberg-High Performance Computer	45
8.1.2	Measures of Performance	45
8.1.3	Preliminary Upper Bounds on Performance – The STREAM Benchmark	46
8.1.4	Speed-Up Measurement and Impact of Coupling Algorithm for Analytic Test Case	49
8.2	Weak Scaling Comparison of Coupled and Segregated Solution Algorithm	50
8.3	Realistic Testing Scenario – Complex Geometry	51
8.4	Classical Benchmarking Case – Temperature-Driven Cavity Flow	53
8.5	Effect of Different Non-Orthogonal Correctors on Solver Convergence	56
9	Conclusion and Outlook	58
	References	59

List of Figures

1	Vertex centered, cell centered and staggered variable arrangement. For the staggered arrangement, the arrows denote the velocity component and their location on the grid. The centered dot denotes the location of the pressure and further scalar quantities	8
3	Block-structured grid consisting of one $2 \times 2 \times 2$ block and one $3 \times 3 \times 3$ block	9
4	Notation used for one control volume, its boundary faces and its neighbors	9
4	Minimum correction, orthogonal correction and over-relaxed approach	11
5	Possible interpretation of a virtual control volume (grey) located between nodes P and Q	13
6	Non-matching grid cells with hanging nodes at a two-dimensional block boundary. Indexing is based on the face segments S_l	24
7	Non-zero structure of the linear systems used in the SIMPLE-algorithm for a block-structured grid consisting of one $2 \times 2 \times 2$ cell block and one $3 \times 3 \times 3$ cell block	27
8	Non-zero structure of block submatrices of the linear systems used in the coupled solution algorithm for a block-structured grid consisting of one $2 \times 2 \times 2$ cell block and one $3 \times 3 \times 3$ cell block. The blue coefficients represent the pressure-velocity coupling, the red coefficients correspond to the velocity-to-temperature coupling due to the implicit Boussinesq approximation and the green coefficients result from the Newton-Raphson linearization technique addressing the temperature-to-velocity/pressure coupling.	33
9	Non-zero structure of the linear system used in the coupled solution algorithm for a block-structured grid consisting of one $2 \times 2 \times 2$ cell block and one $3 \times 3 \times 3$ cell block. The variables have been interlaced and the matrix consists of blocks as shown in figure 8.	34
10	Non-zero structure of the linear system used in the coupled solution algorithm for a block-structured grid consisting of one $2 \times 2 \times 2$ cell block and one $3 \times 3 \times 3$ cell block. The variables have been ordered in a way that each major matrix block refers to only one variable or the coupling between exactly two variables as it is denoted by the corresponding subindices.	35
11	Storage and update of ghost values in vectors related to variables on multi block domains. The blocks have been assigned to two different processes, <i>Proc 1</i> and <i>Proc 2</i> . The control volumes of the two-dimensional problem domain are indexed with respect to the process local indexing.	39
12	Comparison of calculated error for different under-relaxation factors α_u on grids with $16 \times 16 \times 16$, $32 \times 32 \times 32$ and $64 \times 64 \times 64$ cells	44
13	Sustainable memory bandwidth as determined by the STREAM benchmark (Triad) for different process binding options on one node of the MPI1 section	47
14	Sustainable memory bandwidth as determined by the STREAM benchmark (Triad) for different process binding options on one node of the MPI2 section	48

15	Wall-clock time comparison for the SIMPLE-algorithm, solving for the analytic solution presented in section 7.2 on a structured grid with 64^3 cells using different under-relaxation factors up to $\alpha_u = 0.999$	49
16	Wall-clock time and Speed-Up comparison for segregated and fully coupled solution algorithm, solving for the analytic solution presented in section 7.2 on a structured grid with 128^3 cells	50
17	Comparison of the number of outer iterations and the corresponding wall-clock time, needed to achieve a reduction of $1E-8$ of the initial residual, for different methods to resolve pressure-velocity coupling	51
18	Average number of inner iterations for different process counts solving for the analytic solution presented in section 7.2 with the fully coupled solution algorithm	51
19	Sketch of the channel flow problem domain	52
20	West and east boundary of the numerical grid for the channel flow problem	53
21	Blocking for the two different obstacles within the problem domain of the channel flow	53
22	Temperature field and velocity field in the coordinate direction of the gravitational force of the temperature-driven cavity flow problem for a cross section in the middle of the problem domain	54
23	Skewing of an initially equidistant, structured and orthogonal grid via random movement of inner grid points within a small neighborhood of their original location. The neighborhood with maximal diameter is indicated by a circle around one inner grid point	56
24	Number of needed outer iterations for different relative diameters of the neighborhood in which the movement of grid points takes place, parametrized by the orthogonal corrector	57

List of Tables

1	Comparison of the errors of the velocity, calculated by the segregated and the coupled solver for different grid resolutions and the resulting order of accuracy	43
2	Comparison of the errors of the pressure, calculated by the segregated and the coupled solver for different grid resolutions and the resulting order of accuracy	43
3	Comparison of the errors of the temperature, calculated by the segregated and the coupled solver for different grid resolutions and the resulting order of accuracy	43
4	Characteristic problem properties used in the performance measurements, solving for the analytic solution presented in section 7.2	50
5	Characteristic problem properties used in the channel flow test case	52
6	Performance analysis results of the channel flow problem for different numbers of unknowns comparing the segregated (SEG) to the fully coupled (CPLD) solution algorithm using one process on the MPI2 section of the HHLR supercomputer.	54
7	Characteristic problem properties used in the heated cavity flow test case	55
8	Performance analysis results of the temperature-driven cavity flow problem comparing the SIMPLE-algorithm with segregated temperature solve (SEG), the fully coupled solution algorithm with segregated temperature solve (CPLD), the fully coupled solution algorithm with an implicit Boussinesq approximation (TCPLD) and the fully coupled solution algorithm using an implicit Boussinesq approximation and a semi-implicit Newton-Raphson linearization of the convective part of the temperature equation (NRCPLD).	55

List of Algorithms

1	SIMPLE-Algorithm	17
2	Fully Coupled Solution Algorithm	29



1 Introduction

Due to the advancements in computer architecture and algorithms, the involvement of numerical methods in sciences and industrial practice has increased significantly during the past decades, having a great impact on academic research and processes in industrial development. Besides theoretical or experimental approaches, numerical methods nowadays form an integral part in fostering a better understanding of physical phenomena and in modeling the behaviour of technical systems. Numerical methods are used for the simulation of fluid flow problems, a branch of numerics also known as *CFD* (*Computational Fluid Dynamics*). CFD continuously reaches the limits of current simulation environments, creating new challenges to further improvements of the deployed hardware and software. CFD methods not only have to respond to a higher level of complexity, which can include complex geometries or multiphysical phenomena, they are also expected to generate results after shorter periods of simulation. To overcome these challenges, research is done into the implementation of novel and more efficient algorithms; apart from that, CFD applications often rely on the efficient usage of high performance computer clusters.

In the field of CFD, the incompressible Navier-Stokes equations represent the commonly used mathematical model for incompressible fluid flows. Among other properties, this model couples two variables, velocity and pressure. Thus, algorithms that tend to solve for incompressible flow fields have to address the so called pressure-velocity coupling. A common approach is to resolve this coupling through segregated solution methods. By these methods, the algebraic equations for each of the dependent variables are solved sequentially, using approximate values for the other involved variables. One common representative of this type of algorithms is the SIMPLE-algorithm, being widely used in scientific codes or industrial solvers due to its small memory requirements and the efficient iterative solvers which exist for the resulting linear systems. Reference [1] reviews the evolution of this type of algorithms.

The use of segregated solution algorithms entails significant deficiencies. In order to stabilize the iteration for the coupling process, the obtained solutions have to be under-relaxed. In doing so, the degree of under-relaxation also depends on the problem to be solved. Not only does this degrade the solver's performance, but, making things even more difficult, there also don't exist strict rules for the best selection of the appropriate amount of under-relaxation. Furthermore, algorithms like SIMPLE lack scalability with the mesh size. Resolving the pressure-velocity coupling defines the flow solver's performance. Consequently, new coupling algorithms, more robust or even independent of parameters, and, at the same time, resource efficient in their application, have to be developed. A helpful analysis of pressure-velocity coupling can be found in [48].

1.1 State of Research and Knowledge

Considering the solution approach, pressure based solution algorithms for incompressible Navier-Stokes equations can be divided into segregated solution methods and semi-direct or coupled solution methods. Segregated solution methods, like the popular SIMPLE-algorithm [46], have been actively used for the past decades. One common property of segregated solution algorithms, is the two-step solution approach, involving a predictor step and a corrector step. In the first step, a velocity field is calculated based on a guessed pressure field, in the second step a pressure or pressure correction equation is solved. After this, the velocities are corrected. Many different variants of segregated solution methods have been developed over the time, aiming to increase the robustness and efficiency of the SIMPLE-algorithm. The most popular ones, besides the SIMPLE-algorithm, are SIMPLER (SIMPLE Revised) [47] and SIMPLEC (SIMPLE Consistent) [62], which modify the SIMPLE-algorithm and the PISO-algorithm which is based on operator-splitting. The main advantages of segregated solution algorithms are low memory requirements and the straightforward extension to the solution of additional partial differential equations. An overview over segregated solution algorithms can be found in [13].

The SIMPLER-algorithm addresses the inherent problem of the need to under-relax the variables throughout the solution process. Based on the observation, that the pressure correction based velocity corrections are satisfactory with respect to the improvement of the solution, the pressure corrections have to be under-relaxed before they are applied to the pressure itself, because they tend to be too high. In order to remove the need to under-relax the pressure the pressure correction from the SIMPLE algorithm is only used for velocity corrections, whereas a separate pressure equation is derived, solving for the pressure based on the corrected velocities. The SIMPLEC-algorithm proposes another remedy to the under-relaxation problematic by introducing a consistent simplification in the derivation of the pressure equation that eliminates the need to under-relax pressure. The SIMPLER-algorithm is found to efficiently reduce the number of iterations for convergence, while increasing the computational effort and having higher memory requirements. The SIMPLEC-algorithm also leads to a smaller, compared to the SIMPLER-algorithm, decrease in the number of iterations for convergence, while there is no notable increase in the computational effort.

The PISO-algorithm, as does the SIMPLE-algorithm, uses a pressure correction step to correct the velocities, which is then, different to the SIMPLE-algorithm, followed by further correction steps to compensate the simplifications made in the derivation of the pressure-correction equation. Similar to the other modifications, the PISO-algorithm leads to a reduced iteration count for convergence and lacks the need to under-relax the pressure correction. One known drawback of the PISO-algorithm is the higher computational effort and the higher memory requirements.

An alternative solution approach to the problem of numerically solving the Navier-Stokes equations is by means of coupled solvers, which simultaneously solve the momentum and mass balance equations. Algorithms of this type, using finite-volume methods, have been presented in [9, 12, 18, 20, 30, 36, 61]. It has been shown [12] that the principal advantages of coupled solution approaches compensate for the deficiencies of segregated solution algorithms. All mentioned references report that on the one hand, the computational time is reduced, while, on the other hand, the robustness of the solution algorithm increases. According to [12] no under-relaxation is necessary at all to obtain a robust implementation, [20] reported insensitivity towards the relaxation allowing faster convergence compared to popular segregated solution algorithms. In all of the mentioned references the two-step, predictor-corrector solution procedure was reduced to one step. An equation for the pressure was deduced using a momentum interpolation scheme [51] in all cited references but [20].

[9] extended the linear system for velocities and pressure by an additional decoupled block accounting for temperature. So far, only [20, 61] have used coupled solution processes for flow problems implicitly involving temperature transport. The mentioned references extended the scope of a coupled solution process for the two-dimensional Navier-Stokes equations to include a temperature equation. They investigated the effect of different methods for implicit temperature-to-velocity/pressure coupling via a Newton-Raphson linearization on the solver's performance, showing that the number of iterations until convergence is significantly reduced by the use of coupled methods. The objectives of the present thesis are to implement a fully-coupled solution algorithm for the three-dimensional Navier-Stokes equations and the temperature equation on co-located grids, and to analyse their single and multi-processor performance. The methods to achieve temperature coupling are similar to the ones presented in [61]. Different from the implementation in the present thesis is the treatment of the interpolation of the volumes in the pressure weighted interpolation method. Another difference results from the used interpolation technique for pressure gradients at faces, for which [61] uses local coordinate systems and a wider computational stencil.

1.2 Structure of the Thesis

In section 2, the thesis introduces the mathematical and physical fundamentals of flow problems. Section REF presents the main frame of partial differential equations that have to be solved numerically. Each subsection outlines a derivation from the continuum mechanical point of view, and shows common simplifications, comprising not only the Navier-Stokes equations for incompressible fluids but also the Boussinesq approximation, which is widely used for incompressible flows experiencing buoyant forces.

Section number 3 outlines the fundamentals of finite-volume methods in general. At first, the necessary terminology for numerical grids is introduced. After that, the main approximation principles for the discretization of integral or differential operators are presented, placing particular emphasis on the modifications necessary for discretizing them on non-orthogonal grids. At the end, the chapter outlines the main method of linearization for the Navier-Stokes equations used throughout the work.

Section 4 focuses on segregated solution methods, because many ideas of the commonly used SIMPLE-algorithm are also relevant for the implementation of coupled solution methods. The section starts by discretizing the mass balance coming directly from the continuum mechanical introduction of the first section, outlining its drawbacks on their applicability to co-located variable arrangements. Subsection 4.2 introduces a remedy to the *checker-boarding* effect surging from an unmodified discretization of the mass balance, by introducing a modification of the commonly used Rhie-Chow [51] momentum interpolation scheme. This modification yields results which do not depend on under-relaxation. This is important for later comparison studies with the coupled solution algorithm. In subsection 4.3, the popular SIMPLE-algorithm is derived, using the previously demonstrated interpolation scheme as a basis for the construction of a pressure correction equation. In subsections 4.4, 4.5 and 4.6, the pressure equation and the two remaining sets of equations, i.e. the three momentum balances and the temperature equation, are discretized. Each section concludes with a complete list of the resulting discretized equations and calculations of the coefficients as they have been implemented in the developed solver framework. Subsection 4.7 is dedicated to the implementation of boundary conditions, presenting the necessary modifications to the discretized equations for each boundary condition. This includes the treatment of block boundaries. Since in most cases the boundary conditions lead to a singular system for the pressure correction, subsection 4.8 describes methods to constrain the pressure in incompressible flows. This section concludes by presenting the non-zero structure of the resulting assembled linear system, taking into account the special treatment of block boundaries with hanging nodes.

The structure of section 5, describing the implemented coupled solution algorithm, resembles section 4, since both solution approaches share a great part of their components. This is why special emphasis is put on the differences between both algorithms. Subsection 5.1 therefore reconsiders the pressure velocity coupling that accounts for the implicit consideration of the pressure in the momentum balances. A major difference to segregated solution methods is then presented in 5.2: the coupling to the temperature equation. Here, different methods, varying in their relation of implicit to explicit coupling of velocities, pressure, and temperature, are deduced. After that, the necessary modifications for boundary conditions are presented. Subsection 5.4 shows the final form of the discretized equations and the resulting

structure of the assembled linear systems, placing special emphasis on the effects of the different methods leading to implicit coupling of the velocities to the temperature equation and vice versa.

Section 6 introduces the developed *CAFFA* (*Computer Aided Fluid Flow Analysis*) framework, having its theoretical basis in the preceding chapters. The framework is linked to the PETSc library, being a sophisticated framework for the parallelisation of solver applications surging in scientific computing, and providing a set of efficient preconditioners and Krylov subspace solvers for the solution of linear systems. Subsection REF then introduce the other building blocks of the framework, beginning with a grid generator for arbitrarily refined block-structured grids and a preprocessor program, concluding with implementation details of the CAFFA solver. These details comprise the used message passing model, the control of convergence and the data decomposition in the context of parallel computation.

Section 7 presents the results of a conducted convergence study to verify the CAFFA framework. After a brief introduction to the theory behind the verification of scientific codes via manufactured solutions, a manufactured solution for the Navier-Stokes equations using the Boussinesq approximation is proposed by the author, followed by the results of a grid convergence study to prove the order of accuracy of the developed solvers within the CAFFA framework. The section concludes with the results of a short study on the effect the under-relaxation factors of the velocities have on the final result of calculations.

Section 8 presents the results of the performance analyses conducted on the high performance cluster *HHLR* (*Hessischer Hochleistungsrechner*) of TU Darmstadt. After introducing the used hardware and deployed software configuration and the relevant performance metrics, program runtime is measured and application speedup is calculated. Furthermore, a study on the weak scalability of the developed fully coupled solution algorithm is conducted. The section furthermore presents performance analyses regarding the efficiency of coupled and segregated solution algorithms for flow problems, considering, among other test cases, complex geometries and buoyancy forces. The section concludes showing the results of a conducted study on the effect of different types of correctors addressing the non-orthogonality of the grid on the convergence of the coupled solution algorithm.

The last section 9 summarizes the results of the conducted research work, giving an outlook on further research.

2 Fundamentals of Continuum Physics for Thermo-Hydrodynamical Problems

This section covers the set of fundamental equations for thermo-hydrodynamical problems which the numerical solution techniques of the following chapters are aiming to solve. Furthermore, the notation regarding the physical quantities to be used throughout this thesis is introduced. The following paragraphs are based on [2, 19, 31, 58]. For a thorough derivation of the matter to be presented the reader may consult the mentioned sources. Since the present thesis focuses on the application of finite-volume methods, the focus lays on stating the integral forms of the relevant conservation laws. However, the process of deriving the final set of equations requires the use of differential formulations of the stated laws. Einstein's convention for taking sums over repeated indices simplifies certain expressions. The remainder of this thesis uses non-moving inertial frames in a Cartesian coordinate system with the coordinates $x_i, i = 1, \dots, 3$. This approach is also known as *Eulerian approach*.

2.1 Conservation of Mass – Continuity Equation

The conservation law of mass, also known as the continuity equation, embraces the physical concept that, neglecting relativistic effects and nuclear reactions, mass cannot be created or destroyed. Using the notion of a mathematical control volume to denote a constant domain of integration, one can state the integral mass balance of a control volume V with control surface S with surface unit normal vector $\mathbf{n} = (n_i)_{i=1,\dots,3}$ using Gauss's theorem as

$$\iiint_V \frac{\partial \rho}{\partial t} + \frac{\partial}{\partial x_i} (\rho u_i) dV = \iiint_V \frac{\partial \rho}{\partial t} dV + \iint_S \rho u_i n_i dS = 0,$$

where ρ denotes the material density, t denotes the independent variable of time and $\mathbf{u} = (u_i)_{i=1,\dots,3}$ is the velocity vector field. Since this equation remains valid for arbitrary control volumes, the equality has to hold for the integrands as well. In this sense, the differential form of the conservation law of mass can be formulated as

$$\frac{\partial \rho}{\partial t} + \frac{\partial}{\partial x_i} (\rho u_i) = 0. \quad (1)$$

2.2 Conservation of Momentum – Cauchy-Equations

The conservation law of momentum, also known as Newton's Second Law, axiomatically demands the balance of the temporal change of momentum and the sum of all attacking forces on a body. These forces can be divided into body forces and surface forces. Let $\mathbf{k} = (k_i)_{i=1,\dots,3}$ denote a mass specific force and $\mathbf{t} = (t_i)_{i=1,\dots,3}$ the stress vector. A first form of the integral momentum balance in the direction of x_i can be formulated as

$$\iiint_V \frac{\partial}{\partial t} (\rho u_i) dV + \iint_S \rho u_i (u_j n_j) dS = \iiint_V \rho k_i dV + \iint_S t_i dS. \quad (2)$$

In general, the stress vector \mathbf{t} is a function not only of the location $\mathbf{x} = (x_i)_{i=1,\dots,3}$ and of the time t but also of the surface unit normal vector \mathbf{n} . A central simplification can be introduced, namely Cauchy's stress theorem, which states that the stress vector is the image of the unit normal vector under a linear mapping \mathbf{T} . With respect to the Cartesian canonical basis $(\mathbf{e}_i)_{i=1,\dots,3}$, the mapping \mathbf{T} is represented by the coefficient matrix $(\tau_{ji})_{i,j=1,\dots,3}$, and Cauchy's stress theorem reads

$$\mathbf{t}(\mathbf{x}, t, \mathbf{n}) = \mathbf{T}(\mathbf{x}, t, \mathbf{n}) = (\tau_{ji} n_j)_{i=1,\dots,3}.$$

Assuming the validity of Cauchy's stress theorem, one can derive Cauchy's first law of motion, which in differential form can be formulated as

$$\frac{\partial}{\partial t} (\rho u_i) + \frac{\partial}{\partial x_j} (\rho u_i u_j) = \rho k_i + \frac{\partial \tau_{ji}}{\partial x_j} \quad (3)$$

representing the starting point for the modeling of fluid mechanical problems. One should note that Cauchy's first law of motion does not make any assumptions regarding material properties, which is why the set of equations (1,3) is not closed, meaning that there does not exist an independent equation for each of the dependent variables.

2.3 Closing the System of Equations – Newtonian Fluids

As a result of Cauchy's theorem, the stress vector \mathbf{t} can be specified once the nine components τ_{ji} of the coefficient matrix are known. As shown in [31, 58], by formulating the conservation law of angular momentum the coefficient matrix is symmetric,

$$\tau_{ji} = \tau_{ij}, \quad (4)$$

hence the number of unknown coefficients may be reduced to six unknown components. In a first step, it is assumed that the coefficient matrix can be decomposed into fluid-static and fluid-dynamic contributions,

$$\tau_{ij} = -p\delta_{ij} + \sigma_{ij},$$

where p is the thermodynamic pressure, δ_{ij} is the *Kronecker-Delta* and σ_{ij} is the so called *deviatoric stress tensor*.

In the present thesis, viscous fluids are modeled using a linear relation between the components of the deviatoric stress tensor and the symmetric part of the transpose of the Jacobian of the velocity field $(S_{ij})_{i,j=1,\dots,3}$,

$$S_{ij} = \frac{1}{2} \left(\frac{\partial u_i}{\partial x_j} + \frac{\partial u_j}{\partial x_i} \right).$$

If one now imposes material-isotropy and the mentioned stress-symmetry (4) restriction, it can be shown [3] that the constitutive equation for the deviatoric stress tensor reads

$$\sigma_{ij} = 2\mu S_{ij} + \lambda S_{mm} \delta_{ij},$$

where λ and μ denominate scalars which depend on the local thermodynamical state. Taking everything into account, (3) can be formulated as the differential conservation law of momentum for Newtonian fluids, better known as the *Navier-Stokes equations* in differential form:

$$\frac{\partial}{\partial t}(\rho u_i) + \frac{\partial}{\partial x_j}(\rho u_i u_j) = \rho k_i - \frac{\partial p}{\partial x_i} + \frac{\partial}{\partial x_j} \left(\mu \left(\frac{\partial u_i}{\partial x_j} + \frac{\partial u_j}{\partial x_i} \right) \right) + \frac{\partial}{\partial x_i} \left(\lambda \frac{\partial u_m}{\partial x_m} \right) \quad (5)$$

2.4 Conservation of Scalar Quantities

The modeling of the transport of scalar or vector quantities by a flow field \mathbf{u} , also known as convection, is necessary if the fluid mechanical problem to be analyzed includes, for example, heat transfer. Other scenarios that involve the necessity to model scalar transport surge, when turbulent flows are to be modelled by two-equation models like the k - ε -model [49].

Since this thesis focuses on the transport of the scalar temperature T , this section introduces the conservation law of energy in differential form, formulated in terms of the temperature,

$$\frac{\partial (\rho T)}{\partial t} + \frac{\partial}{\partial x_j} \left(\rho u_j T - \kappa \frac{\partial T}{\partial x_j} \right) = q_T,$$

where κ denotes the thermal conductivity of the modelled material and q_T is a scalar field representing sources and sinks of heat throughout the domain of the problem. This equation is also known as the temperature equation.

2.5 Necessary Simplification of Equations

The purpose of this section is to introduce and justify further common simplifications of the previously presented set of constitutive equations.

2.5.1 Incompressible Flows and Hydrostatic Pressure

A common simplification when modeling low Mach number flows ($Ma < 0.3$), is the assumption of *incompressibility*, or the assumption of an *isochoric* flow. If one furthermore assumes homogeneous density ρ in space and time, a restrictive assumption that will be partially alleviated in the following section, the continuity equation in differential form (1) can be simplified to

$$\frac{\partial u_i}{\partial x_i} = 0.$$

In other words: In order for a velocity vector field \mathbf{u} to be valid for an incompressible flow, it has to be free of divergence, in other terms *solenoidal* [3, 58].

If furthermore one assumes the dynamic viscosity μ to be constant, which can be suitable in the case of isothermal flow or if the temperature differences within the flow are small, the Navier-Stokes equations in differential form can be reduced to

$$\frac{\partial}{\partial t}(\rho u_i) + \rho \frac{\partial}{\partial x_j}(u_i u_j) = \rho k_i - \frac{\partial p}{\partial x_i} + \frac{\partial}{\partial x_j} \left(\mu \left(\frac{\partial u_i}{\partial x_j} + \frac{\partial u_j}{\partial x_i} \right) \right) \quad (6a)$$

$$= \rho k_i - \frac{\partial p}{\partial x_i} + \mu \frac{\partial}{\partial x_j} \left(\frac{\partial u_i}{\partial x_j} \right) \quad (6b)$$

by using *Schwartz's* lemma to interchange the order of differentiation. Another common step to further simplify the set of equations is the assumption of a volume specific force $\rho \mathbf{k}$ that can be modelled by a potential, in such a way that it can be represented as the gradient of a scalar field $\Phi_{\mathbf{k}}$ as

$$-\rho k_i = \frac{\partial \Phi_{\mathbf{k}}}{\partial x_i}.$$

In the case of this thesis, this assumption is valid, since the mass specific force is the mass specific gravitational force $\mathbf{g} = (g_i)_{i=1,\dots,3}$, and the density is assumed to be constant, so that the potential can be modeled as

$$\Phi_g = -\rho g_j x_j.$$

This term can be interpreted as the hydrostatic pressure p_{hyd} and can be added to the thermodynamical pressure p to simplify calculations

$$\begin{aligned} \rho g_i - \frac{\partial p}{\partial x_i} &= \frac{\partial}{\partial x_i}(\rho g_j x_j) - \frac{\partial p}{\partial x_i} \\ &= \frac{\partial}{\partial x_i}(\rho g_j x_j) - \frac{\partial}{\partial x_i}(\hat{p} + p_{hyd}) \\ &= -\frac{\partial \hat{p}}{\partial x_i}. \end{aligned}$$

Since in incompressible fluids only pressure differences matter, this has no effect on the solution. After finishing the calculations p_{hyd} can be calculated and added to the resulting pressure \hat{p} .

2.5.2 Variation of Fluid Properties – The Boussinesq Approximation

Modeling of an incompressible flow with heat transfer has to take into account that fluid properties, as for example density, change with varying temperature. If the variation of temperature is small, one can still assume constant density to maintain the structure of the advection and diffusion terms in (5), and only consider changes of the density in the gravitational term. If linear variation of density with respect to temperature is assumed, this approximation is called *Boussinesq* approximation [22]. This approximation furthermore consists in assuming that all other fluid properties are constant and viscous dissipation can be neglected. In this case, the Navier-Stokes equations are formulated using a reference density ρ_0 at the reference temperature T_0 and the following temperature-dependent density ρ

$$\rho(T) = \rho_0(1 - \beta(T - T_0)).$$

Here, β denotes the coefficient of thermal expansion. Using the Boussinesq approximation, the incompressible Navier-Stokes equations in differential form can be formulated as

$$\begin{aligned}
\rho_0 \frac{\partial (u_i)}{\partial t} + \rho_0 \frac{\partial}{\partial x_j} (u_i u_j) &= \rho_0 g_i + (\rho - \rho_0) g_i - \frac{\partial p}{\partial x_i} + \mu \frac{\partial}{\partial x_j} \left(\frac{\partial u_i}{\partial x_j} + \frac{\partial u_j}{\partial x_i} \right) \\
&= \frac{\partial}{\partial x_i} (\rho_0 g_j x_j) + (\rho - \rho_0) g_i - \frac{\partial p}{\partial x_i} + \mu \frac{\partial}{\partial x_j} \left(\frac{\partial u_i}{\partial x_j} + \frac{\partial u_j}{\partial x_i} \right) \\
&= - \frac{\partial \hat{p}}{\partial x_i} + (\rho - \rho_0) g_i + \mu \frac{\partial}{\partial x_j} \left(\frac{\partial u_i}{\partial x_j} + \frac{\partial u_j}{\partial x_i} \right) \\
&= - \frac{\partial \hat{p}}{\partial x_i} - \rho_0 \beta (T - T_0) g_i + \mu \frac{\partial}{\partial x_j} \left(\frac{\partial u_i}{\partial x_j} + \frac{\partial u_j}{\partial x_i} \right)
\end{aligned}$$

using $\rho \mathbf{g}$ as the mass specific force.

2.6 Final Form of the Set of Equations

In the previous subsections different simplifications have been introduced, which will be used throughout the thesis. The final form of the set of equations to be used is thereby presented. As a further simplification, the modified pressure \hat{p} will be treated as p , and since the use of the Boussinesq approximation replaces the variable ρ by a linear function of the temperature T , the reference density ρ_0 for the remainder of this thesis will be referred to as ρ . It should be noted that incompressibility has been taken into account

$$\frac{\partial u_i}{\partial x_i} = 0. \quad (7a)$$

$$\rho \frac{\partial (u_i)}{\partial t} + \rho \frac{\partial}{\partial x_j} (u_i u_j) = - \frac{\partial p}{\partial x_i} - \rho \beta (T - T_0) g_i + \mu \frac{\partial}{\partial x_j} \left(\frac{\partial u_i}{\partial x_j} + \frac{\partial u_j}{\partial x_i} \right) \quad (7b)$$

$$\frac{\partial (\rho T)}{\partial t} + \frac{\partial}{\partial x_j} \left(\rho u_j T - \kappa \frac{\partial T}{\partial x_j} \right) = q_T. \quad (7c)$$

3 Finite-Volume Methods for Incompressible Flows – Theoretical Basics

This section deals with the fundamentals of the numerical solution via a finite-volume method of the formerly presented set of partial differential equations (7). The focus of this section is on providing an overview over the methods to be used in the present thesis. The information contained in this section is based on [19, 28, 43, 55]. The overview starts by mentioning the different grid types to be used and the discretization techniques to be applied. On the basis of integral formulations of the equations to be solved, the therein contained integrals and differential operators have to be discretized. This chapter furthermore presents different approaches to handle corrections for cases of increased non-orthogonality causing degraded grid quality.

The finite-volume method aims at providing particular linear algebraic equations to be used in determining an approximate solution of a partial differential equation. However, since the Navier-Stokes equations are non-linear, an intermediate step, before solvers for linear equation systems can be applied, has to be taken, by linearizing the discrete equations. This leads to the need of an iteration process, the *Picard iteration*, which will also be briefly explained.

3.1 Numerical Grid

This subsection provides a brief overview of the general grid structure to be used in the present thesis. The main idea behind finite-volume methods is to solve partial differential equations by integrating them over the specified continuous problem domain and dividing this domain into a finite number of subdomains, the so called *control volumes*. The result of this finite partition of a continuous problem domain is called numerical grid. Every grid consists of a finite number of control volumes representing the boundaries of a discrete domain of integration. Depending on whether the numerical solution of an equation is to be calculated on the boundary vertices or in the center of a grid cell, the variable arrangement is denoted to be vertex-oriented or cell center-oriented. The remainder of the present thesis assumes that the variables reside in the cell centers of the respective control volumes.

Another option to arrange variables specifically for algorithms that solve the Navier-Stokes equations is the so called *staggering*. This method stores the variable values of different variables at different grid locations, which facilitates the interpolation of the respective values. As the methods employed in the present thesis are intended to be generally applicable to complex geometries, the cell-centered non-staggered approach offers more flexibility. Figure 1 shows the different variable arrangement options on a two dimensional grid.




Figure 1: Vertex centered, cell centered and staggered variable arrangement. For the staggered arrangement, the arrows denote the velocity component and their location on the grid. The centered dot denotes the location of the pressure and further scalar quantities

Regarding the treatment of domain boundaries and the ordering of the cells within the problem domain, different types of numerical grids can be distinguished. The present thesis makes use of so called boundary-fitted, block-structured grids with hexahedron cells. A structured grid is characterized by a constant amount of grid cells in each coordinate direction. The high regularity of structured grids benefits the computational efficiency of the algorithms to be used on this type of grid. In addition to that, a block-structured grid consists of different grid blocks, of which every single one, considered individually, is structured; but considering the grid as a whole, the grid is unstructured. An example of a block-structured grid with distinguishable grid blocks is given in figure 3. The use of block-structured grids is motivated by the need to increase the adaptivity of structured grids while maintaining high computational efficiency. Furthermore,

it embraces the concept of domain decomposition which facilitates the implementation of parallel algorithms for the decomposed computational domain. Boundary-fitted grids represent domain boundaries by means of the geometry of the control volumes residing on those boundaries. This approach may lead to skewed cells, which affects the accuracy of the results obtained in these boundary cells. Using local grid refinement techniques to refine the grid in regions of geometrically complex boundaries can alleviate these effects.

Inside a structured grid block, cells with the shape of hexahedrons are used. In addition to the geometric boundaries of each control volume, a numerical grid also provides a mapping that assigns to each control volume with index P a set of indexes of neighboring control volumes $NB(P) := \{W, S, B, T, N, E\}$, which are named after the geographic directions. Figure 4 shows a single grid cell with its direct neighbors. The faces $\{S_w, S_s, S_b, S_t, S_n, S_e\}$ of each hexahedral control volume represent the mentioned geometric boundaries.




Figure 3: Block-structured grid consisting of one $2 \times 2 \times 2$ block and one $3 \times 3 \times 3$ block




Figure 4: Notation used for one control volume, its boundary faces and its neighbors

3.2 Approximation of Integrals and Derivatives

In the course of transforming a partial differential equation into a system of linear algebraic equations, integrals and derivatives have to be approximated. The simplest method for approximating an integral is by using the *midpoint rule*. This rule is similar to the mean value theorem of integration, which states that there exists a point $\xi \in V$ for a Riemann integrable function ϕ such that $\phi(\xi) \int_V dV = \int_V \phi(x) dV$. For the midpoint rule ξ is taken to be the center of mass of V . If the integration domain V is indeed a volume, fortunately the calculation of $\phi(\xi)$ with $\xi := (\int_V x_i dV / \int_V dV)_{i=1,\dots,3}$ presents no difficulties since, due to the collocated variable arrangement, the value of ϕ is stored in the cell center, which corresponds to the location ξ . However, if the domain of integration is a surface, a preceding interpolation step is necessary.

On the other hand, in order to transform a partial differential equation into a linear algebraic equation, it is necessary to discretize the differential operators of the equations. For numerical reasons, two different discretization techniques are used in this thesis. A common task is discretizing expressions of the form

$$(\nabla \phi)_e \cdot \mathbf{n}_e,$$

where $(\nabla \phi)_e$ is the gradient of ϕ on a boundary face S_e . One method is to interpret this expression directly as a directional derivative and approximate it with a central difference

$$(\nabla \phi)_e \cdot \mathbf{n}_e \approx \frac{\phi_P - \phi_E}{\|\mathbf{x}_P - \mathbf{x}_E\|_2}. \quad (8)$$

Another method would be to first calculate the cell center gradients $(\nabla \phi)_P$ and $(\nabla \phi)_E$ and interpolate them linearly before calculating the projection onto \mathbf{n}_e

$$(\nabla \phi)_e \cdot \mathbf{n}_e \approx [\gamma_e (\nabla \phi)_P + (1 - \gamma_e) (\nabla \phi)_E] \cdot \mathbf{n}_e, \quad (9)$$

where $\gamma_e := \|\mathbf{x}_p - \mathbf{x}_e\|_2 / \|\mathbf{x}_p - \mathbf{x}_E\|_2$ is a geometric interpolation factor. For calculating the cell center gradients, a method based on Gauss's integration theorem and the midpoint rule for volume integration is employed

$$(\nabla \phi)_{i,p} = \left(\frac{\partial \phi}{\partial x_i} \right)_p \approx \sum_f \phi_f n_i S_f. \quad (10)$$

3.3 Treatment of Non-Orthogonality of Grid Cells

Unfortunately, real applications involve complex geometries, a fact which in turn affects the grid's orthogonality. On non-orthogonal meshes, the directional derivative in direction of the face's unit normal vector \mathbf{n}_e can no longer be approximated as in (8). On the other hand, the exclusive usage of (10) is not desirable due to the bigger truncation error following from this approximation [28]. Hence, a compromise is made and the surface vector $\mathbf{S}_e := S_e \mathbf{n}_e$ is decomposed as

$$\mathbf{S}_e = \Delta + \mathbf{k}, \quad (11)$$

where Δ is parallel to the vector $\mathbf{d}_e := (\mathbf{x}_E - \mathbf{x}_p)$ that directly connects the center \mathbf{x}_p of the control volume P to the center \mathbf{x}_E of its neighbor E . This vector controls the *orthogonal* contribution to the directional derivative. Vector \mathbf{k} controls the influence of the *non-orthogonal* contribution. In the next paragraphs, the three main decompositions of the surface vector \mathbf{S}_e will be presented by stating the respective expression for Δ . The resulting vector \mathbf{k} can be calculated by using (11). One important characteristic common to all the presented approaches, is that the non-orthogonal contribution vanishes, when an orthogonal grid is used. For reasons of graphical representation, the decompositions are chosen to be represented two-dimensionally. Figure 24 gives a geometrical interpretation of the three approaches. The last subsection deals with the integration of one generic approach into the discretization process.

The results of a study on the effect of different correctors on the convergence behaviour of a solver for partial differential equations can be found in section 8.5.

3.3.1 Minimum Correction Approach

This is the approach as proposed in [43]. It should be noted that even though [19] references [43], [19] uses a different approach, which will be presented in the next paragraph. This method is designed to keep the non-orthogonal contribution minimal by always choosing \mathbf{k} to be orthogonal to Δ , which leads to

$$\Delta = (\mathbf{d} \cdot \mathbf{S}_e) \frac{\mathbf{d}}{\mathbf{d} \cdot \mathbf{d}}.$$

It should be noted that the influence of the orthogonal contribution decreases with increasing non-orthogonality of the grid.

3.3.2 Orthogonal Correction Approach

The following method for decomposing the surface normal vector is presented in [19], an approach implemented in the developed solvers. This approach uses a simple projection, independent of the non-orthogonality of the grid. As a result, the orthogonal contribution $\|\Delta\|_2 = \|\mathbf{S}_e\|_2$ and is thus modeled as

$$\Delta = S_e \frac{\mathbf{d}}{\|\mathbf{d}\|_2}.$$

3.3.3 Over-Relaxed Approach

The last approach is used in [12, 28] and is characterized by an increasing influence of the orthogonal contribution with increasing grid non-orthogonality, as opposed to the minimum correction approach. The orthogonal contribution is calculated as

$$\Delta = S_e^2 \frac{\mathbf{d}}{\mathbf{d} \cdot \mathbf{S}_e}.$$

3.3.4 Deferred Non-Orthogonal Correction

In order to reduce the computational stencil that would be necessary to handle the non-orthogonal correction implicitly, the correction will be treated explicitly, using a deferred correction which guarantees, that in the case of a fully converged




Figure 4: Minimum correction, orthogonal correction and over-relaxed approach

solution, only the face normal derivative is taken into account. Generally, the discretization using a non-orthogonal correction would yield

$$(\nabla\phi)_e \cdot \mathbf{S}_e \approx (\nabla\phi)_e \cdot \Delta + (\nabla\phi)_e \cdot \mathbf{k},$$

where the first term can be approximated using a central differencing scheme for the directional derivative, and the second can be approximated by interpolating cell center gradients. If one furthermore makes use of the fact that this method is applied within a solution algorithm for a non-linear system of partial differential equations, a deferred correction can be implemented, which ensures a smaller error from the non-orthogonality. In the case of the previously mentioned discretization techniques for partial derivatives, a possible deferred correction approach reads

$$(\nabla\phi)_e \cdot \mathbf{S}_e \approx \|\Delta\|_2 \frac{\phi_P - \phi_E}{\|\mathbf{x}_P - \mathbf{x}_E\|_2} - (\nabla\phi)_e^{(n-1)} \cdot (\Delta - \mathbf{S}_e).$$

It should be noted that the use of a deferred correction, in conjunction with the requirement that the non-orthogonal correction vanishes on orthogonal grid, introduces an inconsistent discretization of $(\nabla\phi)_e \cdot \Delta$.

3.4 Numerical Solution of Non-Linear Systems – Picard Iteration

In the process of solving non-linear systems, normally two levels of iterations are distinguished: The *inner* iterations, which account for the iterations needed to solve a given system of linear equations, and the *outer* iterations which deal with a possible non-linearity or coupling of the equations. As shown in section 3.3, outer iterations are also necessary if one chooses the deferred correction approach in order to blend higher and lower order discretization schemes, or in order to apply non-orthogonal correctors.

The system of partial differential equations (7) is non-linear due to the terms $\rho u_i u_j$ which often are denoted as the convective terms of the Navier-Stokes equations. It is a common approach to linearize those terms according to the *Picard-Iteration*. For the convective terms of the Navier-Stokes equations or the temperature equation, this means to guess the mass flux ρu_j and approximate as

$$\rho u_j u_i \approx (\rho u_j)^0 u_i.$$

The same procedure applies to the convective terms of the temperature equation $\rho u_j T$. The advantages of this linearization technique are the low memory requirements and the small computational effort for one iteration. This is why this method is commonly found to be implemented in solution programs for flow problems that use a segregated solution algorithm similar to the one used in this thesis. However, compared to Newton-like methods, this approach needs more iterations to converge. Another method to linearize convective terms, applicable when using a fully coupled solution algorithm, will be presented in section 5.2.3.

4 Implicit Finite-Volume Method for Incompressible Flows – Segregated Approach

In this section, an algorithm, addressed to resolve the pressure velocity coupling, is introduced. Furthermore, an under-relaxation factor independent method of calculating mass fluxes by interpolation is introduced and the detailed derivation of all coefficients that result from the discretization process is presented. Finally, the boundary conditions, that are relevant for the present thesis will be shown.

4.1 Discretization of the Mass Balance

Integration of equation (7a) over the integration domain of a single control volume P , after the application of Gauss's integration theorem and the additivity of the Riemann integral, yields

$$\iint_{S_f} u_i n_i dS = \sum_{f \in \{w,s,b,t,n,e\}} \iint_{S_f} u_i n_i dS = 0.$$

In the present work the mass balance is discretized using the midpoint integration rule for the surface integrals and linear interpolation of the velocity to the center of mass of the surface. This leads to the following form of the mass balance:

$$\sum_{f \in \{w,s,b,t,n,e\}} u_{i,f} n_{fi} S_f = 0, \quad (12)$$

where no interpolation to attain the values of u_i at the face S_f has been performed yet, since the straightforward linear interpolation will lead to undesired oscillations in the solution fields. An interpolation method to circumvent this so called *checker-boarding* effect is presented in subsection 4.2.

4.2 A Pressure-Weighted Interpolation Method for Velocities

The advantages of using a cell-centered variable arrangement are evident: The treatment of non-orthogonality is simplified and the conservation property of finite-volume methods is retained [10, 35, 42, 63]. A major deficiency of cell centered variable arrangements is, that the pressure field may delink, which will then lead to unphysical oscillations in both the pressure and the velocity results. If the oscillations are severe enough, the solution algorithm might even get unstable and diverge. The described decoupling occurs, when the pressure gradient in the momentum balances, and the mass fluxes in the continuity equation are discretized using central differences.

A common practice to eliminate this behaviour is the use of a momentum interpolation technique, also known as *Rhie-Chow Interpolation* [51]. However, the original interpolation scheme does not guarantee a unique solution, independent of the amount of under-relaxation. The performance of one of the algorithms that are used in the present thesis heavily relies on the under-relaxation of variables to accomplish stability. Furthermore the original method as proposed by [51] does not account for large body forces which also may lead to unphysical results. These issues will be addressed in this subsection which, at the end, will present an interpolation method that assures an under-relaxation independent solution: the *pressure-weighted interpolation method* [42].

The starting point of the pressure-weighted interpolation method is formed by the discretized momentum balances at node P and an arbitrary neighboring node Q . The discretization for finite-volume methods and details including the incorporation of under-relaxation factors will be handled in subsection 4.5. The semi-discrete implicit momentum balances, if one solves for the velocity at node P or Q , read

$$u_{i,P}^{(n)} = -\frac{\alpha_{u_P}}{a_P^{u_i}} \left(\sum_{F \in NB(P)} a_F^{u_i} u_{i,F}^{(n)} + b_{P,u_i}^{(n-1)} - V_P \left(\frac{\partial p}{\partial x_i} \right)_P^{(n-1)} \right) + (1 - \alpha_u) u_{i,P}^{(n-1)}$$

and $u_{i,Q}^{(n)} = -\frac{\alpha_{u_Q}}{a_Q^{u_i}} \left(\sum_{F \in NB(Q)} a_F^{u_i} u_{i,F}^{(n)} + b_{Q,u_i}^{(n-1)} - V_Q \left(\frac{\partial p}{\partial x_i} \right)_Q^{(n-1)} \right) + (1 - \alpha_u) u_{i,Q}^{(n-1)}$,

where the superscript $(n-1)$ denotes the previous outer iteration number. The reader should note, that the pressure gradient has not been discretized yet. This has the advantage that the selective interpolation technique [55] can be applied, which is crucial for the elimination of the mentioned oscillations. In almost the same way a semi-discrete implicit momentum balance can be formulated for a virtual control volume located between nodes P and Q .

$$u_{i,f}^{(n)} = -\frac{\alpha_{u_f}}{a_f^{u_i}} \left(\sum_{F \in NB(f)} a_F^{u_i} u_{i,F}^{(n)} + b_{f,u_i}^{(n-1)} - V_f \left(\frac{\partial p}{\partial x_i} \right)_f^{(n-1)} \right) + (1 - \alpha_u) u_{i,f}^{(n-1)}. \quad (13)$$

Figure 5 gives an interpretation of the virtual control volume.




Figure 5: Possible interpretation of a virtual control volume (grey) located between nodes P and Q

To guarantee convergence of this expression for $u_{i,f}$, under-relaxation is necessary [35]. To eliminate the undefined artifacts surging from the virtualization of a control volume, the following assumptions have to be made to derive a closed expression for the velocity on the boundary face S_f :

$$\frac{\alpha_u}{a_f^{u_i}} \left(\sum_{F \in NB(f)} a_F^{u_i} u_{i,F}^{(n)} \right) \approx (1 - \gamma_f) \frac{\alpha_u}{a_p^{u_i}} \left(\sum_{F \in NB(P)} a_F^{u_i} u_{i,F}^{(n)} \right) + \gamma_f \frac{\alpha_u}{a_Q^{u_i}} \left(\sum_{F \in NB(Q)} a_F^{u_i} u_{i,F}^{(n)} \right), \quad (14a)$$

$$\frac{\alpha_u}{a_f^{u_i}} b_{f,u_i}^{(n-1)} \approx (1 - \gamma_f) \frac{\alpha_u}{a_p^{u_i}} b_{p,u_i}^{(n-1)} + \gamma_f \frac{\alpha_u}{a_Q^{u_i}} b_{Q,u_i}^{(n-1)} \quad (14b)$$

$$\text{and } \frac{\alpha_u}{a_f^{u_i}} \approx (1 - \gamma_f) \frac{\alpha_u}{a_p^{u_i}} + \gamma_f \frac{\alpha_u}{a_Q^{u_i}}, \quad (14c)$$

where γ_f is a geometric interpolation factor.

Using the assumptions made in equation (14) the expression in equation (13) can be closed in a way that it only depends on the variable values in node P and Q

$$\begin{aligned} u_{i,f}^{(n)} &\approx (1 - \gamma_f) \left(-\frac{\alpha_u}{a_p^{u_i}} \sum_{F \in NB(P)} a_F^{u_i} u_{i,F}^{(n)} \right) + \gamma_f \left(-\frac{\alpha_u}{a_Q^{u_i}} \sum_{F \in NB(Q)} a_F^{u_i} u_{i,F}^{(n)} \right) \\ &\quad + \frac{\alpha_u}{a_f^{u_i}} b_{f,u_i}^{(n-1)} - \frac{\alpha_u}{a_f^{u_i}} V_f \left(\frac{\partial p}{\partial x_i} \right)_f^{(n-1)} + (1 - \alpha_u) u_{i,f}^{(n-1)} \\ &= (1 - \gamma_f) u_{i,p}^{(n)} - (1 - \gamma_f) \frac{\alpha_u}{a_p^{u_i}} \left(b_{p,u_i}^{(n-1)} - V_p \left(\frac{\partial p}{\partial x_i} \right)_p^{(n-1)} \right) + \gamma_f u_{i,Q}^{(n)} - \gamma_f \frac{\alpha_u}{a_Q^{u_i}} \left(b_{Q,u_i}^{(n-1)} - V_Q \left(\frac{\partial p}{\partial x_i} \right)_Q^{(n-1)} \right) \\ &\quad + \frac{\alpha_u}{a_f^{u_i}} b_{f,u_i}^{(n-1)} - \frac{\alpha_u}{a_f^{u_i}} V_f \left(\frac{\partial p}{\partial x_i} \right)_f^{(n-1)} + (1 - \alpha_u) u_{i,f}^{(n-1)} \\ &= \left[(1 - \gamma_f) u_{i,p}^{(n)} + \gamma_f u_{i,Q}^{(n)} \right] \\ &\quad - \left[\left((1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_i}} \right) \left(\frac{\partial p}{\partial x_i} \right)_f^{(n-1)} - (1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_i}} \left(\frac{\partial p}{\partial x_i} \right)_p^{(n-1)} - \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_i}} \left(\frac{\partial p}{\partial x_i} \right)_Q^{(n-1)} \right] \end{aligned}$$

$$\begin{aligned}
& + (1 - \alpha_u) \left[u_{i,f}^{(n-1)} - (1 - \gamma_f) u_{i,P}^{(n-1)} - \gamma_f u_{i,Q}^{(n-1)} \right] \\
& \approx \left[(1 - \gamma_f) u_{i,P}^{(n)} + \gamma_f u_{i,Q}^{(n)} \right] \\
& - \left((1 - \gamma_f) \frac{\alpha_u V_P}{a_P^{u_i}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_i}} \right) \left[\left(\frac{\partial p}{\partial x_i} \right)_f^{(n-1)} - (1 - \gamma_f) \left(\frac{\partial p}{\partial x_i} \right)_P^{(n-1)} - \gamma_f \left(\frac{\partial p}{\partial x_i} \right)_Q^{(n-1)} \right] \\
& + (1 - \alpha_u) \left[u_{i,f}^{(n-1)} - (1 - \gamma_f) u_{i,P}^{(n-1)} - \gamma_f u_{i,Q}^{(n-1)} \right]. \tag{15}
\end{aligned}$$

It should be noted that the argumentation that led to the last expression, is, that the task of the underlined pressure gradient corrector in equation (15) is to suppress oscillations in the converged solution for the pressure field. If there are no oscillations this part should not become active. As long as the behaviour of this corrector remains consistent, i.e. that there are no oscillations in the pressure field within a converged solution, it can be multiplied with arbitrary constants [19]. This is however true on equidistant grids, where $\gamma_f = 1/2$ and central differences are used to calculate the gradients. On arbitrary non-orthogonal grids another modification has to be performed, which is based on a special case of the mean value theorem of differential calculus and the following

Proposition. Let $x_1, x_2 \in \mathbb{R}$ with $x_1 \neq x_2$ and $p(x) = a_0 + a_1 x + a_2 x^2$ a real polynomial function. Then

$$\frac{dp}{dx} \left(\frac{x_1 + x_2}{2} \right) = \frac{p(x_2) - p(x_1)}{x_2 - x_1},$$

i.e. the slope of the secant equals the value of the first derivative of p exactly half the way between x_1 and x_2 .

Proof. Evaluation of the derivative yields

$$\frac{dp}{dx} \left(\frac{x_1 + x_2}{2} \right) = a_1 + 2a_2 \frac{x_1 + x_2}{2} = a_1 + a_2(x_1 + x_2).$$

On the other hand the slope of the secant, using the third binomial rule can be expressed as

$$\begin{aligned}
\frac{p(x_2) - p(x_1)}{x_2 - x_1} &= \frac{a_0 + a_1 x_2 + a_2 x_2^2 - (a_0 + a_1 x_1 + a_2 x_1^2)}{x_2 - x_1} \\
&= \frac{a_1(x_2 - x_1) + a_2(x_2^2 - x_1^2)}{x_2 - x_1} \\
&= a_1 + a_2(x_2 + x_1).
\end{aligned}$$

The comparison of both expressions completes the proof. \square

It is desirable for the pressure corrector to vanish independent of the grid spacing if the profile of the pressure is quadratic and hence does not exhibit oscillations. According to the preceding proposition this can be accomplished by modifying equation (15) to average the pressure gradients from node P and Q instead of interpolating linearly

$$\begin{aligned}
u_{i,f}^{(n)} &= \left[(1 - \gamma_f) u_{i,P}^{(n)} + \gamma_f u_{i,Q}^{(n)} \right] \\
& - \left((1 - \gamma_f) \frac{\alpha_u V_P}{a_P^{u_i}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_i}} \right) \left[\left(\frac{\partial p}{\partial x_i} \right)_f^{(n-1)} - \frac{1}{2} \left(\frac{\partial p}{\partial x_i} \right)_P^{(n-1)} - \frac{1}{2} \left(\frac{\partial p}{\partial x_i} \right)_Q^{(n-1)} \right] \\
& + (1 - \alpha_u) \left[u_{i,f}^{(n-1)} - (1 - \gamma_f) u_{i,P}^{(n-1)} - \gamma_f u_{i,Q}^{(n-1)} \right]. \tag{16}
\end{aligned}$$

Comparing this final expression with the standard interpolation scheme, it is evident that, normally, the underlined term is not taken into consideration [19]. However section 7.4 shows, that neglecting this term indeed creates under-relaxation factor dependent results. This section concludes with a final

Proposition. The pressure weighted momentum interpolation scheme (16) guarantees the converged solution for $u_{i,f}$ to be independent of the velocity under-relaxation α_u .

Proof. An equivalent formulation of (16) is given by

$$\begin{aligned} \alpha_u u_{i,f}^{(n-1)} + u_{i,f}^{(n-1)} - u_{i,f}^{(n)} &= \alpha_u \left[(1 - \gamma_f) u_{i,p}^{(n-1)} + \gamma_f u_{i,Q}^{(n-1)} \right] \\ &+ \left[(1 - \gamma_f) (u_{i,p}^{(n)} - u_{i,p}^{(n-1)}) + \gamma_f (u_{i,Q}^{(n)} - u_{i,Q}^{(n-1)}) \right] \\ &- \alpha_u \left((1 - \gamma_f) \frac{V_p}{a_p^{u_i}} + \gamma_f \frac{V_Q}{a_Q^{u_i}} \right) \left[\left(\frac{\partial p}{\partial x_i} \right)_f^{(n-1)} - \frac{1}{2} \left(\frac{\partial p}{\partial x_i} \right)_p^{(n-1)} - \frac{1}{2} \left(\frac{\partial p}{\partial x_i} \right)_Q^{(n-1)} \right]. \end{aligned}$$

Upon convergence $u_{i,p}^{(n)} = u_{i,p}^{(n-1)}$ and $u_{i,f}^{(n)} = u_{i,f}^{(n-1)}$. This leads to

$$\begin{aligned} \alpha_u u_{i,f}^{(n-1)} &= \alpha_u \left[(1 - \gamma_f) u_{i,p}^{(n-1)} + \gamma_f u_{i,Q}^{(n-1)} \right] \\ &- \alpha_u \left((1 - \gamma_f) \frac{V_p}{a_p^{u_i}} + \gamma_f \frac{V_Q}{a_Q^{u_i}} \right) \left[\left(\frac{\partial p}{\partial x_i} \right)_f^{(n-1)} - \frac{1}{2} \left(\frac{\partial p}{\partial x_i} \right)_p^{(n-1)} - \frac{1}{2} \left(\frac{\partial p}{\partial x_i} \right)_Q^{(n-1)} \right], \end{aligned}$$

which shows, after division by $\alpha_u > 0$, that $u_{i,f}$ is independent of the under-relaxation factor. \square

4.3 Implicit Pressure Correction and the SIMPLE-Algorithm

Finite-volume methods aim at deducing a particular system of linear algebraic equations from a partial differential equation. In the case of the momentum balances the general structure of this linear equations is

$$u_{i,p}^{(n)} = -\frac{\alpha_{u_p}}{a_p^{u_i}} \left(\sum_{F \in NB(P)} a_F^{u_i} u_{i,F}^{(n)} + b_{p,u_i}^{(n-1)} - V_p \left(\frac{\partial p}{\partial x_i} \right)_p^{(n-1)} \right) + (1 - \alpha_u) u_{i,p}^{(n-1)},$$

where the pressure gradient has been discretized only symbolically and $b_{p,u_i}^{(n-1)}$ denotes the source term evaluated at the previous outer iteration.

At this stage the equations are still coupled and non-linear. As described in section 3.4, the Picard iteration process can be used to linearize the equations. After this, every momentum balance equation only depends on the one dominant variable u_i . Furthermore, the coupling of the momentum balances through the convective term ($u_i u_j$) is resolved in the process of linearization. In the next step, the decoupled momentum balances can be solved sequentially for the dominant variable u_i . All coefficients $a_{\{p,F\}}^{u_i}$, the source term and the pressure gradient will be evaluated explicitly by using results of the preceding outer iteration ($n - 1$). For the pressure gradient this means to take the pressure of the antecedent outer iteration as a first guess for the following iteration. This guess has to be corrected in an iterative way until all the non-linear equations are fulfilled up to a certain tolerance. Section (6.3.2) presents a suitable convergence criterion and its implementation. This linearization process in conjunction with the pressure guess leads to the linear equation

$$u_{i,p}^{(n*)} = -\frac{\alpha_u}{a_p^{u_i}} \left(\sum_{F \in NB(P)} a_F^{u_i} u_{i,F}^{(n*)} + b_{p,u_i}^{(n-1)} - V_p \left(\frac{\partial p}{\partial x_i} \right)_p^{(n-1)} \right) + (1 - \alpha_u) u_{i,p}^{(n-1)}. \quad (17)$$

Here (*) indicates, that the solution of this equation still needs to be corrected to also fulfill the discretized mass balance

$$\sum_{F \in NB(P)} (u_i)_f^{(n)} n_i S_f = 0. \quad (18)$$

Using the procedure that has been applied in section 4.2 to equation (17) results in the following expression for the face velocities after solving the discretized momentum balances

$$\begin{aligned}
u_{i,f}^{(n*)} = & \left[(1 - \gamma_f) u_{i,p}^{(n*)} + \gamma_f u_{i,Q}^{(n*)} \right] \\
& - \left((1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_i}} \right) \left[\left(\frac{\partial p}{\partial x_i} \right)_f^{(n-1)} - (1 - \gamma_f) \left(\frac{\partial p}{\partial x_i} \right)_p^{(n-1)} - \gamma_f \left(\frac{\partial p}{\partial x_i} \right)_Q^{(n-1)} \right] \\
& + (1 - \alpha_u) \left[u_{i,f}^{(n-1)} - (1 - \gamma_f) u_{i,p}^{(n-1)} - \gamma_f u_{i,Q}^{(n-1)} \right],
\end{aligned} \tag{19}$$

where the pressure from the previous outer iteration has been used as a pressure guess.

The lack of an equation having the pressure as dominant variable leads to the necessity to alter the mass balance as the only equation left. Methods of this type are called projection methods [19]. A common class of algorithms of this family of methods uses an equation for the additive pressure correction p' instead of the pressure itself and enforces continuity by correcting the velocities with an additive corrector u'_i , such that

$$u_{i,p}^{(n)} = u_{i,p}^{(n*)} + u'_{i,p}, \quad u_{i,f}^{(n)} = u_{i,f}^{(n*)} + u'_{i,f} \quad \text{and} \quad p_p^{(n)} = p_p^{(n-1)} + p'_p.$$

Now it is possible to formulate the discretized momentum balance for the corrected velocities and the corrected pressure as

$$u_{i,p}^{(n)} = -\frac{\alpha_u}{a_p^{u_i}} \left(\sum_{F \in NB(P)} a_F^{u_i} u_{i,F}^{(n)} + b_{p,u_i}^{(n-1)} - V_p \left(\frac{\partial p}{\partial x_i} \right)_p^{(n)} \right) + (1 - \alpha_u) u_{i,p}^{(n-1)}. \tag{20}$$

It should be noted that the only difference to the equation which will be solved in the next outer iteration is, that the source term b_{p,u_i} has not been updated yet. In the same way, the discretized momentum balance for the face velocity $u_{i,f}$ can be formulated as

$$\begin{aligned}
u_{i,f}^{(n)} = & \left[(1 - \gamma_f) u_{i,p}^{(n)} + \gamma_f u_{i,Q}^{(n)} \right] \\
& - \left((1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_i}} \right) \left[\left(\frac{\partial p}{\partial x_i} \right)_f^{(n)} - (1 - \gamma_f) \left(\frac{\partial p}{\partial x_i} \right)_p^{(n)} - \gamma_f \left(\frac{\partial p}{\partial x_i} \right)_Q^{(n)} \right] \\
& + (1 - \alpha_u) \left[u_{i,f}^{(n-1)} - (1 - \gamma_f) u_{i,p}^{(n-1)} - \gamma_f u_{i,Q}^{(n-1)} \right].
\end{aligned} \tag{21}$$

To couple velocity and pressure correctors one can subtract equation (17) from (20) and equation (19) from (21) to get

$$u'_{i,p} = -\frac{\alpha_u}{a_p^{u_i}} \left(\sum_{F \in NB(P)} a_F^{u_i} u'_{i,F} - V_p \left(\frac{\partial p'}{\partial x_i} \right)_p^{(n)} \right) \quad \text{and} \tag{22}$$

$$u'_{i,f} = \left[(1 - \gamma_f) u'_{i,p} + \gamma_f u'_{i,Q} \right] - \left((1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_i}} \right) \left[\left(\frac{\partial p'}{\partial x_i} \right)_f' - (1 - \gamma_f) \left(\frac{\partial p'}{\partial x_i} \right)_p' - \gamma_f \left(\frac{\partial p'}{\partial x_i} \right)_Q' \right]. \tag{23}$$

The majority of the class of pressure correction algorithms has these equations as a common basis. Each algorithm then introduces special, distinguishable approximations of the velocity corrections that are, at the moment of solving for the pressure correction, still unknown. The method used in the present work is the *SIMPLE*-algorithm (*Semi-Implicit Method for Pressure-Linked Equations*) [46]. The approximation this algorithm performs is severe, since the term containing the unknown velocity corrections is dropped entirely. The respective term has been underlined in equation (22). Since the global purpose of the presented method is to enforce continuity by implicitly calculating a pressure correction, the velocity correction has to be expressed solely in terms of the pressure correction. This can be accomplished by inserting equation (22) into equation (23). This gives an update formula

$$u'_{i,f} = - \left((1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_i}} \right) \left(\frac{\partial p'}{\partial x_i} \right)_f', \tag{24}$$

which is then, together with (19), inserted into the discretized continuity equation (18) to obtain

$$\sum_{F \in NB(P)} \left((1 - \gamma_f) \frac{\alpha_u V_P}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_F}{a_F^{u_i}} \right) \left(\frac{\partial p}{\partial x_i} \right)_f' n_i S_f = b_{Pp} \quad , \quad (25)$$

where the right hand side b_{Pp} is defined as

$$b_{Pp} := \sum_{F \in NB(P)} u_{i,f}^{(n*)} n_i S_f. \quad (26)$$

The complete discretization of this equation, using central differences as approximation for the gradient of the pressure correction, is straightforward and will be presented in subsection 4.4.

The approximation performed in the SIMPLE-algorithm affects convergence in a way, that the pressure correction has to be under-relaxed with a parameter $\alpha_p \in [0, 1]$

$$p_p^{(n)} = p_p^{(n-1)} + \alpha_p p_p'. \quad (27)$$

It should be noted that there exist better approximations for the pressure correction which do not rely on pressure under-relaxation for convergence. On example that uses a more consistent approximation is the *SIMPLEC*-algorithm (*SIMPLE Consistent*) [62]. A similar derivation of the pressure weighted interpolation method and its implementation for the SIMPLEC-algorithm can be found in [42].

As shown in section 4.2, the behaviour of the pressure weighted interpolation method on non-equidistant grids can be improved by replacing the linear interpolation of pressure gradients in equation (19) with a simple average. This leads to the following equation for calculating mass fluxes from face velocities

$$\begin{aligned} u_{i,f}^{(n*)} = & \left[(1 - \gamma_f) u_{i,p}^{(n*)} + \gamma_f u_{i,Q}^{(n*)} \right] \\ & - \left((1 - \gamma_f) \frac{\alpha_u V_P}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_i}} \right) \left[\left(\frac{\partial p}{\partial x_i} \right)_f^{(n-1)} - \frac{1}{2} \left(\frac{\partial p}{\partial x_i} \right)_P^{(n-1)} - \frac{1}{2} \left(\frac{\partial p}{\partial x_i} \right)_Q^{(n-1)} \right] \\ & + (1 - \alpha_u) \left[u_{i,f}^{(n-1)} - (1 - \gamma_f) u_{i,p}^{(n-1)} - \gamma_f u_{i,Q}^{(n-1)} \right]. \end{aligned} \quad (28)$$

Generally the SIMPLE-algorithm can be represented by an iterative procedure as shown in algorithm 4.3.

Algorithm 1 SIMPLE-Algorithm

```

INITIALIZE variables
while (convergence criterion not accomplished) do
    SOLVE linearized momentum balances, equation (17)
    CALCULATE mass fluxes using (21) or (28)
    SOLVE pressure correction equation to assure continuity, equation (25)
    UPDATE pressure using (27)
    UPDATE velocities and mass fluxes using (22)
    if (coupled scalar equation) then
        SOLVE scalar equation as described in (4.6)
    end if
end while
if (decoupled scalar equation) then
    SOLVE scalar equation as described in (4.6)
end if

```

4.4 Discretization of the Mass Fluxes and the Pressure Correction Equation

Subsections 4.2 and 4.3 introduced the concept of pressure weighted interpolation to avoid oscillating results and an algorithm to calculate a velocity field that obeys continuity. The derived equations have not been discretized completely, furthermore the approach has not been generalized for non-orthogonal grids.

The discretized mass balance (12) only depends on the normal velocities $u_{i,f} n_{i,f}$. By analogy with equation (28), an interpolated normal face velocity, and thus the mass flux, can be calculated as

$$\begin{aligned} u_{i,f}^{(n*)} n_{i,f} = & \left[(1 - \gamma_f) u_{i,p}^{(n*)} + \gamma_f u_{i,Q}^{(n*)} \right] n_{i,f} \\ & - \left((1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_i}} \right) \left[\left(\frac{\partial p}{\partial n} \right)_f^{(n-1)} - \frac{1}{2} \left(\frac{\partial p}{\partial n} \right)_p^{(n-1)} - \frac{1}{2} \left(\frac{\partial p}{\partial n} \right)_Q^{(n-1)} \right] \\ & + (1 - \alpha_u) \left[u_{i,f}^{(n-1)} - (1 - \gamma_f) u_{i,p}^{(n-1)} - \gamma_f u_{i,Q}^{(n-1)} \right] n_{i,f}, \end{aligned} \quad (29)$$

where the scalar product of pressure gradients and the normal vector has been replaced by a directional derivative in the direction of the face unit normal vector. In the present work, the pressure gradients in (29) and the pressure correction gradients in equation (25) will be discretized by central differences

$$\left(\frac{\partial p}{\partial n} \right)_f \approx \frac{p_p - p_Q}{(\mathbf{x}_p - \mathbf{x}_Q) \cdot \mathbf{n}_f} \quad \text{and} \quad \left(\frac{\partial p'}{\partial n} \right)_f \approx \frac{p'_p - p'_Q}{(\mathbf{x}_p - \mathbf{x}_Q) \cdot \mathbf{n}_f}.$$

In the next step, this discretization can be inserted into the semi-discretized pressure correction equation (25)

$$\sum_{F \in \text{NB}(P)} \left((1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_F}{a_F^{u_i}} \right) \frac{p'_p - p'_F}{(\mathbf{x}_p - \mathbf{x}_F) \cdot \mathbf{n}_f} S_f = b_{p,p}.$$

The resulting coefficients for the pressure correction equation,

$$a_p^{p'} p'_p + \sum_{F \in \text{NB}(P)} a_F^{p'} p'_F = b_{p,p'},$$

can be calculated as

$$a_F^{p'} = - \left((1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_F}{a_F^{u_i}} \right) \frac{S_f}{(\mathbf{x}_p - \mathbf{x}_F) \cdot \mathbf{n}_f} \quad \text{and} \quad a_p^{p'} = - \sum_{F \in \text{NB}(P)} a_F^{p'}. \quad (30)$$

The right hand side can be calculated as in equation (26), if the presented discretization is applied.

4.5 Discretization of the Momentum Balance

If the stationary momentum balance is integrated over a single control volume P , the result reads

$$\underbrace{\iint_S (\rho u_i u_j) n_j dS}_{\text{convective term}} - \underbrace{\iint_S \left(\mu \left(\frac{\partial u_i}{\partial x_j} + \frac{\partial u_j}{\partial x_i} \right) \right) n_j dS}_{\text{diffusive term}} = - \underbrace{\iiint_V \frac{\partial p}{\partial x_i} dV}_{\text{pressure sourceterm}} - \underbrace{\iiint_V \rho \beta (T - T_0) dV}_{\text{temperature sourceterm}}, \quad (31)$$

where the different terms, to be addressed individually in the following sections, are indicated. The reader should note, that the form of this equation has been modified by using Gauss's integration theorem. The terms residing on the left will be treated in an implicit and, due to deferred corrections, also in an explicit way, whereas the terms on the right will be treated exclusively in an explicit way.

4.5.1 Linearization and Discretization of the Convective Term

The convective term $(\rho u_i u_j)$ of the Navier-Stokes equations is the reason for the non-linearity of the equations. In order to deduce a set of linear algebraic equations from the Navier-Stokes equations, this term has to be linearized. As introduced in section (3.4), the non-linearity will be dealt with by means of an iterative process, the Picard iteration. Therefore, the part dependent on the non-dominant dependent variable, will be approximated with its value from the previous iteration as $\rho u_i^{(n)} u_j^{(n)} \approx \rho u_i^{(n)} u_j^{(n-1)}$. However, this linearization will not be directly visible, because it will be covered by the mass flux $\dot{m}_f = \iint_{S_f} \rho u_j^{(n-1)} n_j dS$.

Using the additivity of the Riemann integral, the first step is to decompose the surface integral into individual contributions from each boundary face of the control volume P

$$\iint_S \rho u_i u_j n_j dS = \sum_{f \in \{w,s,b,t,n,e\}} \iint_{S_f} \rho u_i u_j n_j dS = \sum_{f \in \{w,s,b,t,n,e\}} F_{i,f}^c,$$

where $F_{i,f}^c := \iint_{S_f} \rho u_i^{(n)} u_j^{(n-1)} n_j dS$ is the convective flux of the velocity u_i through the boundary face S_f .

To improve diagonal dominance of the resulting linear system, while maintaining the smaller discretization error of a higher order discretization, a blended discretization scheme is combined with a deferred correction. Since, due to the non-linearity of the equations to be solved, an iterative solution process is needed by all means, the overall convergence does not degrade noticeably by the application of a deferred correction [19]. Blending and deferred correction result in a decomposition of the convective flux into a lower order approximation, which is treated implicitly, and an explicit difference between the higher and lower order approximation for the same convective flux. Since for coarse grid resolutions, the use of higher order approximations may lead to oscillations of the solution, which in turn may degrade or even impede convergence, the schemes can be blended by a control factor $\eta \in [0, 1]$. Furthermore, the use of low-order approximations increases the diagonal dominance of the matrix, which can improve the convergence of certain iterative solvers for the linear systems [55].

To show the generality of this approach, all further derivations are presented for the generic boundary face S_f , which separates control volume P from its neighbor $F \in NB(P)$. This decomposition then leads to

$$F_{i,f}^c \approx \underbrace{F_{i,f}^{c,l}}_{\text{implicit}} + \eta \underbrace{[F_{i,f}^{c,h} - F_{i,f}^{c,l}]}_{\text{explicit}}^{(n-1)}. \quad (32)$$

It should be noted, that the convective fluxes carrying an l for *lower* or an h for *higher* as exponent, already have been linearized and discretized. The discretization applied to the convective flux in the present work is using the midpoint integration rule and blends the upwind interpolation scheme with a linear interpolation scheme. Using the decomposition in equation (32), one can derive the following approximations

$$\begin{aligned} F_{i,f}^{c,l} &= u_{i,F} \min(\dot{m}_f, 0) + u_{i,P} \max(0, \dot{m}_f) \\ F_{i,f}^{c,h} &= u_{i,F} \gamma_f + u_{i,P} (1 - \gamma_f), \end{aligned}$$

where the variable values have to be taken from the previous iteration step $(n-1)$ as necessary and the mass flux \dot{m}_f has been used as result of the linearization process. Now, the results can be summarized by presenting the convective contribution to the matrix coefficients $a_{F,u_i}^{u_i}$ and $a_P^{u_i}$ and the right hand side b_{P,u_i} as

$$\begin{aligned} a_F^{u_i,c} &= \min(\dot{m}_f, 0), \quad a_P^{u_i,c} = \sum_{F \in NB(P)} \max(0, \dot{m}_f) \\ \text{and } b_{P,u_i}^c &= \sum_{F \in NB(P)} \eta \left(u_{i,F}^{(n-1)} (\min(\dot{m}_f, 0) - \gamma_f) \right) + \eta \left(u_{i,P}^{(n-1)} (\max(0, \dot{m}_f) - (1 - \gamma_f)) \right). \end{aligned}$$

4.5.2 Discretization of the Diffusive Term

The diffusive term contains the first partial derivatives of the velocity as a result of the material constitutive equation that characterizes the behaviour of Newtonian fluids. As pointed out in section 3.3, directional derivatives can be discretized using central differences on orthogonal grids or, in the more general case of non-orthogonal grids, using central differences implicitly and an explicit deferred correction comprising the non-orthogonality of the grid. As seen in equation (6) the diffusive term of the Navier-Stokes equations can be simplified using the mass balance in the case of an incompressible flow with constant viscosity μ . To sustain the generality of the presented approach, this simplification will be omitted.

As before, by using the additivity and furthermore linearity of the Riemann integral, the integration of the diffusive term will be divided into integration over individual boundary faces S_f

$$\iint_S \left(\mu \left(\frac{\partial u_i}{\partial x_j} + \frac{\partial u_j}{\partial x_i} \right) \right) n_j dS = \sum_{f \in \{w,s,b,t,n,e\}} \left[\iint_{S_f} \mu \frac{\partial u_i}{\partial x_j} n_j dS + \iint_{S_f} \mu \frac{\partial u_j}{\partial x_i} n_j dS \right] = \sum_{f \in \{w,s,b,t,n,e\}} F_{i,f}^d,$$

where $F_{i,f}^d$ denotes the diffusive flux through an individual boundary face. Section 3.3 only covered the non-orthogonal corrector for directional derivatives. Since the velocity is a vector field and not a scalar field, the results of section 3.3 may only be applied to the underlined term. The other term will be treated explicitly since it is considerably smaller than the underlined term, furthermore does not cause oscillations, and thus will not derogate convergence [19]. First, all present integrals will be approximated using the midpoint integration rule. The diffusive flux $F_{i,f}^d$ for a generic face S_f located between the control volumes P and F reads

$$F_{i,f}^d \approx \mu \left(\frac{\partial u_i}{\partial x_j} \right)_f n_j S_f + \mu \left(\frac{\partial u_j}{\partial x_i} \right)_f n_j S_f.$$

Using central differences for the implicit discretization of the directional derivative and using the *orthogonal correction* approach from 3.3.2, the approximation can be derived as

$$\begin{aligned} F_{i,f}^d &\approx \mu \left(\frac{\|\Delta_f\|_2}{\|\mathbf{x}_P - \mathbf{x}_F\|_2} \frac{u_{P_i} - u_{F_i}}{\|\mathbf{x}_P - \mathbf{x}_F\|_2} - (\nabla u_i)_{f_i}^{(n-1)} \cdot (\Delta_f - \mathbf{S}_f) \right) + \mu \left(\frac{\partial u_j}{\partial x_i} \right)_f^{(n-1)} n_{f_i} \\ &= \mu \left(S_f \frac{u_{P_i} - u_{F_i}}{\|\mathbf{x}_P - \mathbf{x}_F\|_2} - \left(\frac{\partial u_i}{\partial x_j} \right)_f^{(n-1)} (\xi_{f_i} - n_{f_i}) S_f \right) + \mu \left(\frac{\partial u_j}{\partial x_i} \right)_f^{(n-1)} n_{f_i}, \end{aligned}$$

where the unit vector pointing in direction of the straight line connecting control volume P and control volume F is denoted as

$$\xi_f = \frac{\mathbf{x}_P - \mathbf{x}_F}{\|\mathbf{x}_P - \mathbf{x}_F\|_2}.$$

The interpolation of the cell center gradients to the boundary faces is performed as in (9). Now, the contribution of the diffusive part to the matrix coefficients and the right hand side can be calculated as

$$\begin{aligned} a_F^{u_i,d} &= -\frac{\mu S_f}{\|\mathbf{x}_P - \mathbf{x}_F\|_2}, \quad a_P^{u_i,d} = \sum_{F \in NB(P)} \frac{\mu S_f}{\|\mathbf{x}_P - \mathbf{x}_F\|_2} \\ b_{F,u_i}^d &= -\sum_{F \in NB(P)} \left(\frac{\partial u_i}{\partial x_j} \right)_f^{(n-1)} (\xi_{f_i} + n_{f_i}) S_f - \mu \left(\frac{\partial u_j}{\partial x_i} \right)_f^{(n-1)} n_{f_i} S_f \\ &= -\sum_{F \in NB(P)} \left(\frac{\partial u_i}{\partial x_j} \right)_f^{(n-1)} \xi_{f_i} S_f + \mu \left(\left(\frac{\partial u_i}{\partial x_j} \right)_f^{(n-1)} + \left(\frac{\partial u_j}{\partial x_i} \right)_f^{(n-1)} \right) n_{f_i} S_f. \end{aligned}$$

4.5.3 Discretization of the Source Terms

Since in the segregated solution approach in every equation all other variables but the dominant one are treated as constants and the source terms in equation (31) do not depend on the dominant variable, the discretization is straightforward. The source terms of the momentum balance are discretized using the midpoint integration rule for volume integrals, which leads to the source term

$$-\iiint_V \frac{\partial p}{\partial x_i} dV - \iiint_V \rho \beta (T - T_0) dV \approx -\left(\frac{\partial p}{\partial x_i} \right)_P^{(n-1)} V_P - \rho \beta (T_P^{(n-1)} - T_0) V_P = b_{P,u_i}^{sc}.$$

4.6 Discretization of the Temperature Equation

The discretization of the temperature equation is performed by the same means as for the momentum balance. The only difference is a simpler diffusion term. The integral form of the temperature equation, after applying the Gauss's theorem of integration, is

$$\underbrace{\iint_S \rho u_j T n_j dS}_{\text{advective term}} - \underbrace{\iint_S \kappa \frac{\partial T}{\partial x_j} n_j dV}_{\text{diffusive term}} = \underbrace{\iiint_V q_T dV}_{\text{source term}}.$$

Now, proceeding as in the previous subsections, one can discretize the advective term, the diffusive term and the source term. Since this process does not provide further insight, only the final results will be presented. The discretization yields the matrix coefficients as

$$\begin{aligned} a_F^T &= \min(\dot{m}_f, 0) + \frac{\kappa S_f}{\|\mathbf{x}_p - \mathbf{x}_F\|_2} \\ a_P^T &= \sum_{F \in NB(P)} \max(0, \dot{m}_f) - \frac{\kappa S_f}{\|\mathbf{x}_p - \mathbf{x}_F\|_2} \\ b_{PT} &= \sum_{F \in NB(P)} \eta \left(T_F^{(n-1)} (\min(\dot{m}_f, 0) - \gamma_f) \right) + \eta \left(T_P^{(n-1)} (\max(0, \dot{m}_f) - (1 - \gamma_f)) \right) \\ &\quad + \sum_{F \in NB(P)} \left(\frac{\partial T}{\partial x_j} \right)_f^{(n-1)} (\xi_{f_j} - n_{f_j}) S_f \\ &\quad + q_{T_P} V_P. \end{aligned}$$

Again it is possible, though not always as necessary as in the case of the velocities, to under-relax the solution of the resulting linear system with a factor α_T . This can be accomplished in the way that will be shown in subsection 4.9.

4.7 Boundary Conditions

As the antecedent subsections showed, it is possible to deduce a linear algebraic equation for each control volume by the finite-volume method. However, the approach presented in the preceding subsections did not cover the treatment of control volumes at domain boundaries. This subsection introduces the boundary conditions which are relevant for the present work and deals with transitional conditions at block boundaries.

4.7.1 Dirichlet Boundary Conditions

The first boundary condition is the Dirichlet boundary condition. This type of boundary condition, when used for the velocity or transported scalars, is used to model inlet conditions for flow problems. For the temperature equation it may also be used at walls, as will be shown in subsection 4.7.2. The Dirichlet boundary condition is characterized by explicitly specifying the value of the variable for which the equation is solved. As a result, boundary fluxes can be calculated directly. In the case of Dirichlet boundary conditions for the velocities, especially the mass flux \dot{m}_f is known and hence does not have to be calculated using the pressure weighted interpolation method. This furthermore implies that the boundary velocities do not have to be corrected at Dirichlet velocity boundaries, i.e. (24) yields

$$0 = u'_{i,f} = - \left((1 - \gamma_f) \frac{\alpha_u V_P}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_i}} \right) \left(\frac{\partial p}{\partial x_i} \right)'_f,$$

which implies

$$0 = \left(\frac{\partial p}{\partial x_i} \right)'_f.$$

This means, that for the pressure correction at Dirichlet velocity boundaries, a Neumann zero gradient boundary condition applies.

Since in the case of Dirichlet boundary conditions for the velocities or the temperature no special modifications have to be made to the discretized equations, because the resulting coefficient for a neighboring control volume laying past

the boundary is considered on the right hand side of the linear system, the implementation approach for inlet boundary conditions will be presented only for the temperature equation. Furthermore there is no boundary condition that fixes the gradient at Dirichlet boundaries, this is why it is assumed, that the partial derivatives of the respective variable are constant and can be extrapolated

$$\left(\frac{\partial T}{\partial x_j}\right)_f \approx \left(\frac{\partial T}{\partial x_j}\right)_p.$$

The modification of the central coefficient of the linear equation can be recursively formulated as

$$a_p^T = a_p^T + \left(\max(0, \dot{m}_f) - \frac{\kappa S_f}{\|\mathbf{x}_p - \mathbf{x}_f\|_2} \right), \quad (33)$$

whereas the contribution to the right hand side reads

$$b_{pT} = b_{pT} - \left(\min(\dot{m}_f, 0) - \frac{\kappa S_f}{\|\mathbf{x}_p - \mathbf{x}_f\|_2} \right) T_f + \left(\frac{\partial T}{\partial x_j} \right)_p^{(n-1)} (\xi_{fj} - n_{fj}) S_f. \quad (34)$$

The reader should note, that, even though the gradient discretization at domain boundaries is realized by a one sided forward differencing scheme instead of a central differencing scheme, this does not drastically affect accuracy, because the distance used in the differential quotient is half the distance used in a central difference for the inside of the domain [55].

The case of Dirichlet boundary conditions for the pressure needs special consideration since no equation for the pressure is solved but instead an equation for the pressure correction. This type of boundary condition can be applied to inlet boundaries or pressure boundaries and is recommendable if on the one hand no data on the velocity distribution is known, but on the other hand the pressure difference between inlet and outlet, or the reference pressure at the outlet is a known quantity. The pressure correction at Dirichlet pressure boundaries is assumed to $p' = 0$, and the modification of the coefficients of the algebraic equation for the corresponding control volumes can be carried out as in (33) and (34). The velocity correction at the boundaries is nonzero and will have to be extrapolated, using the formula for the pressure weighted interpolation applied to boundary values (23). Even though Dirichlet boundary conditions for the pressure would correspond to Neumann boundary conditions for the velocities, a Dirichlet boundary condition for the velocities is used as well, using an extrapolated velocity value as fixed value for the calculation of boundary fluxes in the momentum balances.

4.7.2 Treatment of Wall Boundaries

A common boundary for the solution domain is given by solid walls. For all kind of flows this boundary condition first of all has a kinematic character, since the concept of impermeable walls dictates a zero normal velocity at the wall. In viscous flows, wall boundaries can be interpreted furthermore as a no-slip condition, i.e. a Dirichlet boundary condition for the tangential velocities. Convective fluxes through solid walls are thus zero by definition, however, the diffusive fluxes of the velocities and the temperature require special consideration. To approximate the fluid behavior on a wall boundary correctly, special modifications have to be taken into account to model the normal a shear tension. Furthermore, diffusive fluxes for the temperature can be given by Neumann or Dirichlet boundary conditions.

The derivation of the discretized diffusive flux through wall boundaries starts from the integral momentum balance (2) for the vector \mathbf{u} . Here, only the term for surface forces is needed

$$\mathbf{F}_w = \iint_{S_w} \mathbf{t} dS = \iint_{S_w} \mathbf{T}(\mathbf{n}_w) dS.$$

For the purpose of treating wall boundary conditions, it is appropriate to use a local coordinate system n, t, s , where n denotes the wall normal coordinate, t denotes the coordinate tangential to the wall shear force and, s is the binormal coordinate. With respect to this coordinate system, the wall normal vector is represented by $\mathbf{n}_w = (1, 0, 0)^T$ and the image of the wall normal vector $\mathbf{T}(\mathbf{n}_w)$ is represented by

$$\mathbf{T}(\mathbf{n}_w) = \begin{bmatrix} \tau_{nn} & \tau_{nt} & \tau_{ns} \\ \tau_{nt} & \tau_{tt} & \tau_{ts} \\ \tau_{ns} & \tau_{ts} & \tau_{ss} \end{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} \tau_{nn} \\ \tau_{nt} \\ \tau_{ns} \end{bmatrix}$$

The velocity of the wall is assumed to be constant [55], what implies that the directional derivative of the tangential velocity vanishes on the wall

$$\tau_{tt} = \frac{\partial u_t}{\partial x_t} = 0,$$

which, together with the continuity equation in differential form (1), leads to

$$\frac{\partial u_n}{\partial x_n} + \frac{\partial u_t}{\partial x_t} = \frac{\partial u_n}{\partial x_n} = 0, \quad (35)$$

what is equivalent to $\tau_{nn} = 0$ at the wall. It should be noted that the velocity is assumed to have no component in the binormal direction. A physical interpretation of (35) would be, that the transfer of momentum at the wall occurs by shear forces exclusively. Furthermore, the coordinate direction t is chosen to be parallel to the shear force, which does not represent a restriction, because of the possibility to rotate the coordinate system within the plane. This leads to $\tau_{ns} = 0$. Hence, the absolute value of the surface force only depends on the normal derivative of the velocity tangential to the wall. After transforming the coordinates back to the system (x_1, x_2, x_3) the surface force can be calculated and the integral can be discretized using the midpoint integration rule

$$\mathbf{F}_w = \iint_{S_w} \mathbf{t}_w \tau_{nt} dS = \iint_{S_w} \mathbf{t}_w \mu \frac{\partial u_t}{\partial n} dS = \mathbf{t}_w \mu \left(\frac{\partial u_t}{\partial n} \right)_w S_w, \quad (36)$$

where \mathbf{t}_w denotes the transformed tangential vector $(0, 1, 0)^T$ with respect to the coordinate system (x_1, x_2, x_3) . In the discretization process, this tangential vector will be calculated from the velocity vector as

$$\mathbf{t}_w = \frac{\mathbf{u}_t}{\|\mathbf{u}_t\|_2}, \quad \text{where} \quad \mathbf{u}_t = \mathbf{u} - (\mathbf{u} \cdot \mathbf{n}_w) \mathbf{n}_w.$$

According to [19] this force should not be handled explicitly for convergence reasons. On the other side, if the surface force is expressed in terms of the velocities u_i , the discretization process would lead to different central matrix coefficients, which in turn would negatively affect memory efficiency.

The discretization approach used in the present thesis uses a simpler implicit discretization coupled with a deferred correction, that combines the explicit discretization of the surface force (36) with a central difference. At first, the contained directional derivative is discretized implicitly by

$$\left(\frac{\partial u_t}{\partial n} \right)_{t_{wi}} \approx \left(\frac{\partial u_i}{\partial \xi} \right) \approx \frac{u_{i,p} - u_{i,w}}{\|\mathbf{x}_p - \mathbf{x}_w\|_2}$$

and explicitly by

$$\left(\frac{\partial u_t}{\partial n} \right)_{t_{wi}} \approx \frac{(u_{i,p} - u_{i,w}) - (u_{j,p} - u_{j,w}) n_j n_i}{(\mathbf{x}_p - \mathbf{x}_w) \cdot \mathbf{n}_w}.$$

Therefore, the contributions to the central coefficient and the right hand side of the linear equation, that results from the presented discretization process, are

$$a_p^{u_i} = a_p^{u_i} + \mu \frac{S_f}{\|\mathbf{x}_p - \mathbf{x}_w\|_2} \quad \text{and} \quad b_{p,u_i} = b_p^{u_i} + \mu \frac{S_f}{\|\mathbf{x}_p - \mathbf{x}_w\|_2} u_{i,p}^{(n-1)} + \frac{(u_{i,p}^{(n-1)} - u_{i,w}) - (u_{j,p}^{(n-1)} - u_{j,w}) n_j n_i}{(\mathbf{x}_p - \mathbf{x}_w) \cdot \mathbf{n}_w}.$$

The reader should note that, since the deferred correction uses a Dirichlet boundary condition, no correction of the value of the wall velocity $u_{i,w}$ has to be accounted for.

If the solution of the flow field is coupled to the solution of a temperature equation, different options for the boundary condition at walls may be chosen. If on the one hand the wall temperature is known, a Dirichlet boundary condition for the temperature is the choice. A wall of this type is called to be *isothermal*. If on the other hand, only the heat flux is known, a Neumann boundary condition is used. In the special case of zero heat flux the wall is called to be *adiabatic*. For adiabatic walls, which are besides isothermal walls used for the present thesis, the implementation is straightforward, since no coefficients have to be calculated at all.




Figure 6: Non-matching grid cells with hanging nodes at a two-dimensional block boundary. Indexing is based on the face segments S_l

4.7.3 Treatment of Block Boundaries

If block-structured grids are used to decompose the problem domain, the characterizing property of structured grids, which is the constant amount of grid cells in each direction, gives each inner cell exactly six neighbors. If more than one grid block is used, this property is violated, if the number of grid cells of each block is chosen to be arbitrary. The arbitrariness of the grid resolution is a main benefactor for the adaptivity of block-structured grids. Arbitrary grid resolutions may lead to *hanging* nodes at block boundaries, which are not addressed by the discretization techniques, which have been presented in the preceding chapters. How a solution algorithm for partial differential equations via finite-volume methods can deal with block boundaries, is presented in this subsection.

To maintain the conservation property [55] of finite-volume methods, block boundaries should not be interpreted as boundaries in the classical sense. Instead, conditions have to be formulated to guarantee, that flows through block boundaries are conserved. The method to treat block boundaries, which is used in the present thesis, was presented by [33]. A different method was presented in [32]. However, the first method was chosen, since it allows for a fully implicit discretization of the block boundary fluxes. This method calculates fluxes through separate face segments S_l , that surge from *non-matching* grid blocks. Each of these face segments gets assigned a control volume L and a control volume R , which exclusively share this face segment. An algorithm to provide this needed information has been implemented in the present thesis. Figure 6 shows an example of non-matching block boundaries and the used indexing for control volumes and boundary faces, which is based a numbering of the face segments S_l .

Once the geometric data, including the information regarding the interpolation coefficients at block boundaries, has been provided, the fluxes through these boundaries can be calculated as presented in the antecedent subsections. The coefficients in the linear system corresponding to the connectivity of neighboring control volumes of different grid blocks is represented by matrix coefficients a_L^ϕ and a_R^ϕ , where $\phi \in \{u_1, u_2, u_3, p, T\}$.

4.8 Treatment of the Singularity of the Pressure Correction Equation with Neumann Boundaries

It has to be noted, that the derived pressure correction equation is a Poisson equation. As can be proven [25], the linear $N \times N$ systems surging from the presented discretization on a grid with N control volumes with Neumann boundary conditions have a nullspace of dimension one, i.e.

$$\text{null}(A_{p'}) = \text{span}(\mathbb{1}),$$

where $\mathbb{1} = (1)_{i=1, \dots, N} \in \mathbb{R}^N$ is the vector spanning the nullspace. This singularity accounts for the property of incompressible flows, that pressure can only be determined up to a constant. To fix this constraint, various possibilities exist [19]. A common method is to set the pressure correction to zero in one reference control volume which corresponds to fixing

the pressure at one reference point in the problem domain. This can be done before solving the system by applying this Dirichlet-type condition, or it can be done afterwards when pressure is calculated from the pressure correction. It has to be noted, that this approach is not suitable for grid convergence studies since, without special interpolation, it is not guaranteed that the reference pressure correction is taken at the correct location throughout the different grid levels.

Since some comparisons performed in the present work rely on grid convergence studies, another approach for reducing the lose constraint of the pressure correction system has been used: The reference pressure correction is taken to be the mean value of the pressure correction over the domain

$$p'_{\text{ref}} = \frac{\iiint_V p' dV}{\iiint_V dV} \approx \frac{\sum_{P=1}^N p'_P V_P}{\sum_{P=1}^N V_P},$$

such that the net pressure correction amounts to zero. This modifies equation (27) to read

$$p_P^{(n)} = p_P^{(n-1)} + \alpha_p (p'_P - p'_{\text{ref}}).$$

It should be noted, that as soon as one control volume has been provided with a Dirichlet pressure boundary condition, the systems is not singular anymore. This is for example the case of a fixed pressure at inlet or outlet boundary conditions, which has been presented in section 4.7.1.

4.9 Structure of the Assembled Linear Systems

The objective of a finite-volume method is to create a set of linear algebraic equations from a discretization of partial differential equations. In the case of the discretized momentum balance, taking all contributions together, this leads to the following linear algebraic equation for each control volume P

$$a_p^{u_i} u_{p_i} + \sum_{F \in NB(P)} a_F u_{F_i} = b_{p_{u_i}},$$

where the coefficients are composed as

$$\begin{aligned} a_p^{u_i} &= a_p^{u_i,c} - a_p^{u_i,d} \\ a_F^{u_i} &= a_F^{u_i,c} - a_F^{u_i,d} \\ b_{p_{u_i}} &= b_{p_{u_i},c} - b_{p_{u_i},d} + b_{p_{u_i}}^{sc}. \end{aligned}$$

Similar expressions for the pressure correction equation and the temperature equation can be deduced. In the case of control volumes located at boundaries some of the coefficients will be calculated in a different way. This aspect has been addressed in section 4.7.

For the decoupled iterative solution process of the Navier-Stokes equations it is necessary to reduce the change of each dependent variable in each iteration. Normally this is done by an *under-relaxation* technique, a convex combination of the solution of the linear system for the present iteration (n) and from the previous iteration ($n-1$), with the under-relaxation parameter $\alpha_{u_i} \in (0, 1]$, where $\alpha_{u_i} = 1$ corresponds to no under-relaxation. Generally speaking, this parameter can be chosen individually for each equation. Since there are no rules for choosing this parameters in a general setting, the under-relaxation parameter for the velocities is chosen to be equal for all three velocities, $\alpha_{u_i} = \alpha_u$ [55]. This has the further advantage that, in case the boundary conditions are implemented with the same intention, the linear system for each of the velocities remains unchanged except for the right hand side. This helps to increase memory efficiency.

Let the solution for the linear system without under-relaxation be denoted as

$$\tilde{u}_{p_i}^{(n)} := \frac{b_{p_{u_i}} - \sum_{F \in NB(P)} a_F u_{F_i}}{a_p^{u_i}},$$

which is only a formal expression for the case, in which the underlying linear system is solved exactly. The described convex combination yields

$$\begin{aligned} u_{p_i}^{(n)} &:= \alpha_u \tilde{u}_{p_i}^{(n)} + (1 - \alpha_u) u_{p_i}^{(n-1)} \\ &= \alpha_u \frac{b_{p_i} - \sum_{F \in NB(P)} a_F u_{F_i}}{a_p^{u_i}} + (1 - \alpha_u) u_{p_i}^{(n-1)}, \end{aligned}$$

an expression that can be modified to derive a linear system whose solution is the under-relaxed velocity

$$\frac{a_p^{u_i}}{\alpha_u} u_{i,p} + \sum_{F \in NB(P)} a_F^{u_i} u_{i,F} = b_{p_i} + \frac{(1 - \alpha_u) a_p^{u_i}}{\alpha_u} u_{i,p}^{(n-1)}.$$

After all matrix coefficients have been successfully calculated, the linear system can be represented by a system matrix A and a right hand side vector b . Figure 6.3.4 shows the non-zero structure of a linear system for a grid consisting of a $2 \times 2 \times 2$ cell block and a $3 \times 3 \times 3$ cell block.




Figure 7: Non-zero structure of the linear systems used in the SIMPLE-algorithm for a block-structured grid consisting of one $2 \times 2 \times 2$ cell block and one $3 \times 3 \times 3$ cell block

5 Implicit Finite-Volume Method for Incompressible Flows – Fully Coupled Approach

Since the antecedent section 4 already discussed the discretization details of the involved equations, this section aims at a comparison on the algorithmic level of the SIMPLE-algorithm, presented in section 4.3, with an implementation of a fully coupled solution algorithm. It should be noted, that the discretization of the equations to be solved is not changed in any way to maintain comparability, so the presented differences are due to differences in the solution algorithm exclusively. Successful implementations of a fully coupled solution algorithm for incompressible Navier-Stokes equations have been presented in [9, 12, 18, 61]. The presented work will extend the solution approach presented in [18] to three-dimensional domains. Furthermore, this section will present various approaches to incorporate different degrees of velocity-to-temperature and temperature-to-velocity/pressure coupling, as they are presented in [61]. Finally, the structure of the resulting linear system to be solved is discussed.

5.1 The Fully Coupled Algorithm – Pressure-Velocity Coupling Reconsidered

This subsection motivates the use of a fully coupled algorithm to address pressure-velocity coupling and mentions the differences to the approach presented in subsection 4.3. The mentioned subsection presented a common solution approach to solve the incompressible Navier-Stokes equations. After linearizing the equations, the momentum balances are solved, using the pressure from the previous iteration as a guess. Generally the velocity field obtained by solving a momentum balance with a guessed pressure does not obey continuity, hence the velocity field and the pressure field have to be corrected. This in turn leads to an inferior solution with respect to the residual of the momentum balances. To avoid this iterative guess-and-correct solution process, another class of approaches dealing with the pressure-velocity coupling problematic, represented by algorithms that are *fully coupled*, will be introduced now.

The central aspect of fully coupled solution methods for Navier-Stokes equations is, that instead of solving for the velocities and pressure corrections sequentially, the velocity field and the pressure are solved for simultaneously [55]. This way, every velocity field that is calculated, obeys conservation of momentum and mass, without the need to be corrected. As a result, the under-relaxation due to the resolution of the pressure-velocity coupling is no longer required, which accelerates convergence significantly. The only reason for still needing an iterative solution process, is the non-linearity of the Navier-Stokes equations, which is accounted for by the use of the Picard iteration process, which has been presented in subsection 3.4. The iterations fortunately make room for deferred corrections as introduced in subsection 4.5. It should be noted that it is also possible to use the coupled solution algorithm to calculate pressure corrections. Reference [30] uses this approach to solve the Navier-Stokes equations by using a SIMPLE-type method as preconditioner.

On the downside, implementations of fully coupled solution methods require significantly more system memory than segregated methods. This is due to the higher amount of information, that has to be available at the same time, and the resulting linear systems, which also require more memory, due to the increased amount of unknowns to solve for. Furthermore, the bad condition of the linear algebraic system [55] tends to slow down the convergence of the equation solver algorithm.

As in the case of segregated methods, it is not advisable to solve the mass balance equation directly [55]. Instead of deriving an equation for the pressure correction p' , as shown in 4.3, an equation for the pressure itself is derived, using the pressure-weighted interpolation method, which has been introduced in 4.2. Concretely equation (16) is adapted for the use in a fully coupled algorithm

$$u_{i,f}^{(n)} = \left[(1 - \gamma_f) u_{i,p}^{(n)} + \gamma_f u_{i,Q}^{(n)} \right] - \left((1 - \gamma_f) \frac{V_p}{a_p^{u_i}} + \gamma_f \frac{V_Q}{a_Q^{u_i}} \right) \left[\left(\frac{\partial p}{\partial x_i} \right)_f^{(n)} - \frac{1}{2} \left(\left(\frac{\partial p}{\partial x_i} \right)_p^{(n-1)} + \left(\frac{\partial p}{\partial x_i} \right)_Q^{(n-1)} \right) \right].$$

The adaption comprises the removal of last term of equation (16) accounting for the under-relaxation, since according to [12] no under-relaxation is needed in fully coupled algorithms, which is equivalent to an under-relaxation factor $\alpha_u = 1$. Furthermore, the underlined partial derivatives are going to be treated implicitly. This leads to the semi-implicit pressure equation

$$\begin{aligned} & \sum_{F \in NB(P)} \rho \left[(1 - \gamma_f) u_{i,p}^{(n)} + \gamma_f u_{i,F}^{(n)} \right] S_f - \rho \left((1 - \gamma_f) \frac{V_p}{a_p^{u_i}} + \gamma_f \frac{V_F}{a_F^{u_i}} \right) \left[\left(\frac{\partial p}{\partial x_i} \right)_f^{(n)} \right] S_f \\ & = - \sum_{F \in NB(P)} \rho \left((1 - \gamma_f) \frac{V_p}{a_p^{u_i}} + \gamma_f \frac{V_F}{a_F^{u_i}} \right) \left[\frac{1}{2} \left(\left(\frac{\partial p}{\partial x_i} \right)_p^{(n-1)} + \left(\frac{\partial p}{\partial x_i} \right)_F^{(n-1)} \right) \right] S_f. \end{aligned}$$

Even though this equation is solved for the pressure p , while equation (25) is solved for the pressure correction p' , the matrix coefficients from the discretization of the pressure correction equation (30) and the calculation of the right hand

side given in equation (26), remain the same. For the implicit coupling to the velocities u_i , additional matrix coefficients have to be considered. These can be calculated as

$$a_F^{p,u_i} = \rho \gamma_f S_f \quad \text{and} \quad a_p^{p,u_i} = \sum_{F \in NB(P)} \rho (1 - \gamma_f) S_f.$$

On the other side, the discretized momentum balance yields matrix coefficients to account for the implicit velocity-pressure coupling. They are calculated as

$$a_F^{u_i,p} = \gamma_f S_f \quad \text{and} \quad a_p^{u_i,p} = \sum_{F \in NB(P)} (1 - \gamma_f) S_f.$$

If the Boussinesq approximation is used implicitly, an additional coefficient $a_p^{u_i,T}$ to account for the velocity temperature coupling has to be considered. This will be handled in subsection 5.2. The main steps of the fully coupled solution algorithm are shown in algorithm 5.1.

Algorithm 2 Fully Coupled Solution Algorithm

```

INITIALIZE variables
while (convergence criterion not accomplished) do
  if (temperature coupling) then
    SOLVE the linear system for velocities, pressure and temperature
  else
    SOLVE the linear system for velocities and pressure
  end if
  CALCULATE mass fluxes using (16)
  if (coupled scalar equation) then
    SOLVE scalar equation as described in section(4.6)
  end if
end while
if (decoupled scalar equation) then
  SOLVE scalar equation as described in section(4.6)
end if

```

5.2 Coupling to the Temperature Equation

Since the present work also aims at analyzing the efficiency of different methods to couple the temperature equation to the velocities and vice-versa, this subsection discusses different approaches to handle the coupling of the temperature equation to the Navier-Stokes equations if the Boussinesq approximation, as introduced in subsection 2.5.2, is used. The effectiveness of realizing a strong coupling of the temperature equation to the Navier-Stokes equations depends on the flow problem. The physical coupling tends to increase in flow scenarios of natural convection, where the temperature difference, and hence the buoyancy term in the Navier-Stokes equations, dominates the fluid movement [19, 61]. It is supposed that, the higher the physical coupling of temperature and flow, the greater the benefits of using an implementation that uses a strong coupling to the temperature equation. The following subsections present different approaches that can be used for different intensities of coupling.

5.2.1 Decoupled Approach – Explicit Velocity-to-Temperature Coupling

The decoupled approach for addressing velocity-to-temperature coupling is similar to the treatment of the velocity-to-temperature coupling described in subsection 4.5. If this approach is chosen, no special measures have to be taken. The momentum balances receive a contribution $b_p^{u_i,T}$ with

$$b_p^{u_i,T} := -\rho \beta \left(T_p^{(n-1)} - T_0 \right) g_i V_P,$$

that handles the coupling explicitly by using the temperature result of the previous outer iteration. In some cases it might be necessary to under-relax the temperature each iteration [19].

5.2.2 Implicit Velocity-to-Temperature Coupling

In the same way, it is possible to realize the coupling described in the previous subsection by implicitly considering the temperature in the momentum balances, which leads to an additional matrix coefficient $a_p^{u_i, T}$ and an additional contribution to the right hand side, accounting for the coupling to the temperature equation,

$$a_p^{u_i, T} = \rho \beta g_i V_p \quad \text{and} \quad b_p^{u_i, T} = \rho \beta T_0 g_i V_p.$$

5.2.3 Temperature-to-Velocity/Pressure Coupling – Newton-Raphson Linearization

The previous section discussed the velocity-to-temperature coupling and so far, only the momentum balances were affected by the realization of the coupling methods. Independently, it is possible to couple the temperature equation, not only to the momentum balances, but also to the pressure equation. The purpose of this subsection is to present a discretization that achieves this opposite type of coupling, which will be referred to as *temperature-to-velocity/pressure* coupling.

In section 4.6, the non-linear partial differential equations were linearized using a Picard iteration method. Specifically the convective term $\rho u_j T$ was linearized by taking the mass flux from the antecedent solve of the momentum and pressure correction equations. If the decoupled approach is used, this treatment remains valid. But, if an implicit temperature coupling is sought, a coupling of the temperature equation to the momentum balances which considers the mass fluxes is desirable. Methods that exhibit the described property can be found in [20, 45, 56, 61]. A common denomination of the therein used linearization method is called *Newton-Raphson linearization*, which can be interpreted as a bilinear approximation of the convective term by the linear terms of the respective Taylor polynomial, similar to Newton-like methods to solve non-linear equations [19]. It should be noted that, even though Newton-like methods are known for their faster, up to quadratic convergence compared to methods with linear convergence like the Picard iteration, the intention of using a Newton-like linearization is to create coupling to accelerate convergence. For a Newton-like linearization to accelerate convergence, instead of only the convective term, the whole equation should use a Newton method for the non-linear solution algorithm. For this, the *Newton-Raphson linearization* has not been used within segregated, SIMPLE-like solution algorithms [19].

The Newton-Raphson linearization technique is applied to the convective term of the temperature equation as follows. A first order Taylor approximation of the mass specific convective flux through the boundary face S_f around the $(n-1)$ st iteration value yields for the approximation at the n th iteration

$$\begin{aligned} (u_{j,f} T_f)^{(n)} &\approx (u_{j,f} T_f)^{(n-1)} + \frac{\partial}{\partial u_{j,f}} (u_{j,f} T_f)^{(n-1)} (u_{j,f}^{(n)} - u_{j,f}^{(n-1)}) + \frac{\partial}{\partial T_f} (u_{j,f} T_f)^{(n-1)} (T_f^{(n)} - T_f^{(n-1)}) \\ &= (u_{j,f} T_f)^{(n-1)} + T_f^{(n-1)} (u_{j,f}^{(n)} - u_{j,f}^{(n-1)}) + u_{j,f}^{(n-1)} (T_f^{(n)} - T_f^{(n-1)}) \\ &= \underline{u_{j,f}^{(n-1)} T_f^{(n)}} + u_{j,f}^{(n)} T_f^{(n-1)} - u_{j,f}^{(n-1)} T_f^{(n-1)}. \end{aligned}$$

A comparison with section 4.5.1 shows, that the underlined term coincides with the term from the usual linearization. The first of the two new terms will be treated implicitly and explicitly after using the pressure-weighted interpolation method from section 4.2 to interpolate the value of $u_{j,f}$. The second term will be treated explicitly. The use of the pressure-weighted interpolation method does not only create a temperature-to-velocity but also a temperature-to-pressure coupling. Applying the pressure-weighted interpolation, the equation for the mass specific convective flux reads

$$\begin{aligned} (u_{j,f} T_f)^{(n)} &\approx \underline{u_{j,f}^{(n-1)} T_f^{(n)}} + T_f^{(n-1)} \left[(1 - \gamma_f) u_{j,p}^{(n)} + \gamma_f u_{j,Q}^{(n)} \right] \\ &\quad - T_f^{(n-1)} \left((1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_j}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_j}} \right) \left[\left(\frac{\partial p}{\partial x_j} \right)_f^{(n)} - \frac{1}{2} \left(\frac{\partial p}{\partial x_j} \right)_p^{(n-1)} - \frac{1}{2} \left(\frac{\partial p}{\partial x_j} \right)_Q^{(n-1)} \right] \\ &\quad - u_{j,f}^{(n-1)} T_f^{(n-1)}. \end{aligned}$$

If this semi-implicit, semi-discrete equation is discretized completely, new matrix and right hand side contributions can be calculated

$$\begin{aligned}
a_F^{T,u_i} &= \rho T_f^{(n-1)} \gamma_f n_{f,i} S_f \\
a_p^{T,u_i} &= \sum_{F \in NB(P)} \rho T_f^{(n-1)} (1 - \gamma_f) n_{f,i} S_f \\
a_F^{T,p} &= -\rho T_f^{(n-1)} \left((1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_i}} + \gamma_f \frac{\alpha_u V_F}{a_F^{u_i}} \right) \frac{S_f}{(\mathbf{x}_p - \mathbf{x}_F) \cdot \mathbf{n}_f} = \rho T_f^{(n-1)} a_F^p \\
\text{and } a_p^{T,p} &= - \sum_{F \in NB(P)} T_f^{(n-1)} a_F^{T,p}.
\end{aligned}$$

The explicit parts resulting from the Newton-Raphson linearization yield additional contributions to the right hand side

$$b_p^{T,NR} = \rho u_{j,f}^{(n-1)} T_f^{(n-1)} n_{f,j} S_f - \rho T_f^{(n-1)} \left((1 - \gamma_f) \frac{\alpha_u V_p}{a_p^{u_j}} + \gamma_f \frac{\alpha_u V_Q}{a_Q^{u_j}} \right) \left[\frac{1}{2} \left(\frac{\partial p}{\partial x_j} \right)_p^{(n-1)} + \frac{1}{2} \left(\frac{\partial p}{\partial x_j} \right)_Q^{(n-1)} \right] n_{j,f} S_f.$$

5.3 Boundary Conditions on Domain and Block Boundaries

The treatment of boundary conditions and transitional conditions at block boundaries is very similar to the way presented in section 4.7. The main difference results from the fact, that, in addition to the necessary changes from section 4.7, modifications are necessary to the new matrix coefficients from the discretized momentum balance and the pressure equation in the coupled solution algorithm. The purpose of this section is therefore to only focus on these additional changes.

How Dirichlet boundary conditions, also known as inlet boundaries, affect the matrix coefficients that only correspond to the velocities has been presented in section 4.7.1. It is common to use Neumann boundary conditions for the pressure at Dirichlet boundaries with respect to the velocities [12]. Under the assumption that the gradient of the pressure vanishes at inlet boundaries, the matrix coefficients for the velocity-to-pressure coupling have to be modified to

$$a_p^{u_i,p} = a_p^{u_i,p} + n_{f,i} S_f.$$

Wall boundary conditions are treated likewise with the additional modifications to the right hand side contribution as presented in section 4.7.2.

The treatment of velocity Dirichlet boundary conditions in the pressure equation is straightforward, since they correspond to mass fluxes through the boundary faces. Hence, the matrix coefficients remain unchanged and only the right hand side contributions embrace the fact that no implicit interpolation is necessary to calculate the mass flux through the boundary face.

In the case of Dirichlet pressure boundary conditions, an additional contribution to the right hand side of the momentum balances has to be considered

$$b_p^{u_i} = b_p^{u_i} + p_f n_{i,f} S_f.$$

As presented in subsection 4.7 a *lagged* Dirichlet boundary condition is used for the velocities in the discretized momentum balances. Different to the presentation in subsection 4.7 the boundary mass fluxes in the pressure equation are considered implicitly. This creates additional contributions to the coefficients accounting for the pressure-to-velocity coupling in the pressure equation

$$a_p^{p,u_i} = a_p^{p,u_i} + \rho n_{f,i} S_f.$$

The coefficients for the pressure are modified as usual for Dirichlet boundary conditions, which was shown in section 4.7.

Finally, it should be noted that the Newton-Raphson linearization technique does not apply to the velocities on boundary faces, since the value of the velocities is known exactly and a linearization of the convective contribution becomes unnecessary.

At block boundaries the new matrix coefficients $a_L^{u_i,p}$ and $a_R^{u_i,p}$ for the velocity-to-pressure coupling and the matrix coefficients a_L^{p,u_i} and a_R^{p,u_i} have to be taken into account. This does not present any drawbacks on the presented approach.

5.4 Assembly of Linear Systems – Final Form of Equations

Concluding this section, the final form of the derived linear algebraic equations are presented. Based on the presented discretization, each equation only contains variable values of neighboring control volumes. The set of five equations reads

$$a_p^{u_i} u_{p,i} + \sum_{F \in NB(P)} a_F^{u_i} u_{F,i} + \underbrace{a_p^{u_i,p} p_p + \sum_{F \in NB(P)} a_F^{u_i,p} p_F}_{\text{Pressure-velocity coupling}} + \underbrace{a_p^{u_i,T} T_p}_{\text{Boussinesq approximation}} = b_{p,u_i} \quad i = 1, \dots, 3$$

$$a_p^p p_p + \sum_{F \in NB(P)} a_F^p p_F + \underbrace{\sum_{j=1}^3 \left(a_p^{p,u_j} u_{p,j} + \sum_{F \in NB(P)} a_F^{p,u_j} u_{F,j} \right)}_{\text{Pressure-velocity coupling}} = b_{p,p}$$

$$a_p^T T_p + \sum_{F \in NB(P)} a_F^T T_F + \underbrace{\sum_{j=1}^3 \left(a_p^{T,u_j} u_{p,j} + \sum_{F \in NB(P)} a_F^{T,u_j} u_{F,j} \right)}_{\text{Newton-Raphson linearization}} + a_p^{T,p} p_p + \sum_{F \in NB(P)} a_F^{T,p} p_F = b_{p,T}.$$

The different degrees of coupling between the different equations will affect the matrix structure. Depending on the variable, arrangement other matrix structures are possible. Figure 8 shows the resulting block matrix structure if the fields are treated as interlaced quantities. This variable arrangement corresponds to a linear sub system for each control volume:

$$\begin{bmatrix} a_p^{u_1} & 0 & 0 & a_p^{u_1,p} & a_p^{u_1,T} \\ 0 & a_p^{u_2} & 0 & a_p^{u_2,p} & a_p^{u_2,T} \\ 0 & 0 & a_p^{u_3} & a_p^{u_3,p} & a_p^{u_3,T} \\ a_p^{p,u_1} & a_p^{p,u_2} & a_p^{p,u_3} & a_p^p & 0 \\ a_p^{T,u_1} & a_p^{T,u_2} & a_p^{T,u_3} & a_p^{T,p} & a_p^T \end{bmatrix} \begin{bmatrix} u_{1,p} \\ u_{2,p} \\ u_{3,p} \\ p_p \\ T_p \end{bmatrix} + \sum_{F \in NB(P)} \begin{bmatrix} a_F^{u_1} & 0 & 0 & a_F^{u_1,p} & a_F^{u_1,T} \\ 0 & a_F^{u_2} & 0 & a_F^{u_2,p} & a_F^{u_2,T} \\ 0 & 0 & a_F^{u_3} & a_F^{u_3,p} & a_F^{u_3,T} \\ a_F^{p,u_1} & a_F^{p,u_2} & a_F^{p,u_3} & a_F^p & 0 \\ a_F^{T,u_1} & a_F^{T,u_2} & a_F^{T,u_3} & a_F^{T,p} & a_F^T \end{bmatrix} \begin{bmatrix} u_{1,F} \\ u_{2,F} \\ u_{3,F} \\ p_F \\ T_F \end{bmatrix} = \begin{bmatrix} b_p^{u_1} \\ b_p^{u_2} \\ b_p^{u_3} \\ b_p^p \\ b_p^T \end{bmatrix}.$$

Figure 9 shows the non-zero structure of a block matrix for the same case if field interlacing is used. It should be noted, that the variable arrangement only affects the storage of the matrices not the way the matrix is used, since it is possible to arbitrarily reorder the matrix entries through the underlying index sets to obtain other matrix representations. One commonly used matrix representation is shown in figure 10.




Figure 8: Non-zero structure of block submatrices of the linear systems used in the coupled solution algorithm for a block-structured grid consisting of one $2 \times 2 \times 2$ cell block and one $3 \times 3 \times 3$ cell block. The blue coefficients represent the pressure-velocity coupling, the red coefficients correspond to the velocity-to-temperature coupling due to the implicit Boussinesq approximation and the green coefficients result from the Newton-Raphson linearization technique addressing the temperature-to-velocity/pressure coupling.




Figure 9: Non-zero structure of the linear system used in the coupled solution algorithm for a block-structured grid consisting of one $2 \times 2 \times 2$ cell block and one $3 \times 3 \times 3$ cell block. The variables have been interlaced and the matrix consists of blocks as shown in figure 8.




Figure 10: Non-zero structure of the linear system used in the coupled solution algorithm for a block-structured grid consisting of one $2 \times 2 \times 2$ cell block and one $3 \times 3 \times 3$ cell block. The variables have been ordered in a way that each major matrix block refers to only one variable or the coupling between exactly two variables as it is denoted by the corresponding subindices.

6 CAFFA Framework

This section presents the CAFFA framework that has been extended for the present thesis. *CAFFA* is an acronym, which stands for *Computer Aided Fluid Flow Analysis*. The solver framework is based on Perić's CAFFA code [19] that has been modified to handle three-dimensional problem domains and block-structured grids with non-matching blocks. Furthermore, the framework has been parallelized making extensive use of the PETSc library. One characteristic of the solver framework is, that arbitrary block boundaries are handled in a fully implicit way. Details on how those transitional conditions are handled can be found in section 4.7.3. The framework consists of two different solver algorithms, a segregated solver and a fully coupled solver, whose theoretical basis has been introduced in sections 4 and 5 respectively. In addition, the framework features an own implementation of a grid generator for block-structured grids as well as a mapper program to convert grids generated by a commercial meshing tool into a format understandable by the CAFFA code. To prepare the grid data for later use in the solvers, a preprocessor program has been integrated into the framework. Furthermore the solver programs exhibit different export functionalities to visualize the numerical grid and the results within ParaView [27] and the export of binary vectors and matrices to MATLAB ® [38].

The following subsections will now briefly introduce the PETSc framework, present the components of the CAFFA framework.

6.1 PETSc Framework

PETSc is an acronym for *Portable Extensible Toolkit for Scientific Computation* and is a software suite, which comprises data structures and an extensive set of linear solver routines [5, 6]. A great part of the data structures is designed to work in parallel on high performance computers, also the solver routines have been parallelized using the MPI standard for message passing. PETSc provides the tools, that can be used to build large-scale application codes for scientific calculations. The great advantage of PETSc is, that it adds multiple additional layers of abstraction. This makes the use of PETSc in applications straightforward and allows the users to focus on the essential implementations of their own code instead of having to deal with parallel data structures and algorithms, since this represents an additional source for programming errors. PETSc provides different interfaces through which the user can parallelize his code, what furthermore increases readability and maintainability of the developed code. Last but not least, the implementations used for data structures and solvers in PETSc are not only verified and updated on a regular basis, they have furthermore proven to be highly efficient in the creation of large-scale computer applications on high performance clusters [7, 24, 29].

The data structures of PETSc reflect the commonly used elements in scientific codes. Vectors, matrices and index sets are among these basic data structures. Beyond that, PETSc offers a variety of options for its data structures. Vectors can either be used sequentially or in parallel, in which case PETSc distributes the vector components to the processes inside a specified MPI communicator. Furthermore parallel matrix formats are available that allow the processes to assemble different parts of one global matrix simultaneously. PETSc contains other data structures that can be used to fetch and distribute either vector or matrix elements from remote processes.

On the other side PETSc contains a large collection of parallel solvers for linear systems. The major part comprises Krylov subspace methods. An overview over classical Krylov subspace methods can be found in [52]. Among the variety of solver options that are available for further optimization of the performance, the Krylov subspace methods can be combined with elements of another thorough collection of preconditioners and other third party packages like Hypre [17] or ML [21]. For further details on the available data structures and subroutines the reader is referred to [4, 5].

In addition to the mentioned tools, PETSc also comes with an internal profiler tool, that collects a variety of information on the used PETSc objects and bundles them into a human readable log file. The log file and other dynamic output, which can be triggered by command line arguments, facilitate the optimization process of developed applications and allow the accurate determination of performance measures.

6.2 Grid Generation and Preprocessing

A grid generator has been developed to generate hexahedral block-structured grids. To provide flexibility for testing locally refined and skewed grids the grid generator deploys a random number generator which moves grid points within the domain and varies the number of grid cells of each block optionally. Key feature of the developed framework is the handling of block-structured locally refined grids with non-matching block interfaces. The neighboring relations are represented by a special type of boundary condition.

After the grid has been generated, a preprocessor program comes to application, which prepares the numerical grid for the use in the solver program. An essential part of the preprocessing step is the application of various matching algorithms depending on the type of block boundary condition. The preprocessor program differentiates between three block boundary types of which each is a subtype of the other. The most simple block boundary condition is given by a one-to-one relation between two neighboring boundary cells. The generalisation of this condition involves several boundary cells, that share only one neighboring boundary cell in the neighboring block. This relation is also known as a many-to-one relation. The third boundary condition on block boundaries implements the most general, also known as

many-to-many, relation, which needs special consideration in the preprocessing step, because the number of neighbors one cell can have is not known a priori. What makes the third type of boundary conditions more difficult than the first two, is the geometry of the resulting boundary faces. The current implementation uses a polygon clipper to calculate the polygons resulting from the intersection of two boundary faces of neighboring blocks. Based on the polygon of the intersection, all further geometric data like interpolation factors and boundary face area can be calculated for the later use within the solve application.

6.3 Implementation of CAFFA

This section provides an overview of the implementation aspects of the developed CAFFA framework. Since applications for modern computers have to be able to efficiently use the provided resources, the concept for the parallelization of the application is presented. A separate section presents, how convergence is monitored and controlled. Later on, this section presents how the PETSc data structures that store the variables are used to realize the data composition as central part of the parallelization process.

6.3.1 The Message-Passing Model

The PETSc framework makes use of the message-passing computational model [23]. This model assumes, that a set of processes with local memory is able to communicate with other processes by passing, i.e. sending and receiving, messages. *MPI* is an acronym for *Message Passing Interface* and is a software library that collects features of message-passing systems. More importantly, it provides a widely used standard that assures portability of application that make use of this library as does PETSc.

The message-passing model is not only reflected in the application as a programming paradigm, but also in the design of parallel computers and their communication network. Under this perspective the message-passing model is an analogue to the distributed memory design, which encapsulates the idea of local memory and the exchange of data through a network. High performance computers are seldom designed exclusively as a distributed memory system. They rather additionally incorporate features of another model, namely the shared memory model. This memory model is characterized by a common address space, which each of the processes can access uniformly, i.e. with a location independent latency, or non-uniformly. Latter systems are also called *NUMA (Non-Uniform Memory Access)* systems. A widely used system design, as used in the system the performance tests of section 8 are performed on, comprises features of both models by using NUMA computing nodes that share local memory and are able to exchange messages with other NUMA nodes via an interconnect and the passing of messages.

One important advantage of using software that uses the message passing model and hence applications that were parallelized within this programming paradigm, is their compatibility with computers that use combinations of distributed and shared memory.

6.3.2 Convergence Control

One part of the PETSc philosophy is to provide maximum flexibility in choosing solvers and parameters from within the command line. However, the developed framework implements a non-linear iteration process without using the *SNES (Scalable Nonlinear Equations Solvers)* objects provided by PETSc. This creates the need for a hard-coded implementation of convergence control. Convergence control for the used Picard iteration method comprises two parts. One part controls the overall convergence, and determines when the calculations have finished. The other part controls convergence on an outer iteration basis. The convergence criteria are met if the actual residual falls below the upper bound for the residual r_{final} . The bound for final convergence is hereby determined to

$$r_{final} = r_{initial} * 10^{-8},$$

which states that a decrease of the initial residual of 10^{-8} indicates convergence. This criterion will be fixed for all further analyses since it controls the overall accuracy of the solver results and hence maintains comparability of solver performance.

The other convergence criterion can be handled with more flexibility, since it controls the relative decrease of the residual in each outer iteration. This parameter should not affect the overall accuracy but instead convergence speed, within the individual bounds on the solver algorithm. This means that on the one hand, depending on the coupling algorithm, low relative solver tolerances will not always benefit convergence speed. On the other hand, low relative solver tolerances will affect the number of inner iterations and hence the execution time of the solver program. The convergence criterion for each outer iteration is implemented as

$$r_{final}^{(k)} = r_{initial}^{(k)} * rtol,$$

where rtol indicates the relative decrease of the residual in each outer iteration.

6.3.3 Indexing of Variables and Treatment of Boundary Values

All variables needed by the CAFFA solver are represented by one-dimensional arrays. To establish a mapping between a location (i, j, k) of the numerical grid and the respective variable values location (ijk) inside the array the following formula is used

$$(ijk) = N_i N_j (k - 1) + N_j (i - 1) + j,$$

where N_i and N_j are the numbers of grid cells in the respective coordinate direction plus two additional boundary cells. The inclusion of the boundary cells circumvents the need to provide an additional data structure for boundary values. Furthermore, the same mapping rule can be used for the grid vertex coordinates. This variable indexing will be used to assemble the matrices surging from the discretization process. This will lead to rows for the boundary values, that that are decoupled from the rest of the linear system. PETSc provides two different approaches to handle boundary values. In the first approach another abstraction layer is introduced via an object, that redistributes the data and removes the rows containing the boundary values. This approach has been tested with the conclusion that the redistribution not only causes significant overhead but also hides important data for debugging solver convergence. In the present work, a second approach is used, that retains boundary values and adapts the right hand side accordingly. This approach has proven to be more efficient with the drawback of memory overhead and the need to reset the boundary values after a linear system has been solved with a high relative tolerance.

The presented indexing model is not capable of handling block-structured grids or grids that have been distributed across various processors within the MPI programming model. This property is implemented through the use of additional mappings that return a constant offset $(ijk)_b$ for each block and $(ijk)_p$ for each processor

$$(ijk) = (ijk)_p + (ijk)_b + N_i N_j (k - 1) + N_j (i - 1) + j.$$

The introduction of the processor offset is needed to provide a mapping between locally and globally indexed values. Within a process only the local indexing

$$(ijk)_{loc} = (ijk)_b + N_i N_j (k - 1) + N_j (i - 1) + j$$

will be used, since these indexes can be mapped directly to memory addresses. For inter process communication however, global indices have to be provided to the respective PETSc subroutines by adding the processor offset

$$(ijk)_{glo} = (ijk)_{loc} + (ijk)_p. \quad (37)$$

It should be noted that the presented mappings do not include the index mappings from FORTRAN applications, which use 1-based indexing, to the PETSc subroutines which use 0-based indices.

As presented in section 5, the use of a coupling algorithm leads to a global linear system that contains the linear algebraic equations for different variables. The implementation used for the present thesis interlaces this values, which increases data locality and hence improves memory efficiency. The presented mapping is easily extended to handle arrays of values that contain n_{eq} variables and use an interlaced storing approach by multiplying (37) with the number of interlaced values

$$(ijk)_{interglo} = n_{eq} * ((ijk)_{glo} - 1) + c_{eq},$$

where $c_{eq} = 1, \dots, n_{eq}$.

6.3.4 Domain Decomposition and the Exchange of Ghost Values

The developed CAFFA framework is able to solve a given problem across different processors. For this, before each solve, the relevant data has to be distributed among all involved processes. Within the solver parallelization, three types of this data distribution are considered with respect to their frequency of communication and redistribution among all or only a part of the processes involved in the solution process.

The first type of data refers to the distribution of stationary data that will not change throughout the solution process. This refers mostly to problem geometry related data, as for example the coordinates of the grid points. The distribution of the data is already part of the domain decomposition process and hence can be done even before the solver starts, as it is implemented in the developed solver framework. It should be noted that in the case, in which local dynamic refinement strategies are used, the geometric data at block boundaries will have to be redistributed. If additionally dynamic load balancing is used, even geometric data from inside the blocks of the problem domain might have to be redistributed. This however imposes other design guidelines for the involved processes which lay outside of the scope of the present thesis.




Figure 11: Storage and update of ghost values in vectors related to variables on multi block domains. The blocks have been assigned to two different processes, *Proc 1* and *Proc 2*. The control volumes of the two-dimensional problem domain are indexed with respect to the process local indexing.

The distribution of stationary data takes place throughout the preprocessing program within the developed framework before the CAFFA solver is launched. Each processor is assigned a binary file with the respective data. Since the solver is not able to handle adaptive grid refinement or dynamic load balancing this approach is straightforward.

The second type of data has to be distributed regularly, since this data changes at least every outer iteration. The necessity to interchange values, also known as *ghosting*, between processors surges when calculations of variable gradients or coefficients are made for block boundary control volumes. PETSc provides special interfaces for ghosted vectors that allow to perform different types of vector updates through top-level subroutines. The central part are the involved vector scatter and gather routines which allow to calculate contributions to source terms of control volumes located at block boundaries accurately or to update ghost values for the control volumes that rely on variable values from neighboring blocks to calculate fluxes or gradients. It should be noted that the modeling approach for ghosting values is not symmetric in the sense that each block saves ghost values of the neighboring blocks. On the contrary, for each block boundary only one of the neighbored blocks is responsible for the calculation of the fluxes and gradient contributions. This maintains the same amount of data communicated through the interconnect while almost splitting the computational effort into half.

Another advantage of the PETSc environment for ghosting values that reside in the memory space of other processes is the unified approach of data coherence, since on the top level, only one vector object to store all data related to a variable, e.g. the velocity u_1 , has to be created. Through precise information provided in the creation step of these vector objects, PETSc distributes the vector object to all involved processes and uses the provided information to manage ghosting processes through the PETSc internal vector scatter and gather routines.

The present solver framework treats the calculation of matrix coefficients for block boundary control volumes in a special way. To reduce communication overhead and redundant data, only one of the two involved processors calculates this coefficients and then sends them to the respective neighbor when the matrix is about to be assembled. The data needed to calculate the matrix coefficients of a neighboring processor embraces not only the values of the dominant variables from the last outer iteration but also the values of the cell center gradients. For the ghosting of variable values, PETSc offers special vector types, that comprise a user friendly interface to use the gather and scatter routines provided by PETSc. During the creation of these vectors, a set of global indices referring to the values that are to be ghosted, has to be provided. After this, the interchange of values is realized through a single subroutine call to `VecGhostUpdate`, since PETSc handles the communication process, which is necessary to accomplish the value interchange. Hence, the local representation of a ghosted vector, not only comprises the data local to one processor, but also provides space for ghosted values that are updated repeatedly. The parallel vector layout and the process of ghosting are schematically shown in figure 11 for a two dimensional domain consisting of two grid blocks distributed among two processes.

The last type of data distribution refers to global reduce operations, in which a single processes gathers data from all other processes and, after some modification, redistributes the data. A common scenario for global reduce operations is the calculation of the mean pressure throughout the domain, which has been introduced in section 4.8.

7 Verification of the developed CAFFA Framework

The systematical verification of program code is an essential part of the development process of software for scientific calculations, since it insures that the respective equations are solved correctly [44]. Scientific applications that solve partial differential equations can be verified using the method of manufactured solutions [54]. In addition, the method of manufactured solutions will be used to proof that both solvers within the developed framework, namely the segregated and the fully coupled solver, use the same discretization of the underlying partial differential equations. Before the performance of the developed CAFFA framework will be discussed in section 8, this section presents the results of the verification process performed in the present work via the *Method of Manufactured Solutions*. This section begins by mentioning the main aspects of the method of manufactured solutions. The subsections that follow present a concrete manufactured solution and the results of using it within the verification process. Finally, the results of a grid convergence study regarding the influence of the pressure weighted interpolation method from section 4.2 will be shown.

7.1 The Method of Manufactured Solutions for Navier-Stokes Equations

The method of manufactured solutions comprises a systematical, formal procedure for code verification based on analytic solutions to the partial differential equations to be solved [54]. Using this analytic solutions, the accuracy of the produced results can be assessed and, after the establishment of an acceptance criterion, used to verify the computer program. A common acceptance criterion is the *Order-of-Accuracy* criterion, since, in addition to the formal order of accuracy of a developed solver algorithm, it verifies its consistency. It should be noted that by the method of manufactured solutions it is not possible to detect errors in the physical model.

The basic idea of the method of manufactured solutions is the inversion of the solution process of partial differential equations. Instead of trying to find the solution x to a given equation $F(x) = b$ with source term b , a solution x is *manufactured* and the source term is constructed by applying F to x . This has the advantage of choosing a solution that exercises all parts of the solution process and hence provides a thorough testing environment. On the downside, the integration of boundary conditions other than Dirichlet boundary conditions might be challenging [54]. Furthermore, the program code has to be able to handle arbitrary source terms.

The different guidelines that have to be followed in order to apply the method of manufactured solutions successfully, can be found in [54]. According to the reference, manufactured solutions should be composed of infinitely often differentiable analytic functions whose derivatives are bounded by small constants. Furthermore, the solution domain should not be chosen to be symmetric, but as arbitrary as possible within the limitations of the code. This may however conflict with the application of the method of manufactured solutions to verify a code to solve the Navier-Stokes equations. Subsection 7.2 presents the manufactured solution including the resulting source terms for the set of partial differential equations (7).

The verification of a solver for incompressible Navier-Stokes equations comes with additional conditions, that have to be considered: First, if the solver is not able to handle source terms in the continuity equation and respectively the pressure correction equation, a velocity field should be chosen that is inherently divergence free. This can be achieved by defining the velocity field \mathbf{u} through the curl of a vector field Ψ as

$$\mathbf{u} = \nabla \times \Psi.$$

Using the property that the divergence of the rotation of a vector field vanishes, one gets

$$\nabla \cdot \mathbf{u} = \nabla \cdot (\nabla \times \Psi) = 0,$$

i.e. a locally divergence free velocity field, which implies that the continuity equation is fulfilled also globally in an integral sense for arbitrary domains of integration. This is not necessarily true in the discrete sense, as integrals are to be approximated by, in the case of the present thesis, the midpoint rule of integration. In the general case for arbitrary solution domains, global mass conservation cannot be guaranteed, even though the integrands are known quantities at the domain boundaries. Thus in the case of finding manufactured solutions, the problem domain should be fixed before manufacturing a solution. The velocity field for the manufactured solution should then either vanish on the domain boundaries, which is the case of the commonly used Taylor-Green vortex [59], or should exhibit further symmetry that leads to cancelling non-zero mass fluxes across the domain boundaries. The studies performed within the scope of the present showed, that convergence of the pressure correction equation can no longer be guaranteed if the velocity field does not obey continuity in the discrete sense at the domain boundaries.

7.2 Manufactured Solution for the Navier-Stokes Equations and the Temperature Equation

This subsection presents the manufactured solution for the system of partial differential equations (7). The manufactured solutions and the computations needed to derive the respective source terms were performed by the computer algebra

system Maple ® [37]. After the computation, the terms were directly translated into FORTRAN source code, such that the manufactured solution could be directly integrated into the solver framework. In the following paragraphs, the used solutions for the velocity vector (u_1, u_2, u_3) , the pressure p and the temperature T will be formulated. The manufactured solution for the velocities was formulated as

$$\begin{aligned}u_1 &= 2 \cos(x_1^2 + x_2^2 + x_3^2) x_2 + 2 \sin(x_1^2 + x_2^2 + x_3^2) x_3 \\u_2 &= 2 \cos(x_1^2 + x_2^2 + x_3^2) x_3 - 2 \cos(x_1^2 + x_2^2 + x_3^2) x_1 \\u_3 &= -2 \sin(x_1^2 + x_2^2 + x_3^2) x_1 - 2 \cos(x_1^2 + x_2^2 + x_3^2) x_2.\end{aligned}$$

The manufactured solution for the pressure p was created as

$$p = \sin(x_1^2 + x_2^2 + x_3^2) \cos(x_1^2 + x_2^2 + x_3^2).$$

Since the velocity field was created as a solenoidal field, no pressure source had to be calculated. The manufactured solution for the temperature T reads

$$T = \sin(x_1^2) \cos(x_2^2) \sin(x_3^2).$$

The resulting source terms for the momentum balances read

$$\begin{aligned}b_1 &= 8 \cos(x_1^2 + x_2^2 + x_3^2) x_1^2 x_2 + 8 \cos(x_1^2 + x_2^2 + x_3^2) x_2^3 + 8 \cos(x_1^2 + x_2^2 + x_3^2) x_3^2 x_2 \\&\quad + 8 \sin(x_1^2 + x_2^2 + x_3^2) x_1^2 x_3 + 8 \sin(x_1^2 + x_2^2 + x_3^2) x_2^2 x_3 + 8 \sin(x_1^2 + x_2^2 + x_3^2) x_3^3 \\&\quad + 7 \sin(x_1^2) \cos(x_2^2) \sin(x_3^2) + 4 (\cos(x_1^2 + x_2^2 + x_3^2))^2 x_1 + 4 (\cos(x_1^2 + x_2^2 + x_3^2))^2 x_3 \\&\quad - 4 \cos(x_1^2 + x_2^2 + x_3^2) \sin(x_1^2 + x_2^2 + x_3^2) x_2 - 20 \cos(x_1^2 + x_2^2 + x_3^2) x_3 + 20 \sin(x_1^2 + x_2^2 + x_3^2) x_2 - 6 x_1 \\b_2 &= -8 \cos(x_1^2 + x_2^2 + x_3^2) x_1^3 + 8 \cos(x_1^2 + x_2^2 + x_3^2) x_1^2 x_3 - 8 \cos(x_1^2 + x_2^2 + x_3^2) x_2^2 x_1 \\&\quad - 8 \cos(x_1^2 + x_2^2 + x_3^2) x_3^2 x_1 + 8 \cos(x_1^2 + x_2^2 + x_3^2) x_2^2 x_3 + 8 \cos(x_1^2 + x_2^2 + x_3^2) x_3^3 \\&\quad + 13 \sin(x_1^2) \cos(x_2^2) \sin(x_3^2) - 4 (\cos(x_1^2 + x_2^2 + x_3^2))^2 x_2 - 4 \cos(x_1^2 + x_2^2 + x_3^2) \sin(x_1^2 + x_2^2 + x_3^2) x_1 \\&\quad - 4 \cos(x_1^2 + x_2^2 + x_3^2) \sin(x_1^2 + x_2^2 + x_3^2) x_3 - 20 \sin(x_1^2 + x_2^2 + x_3^2) x_1 + 20 \sin(x_1^2 + x_2^2 + x_3^2) x_3 - 2 x_2 \\b_3 &= -8 \cos(x_1^2 + x_2^2 + x_3^2) x_1^2 x_2 - 8 \cos(x_1^2 + x_2^2 + x_3^2) x_2^3 - 8 \cos(x_1^2 + x_2^2 + x_3^2) x_3^2 x_2 \\&\quad - 8 \sin(x_1^2 + x_2^2 + x_3^2) x_1^3 - 8 \sin(x_1^2 + x_2^2 + x_3^2) x_2^2 x_1 - 8 \sin(x_1^2 + x_2^2 + x_3^2) x_3^2 x_1 \\&\quad - 19 \sin(x_1^2) \cos(x_2^2) \sin(x_3^2) + 4 (\cos(x_1^2 + x_2^2 + x_3^2))^2 x_1 + 4 (\cos(x_1^2 + x_2^2 + x_3^2))^2 x_3 \\&\quad - 4 \cos(x_1^2 + x_2^2 + x_3^2) \sin(x_1^2 + x_2^2 + x_3^2) x_2 + 20 \cos(x_1^2 + x_2^2 + x_3^2) x_1 - 20 \sin(x_1^2 + x_2^2 + x_3^2) x_2 - 6 x_3.\end{aligned}$$

Finally the source term for the temperature equation was calculated as

$$\begin{aligned}b_T &= -4 \sin(x_1^2) \cos(x_2^2) \cos(x_3^2) \cos(x_1^2 + x_2^2 + x_3^2) x_2 x_3 - 4 \sin(x_1^2) \cos(x_2^2) \cos(x_3^2) \sin(x_1^2 + x_2^2 + x_3^2) x_1 x_3 \\&\quad + 4 \sin(x_1^2) \sin(x_3^2) \sin(x_2^2) \cos(x_1^2 + x_2^2 + x_3^2) x_1 x_2 - 4 \sin(x_1^2) \sin(x_3^2) \sin(x_2^2) \cos(x_1^2 + x_2^2 + x_3^2) x_2 x_3 \\&\quad + 4 \cos(x_2^2) \sin(x_3^2) \cos(x_1^2) \cos(x_1^2 + x_2^2 + x_3^2) x_1 x_2 + 4 \cos(x_2^2) \sin(x_3^2) \cos(x_1^2) \sin(x_1^2 + x_2^2 + x_3^2) x_1 x_3 \\&\quad + 4 x_1^2 \sin(x_1^2) \cos(x_2^2) \sin(x_3^2) + 4 \sin(x_1^2) \cos(x_2^2) x_2^2 \sin(x_3^2) \\&\quad + 4 x_3^2 \sin(x_1^2) \cos(x_2^2) \sin(x_3^2) - 2 \sin(x_1^2) \cos(x_2^2) \cos(x_3^2) \\&\quad + 2 \sin(x_1^2) \sin(x_3^2) \sin(x_2^2) - 2 \cos(x_2^2) \sin(x_3^2) \cos(x_1^2).\end{aligned}$$

7.3 Measurement of Error and Calculation of Order

In this subsection the result of the verification process via the formerly in subsection 7.2 stated manufactured solution will be presented. To verify the solver framework using the Order-of-Accuracy test criterion, an error measure has to be

chosen on which the evaluation is based upon. According to [54] the error measure for each variable ϕ will be calculated at the end of each computation as the normalized L_2 -norm of the deviation of the exact solution $\tilde{\phi}$, sometimes referred to as the normalized global error or the RMS-error

$$\text{err}(\phi) = \sqrt{\frac{\iiint_V (\phi - \tilde{\phi})^2 dV}{\iiint_V dV}} \approx \sqrt{\frac{\sum_{n=1}^N (\phi_n - \tilde{\phi}_n)^2 V_n}{\sum_{n=1}^N V_n}}.$$

Herein N denotes the number of control volumes of the discretized solution domain V . Each control volume has the volume V_n . The numerical and the exact solution are evaluated at the center of each control volume for the error calculations. In order to calculate the normalized global error of the pressure p , the same modification as in subsection 4.8 has to be performed on the analytic solution \tilde{p} , which in the case of the present work means to subtract the weighted mean value

$$\text{mean}(\tilde{p}) = \frac{\iiint_V \tilde{p} dV}{\iiint_V dV} \approx \frac{\sum_{n=1}^N \tilde{p}_n V_n}{\sum_{n=1}^N V_n}$$

and to calculate the error as

$$\text{err}(p, N) \approx \sqrt{\frac{\sum_{n=1}^N (p_n - \tilde{p}_n - \text{mean}(\tilde{p}))^2 V_n}{\sum_{n=1}^N V_n}}.$$

The following tables will list the calculated error terms for the velocities, the pressure and the temperature for different grid resolutions and both solver algorithms on a problem domain in the form of a cube with side length $2m$ whose center is located at $x_C = (0, 0, 0)$. All involved material properties are chosen to 1, the gravitational acceleration is chosen to $\mathbf{g} = (7, 13, -19)$. As a reference temperature $T_0 = 0$ is used. The grid resolutions ($n \times n \times n$) are chosen such that n , the number of control volumes in each coordinate direction, will be a power of 2. The Errors are evaluated after the maximal relative residual of all linear systems in one outer iteration fell below the threshold 1×10^{-14} . The concrete implementation of the tolerance criterion to accept a converged solution is found in section 6.3.2. It should be noted that no results were calculated with the coupled solver for the case of $256 \times 256 \times 256$ unknowns due to convergence problems of the linear equation solver.

A comparison of the results leads to the first conclusion, that the results from the segregated solver algorithm and the coupled solver algorithm coincide, which implies that the same discretization is used in both cases. According to [54] the next step comprises the calculation of the formal order of accuracy \hat{p}_ϕ for each variable that the program solves for. This can be done by evaluating the quotient of two different resolutions n_1 and n_2 as

$$\hat{p}_\phi = \frac{\log\left(\frac{\text{err}(\phi, n_1)}{\text{err}(\phi, n_2)}\right)}{\log\left(\frac{n_1}{n_2}\right)}.$$

The discretization techniques used in the present thesis yield an asymptotic discretization error of 2 [55]. Comparing with the results in the tables 1, 2 and 3 shows the asymptotic behavior of the calculated order which leads to the conclusion that the discretization has been implemented with success.

7.4 Influence of the Under-Relaxation Factor for the Velocities

In section 4.2, the main purpose of introducing the pressure weighted interpolation method, was to assure comparability of the generated results from the segregated solver algorithm and the coupled solver algorithm. Section 7.3 showed good agreement of the error calculated with the segregated solver, which uses the pressure weighted interpolation method, and the coupled solver. This section presents the results that are attained, if the standard Rhie-Chow interpolation technique with different under-relaxation factors α_u is used to interpolate the velocities to the boundary faces instead of the pressure weighted interpolation technique. Figure 12 shows the normalized L_2 norm of the error of the results for the velocity u_1 for different under-relaxation factors. The L_2 norms have been normalized by the L_2 norm of the respective calculation that uses the pressure weighted interpolation method.

It is evident that, the higher the under-relaxation factor, the smaller the deviation of the results obtained with the standard Rhie-Chow interpolation technique from the results obtained with the pressure weighted interpolation method. This can be justified by the fact that, for an under-relaxation factor $\alpha_u = 1$, the pressure weighted interpolation and the

Table 1: Comparison of the errors of the velocity, calculated by the segregated and the coupled solver for different grid resolutions and the resulting order of accuracy

Resolution	Error of \mathbf{u} from Segregated Solver	Error of \mathbf{u} from Coupled Solver	Observed Order \hat{p}
8x8x8	3.170450230115786E-002	3.1704502290288525E-002	
	3.037011030746674E-002	3.0370110306195783E-002	
	3.163609515028456E-002	3.1636095141306442E-002	
16x16x16	7.914356673006179E-003	7.9143566730065350E-003	2.0021
	7.534314612440339E-003	7.5343146124409056E-003	2.0111
	7.811316741535871E-003	7.8113167415361021E-003	2.0179
32x32x32	1.933302886747634E-003	1.9333250814097063E-003	2.0334
	1.867599197138540E-003	1.8676512523175990E-003	2.0123
	1.894813910088982E-003	1.8948223205995562E-003	2.0435
64x64x64	4.746894199493598E-004	4.7468942001527124E-004	2.0260
	4.629259933395594E-004	4.6292599348776502E-004	2.0123
	4.630008203550293E-004	4.6300082038027822E-004	2.0330
128x128x128	1.172266972290949E-004	1.172266974040543E-004	2.0177
	1.150024026498247E-004	1.150024030604766E-004	2.0091
	1.140471182846354E-004	1.140471183418716E-004	2.0214
256x256x256	2.876932445431025E-005		2.0267
	2.771472789328757E-005	d.n.c.	2.0529
	2.826736500332136E-005		2.0124

Table 2: Comparison of the errors of the pressure, calculated by the segregated and the coupled solver for different grid resolutions and the resulting order of accuracy

Resolution	Error of p from Segregated Solver	Error of p from Coupled Solver	Observed Order \hat{p}
8x8x8	0.183035893121686	0.183035894664680	
16x16x16	8.967727683983406E-002	8.967727683982576E-002	1.0293
32x32x32	3.796485688816108E-002	3.796481581400934E-002	1.2401
64x64x64	1.438780100622541E-002	1.438780100615654E-002	1.3998
128x128x128	5.160705648852031E-003	5.160705649392788E-003	1.4792
256x256x256	1.809811107444374E-003	d.n.c.	1.5117

Table 3: Comparison of the errors of the temperature, calculated by the segregated and the coupled solver for different grid resolutions and the resulting order of accuracy

Resolution	Error of T from Segregated Solver	Error of T from Coupled Solver	Observed Order \hat{p}
8x8x8	4.067532377541200E-003	4.067532377537829E-003	
16x16x16	1.132477318394773E-003	1.132477318394782E-003	1.8447
32x32x32	2.931009801582052E-004	2.931009624340290E-004	1.9500
64x64x64	7.410021974469997E-005	7.410021974334421E-005	1.9838
128x128x128	1.858846421984987E-005	1.858846421622182E-005	1.9951
256x256x256	4.651911366447223E-006	d.n.c.	1.9985

Rhie-Chow interpolation would yield the same result according to equation (16). Furthermore, it can be observed that the deviation decreases with increasing mesh resolution which is the expected behaviour [19].

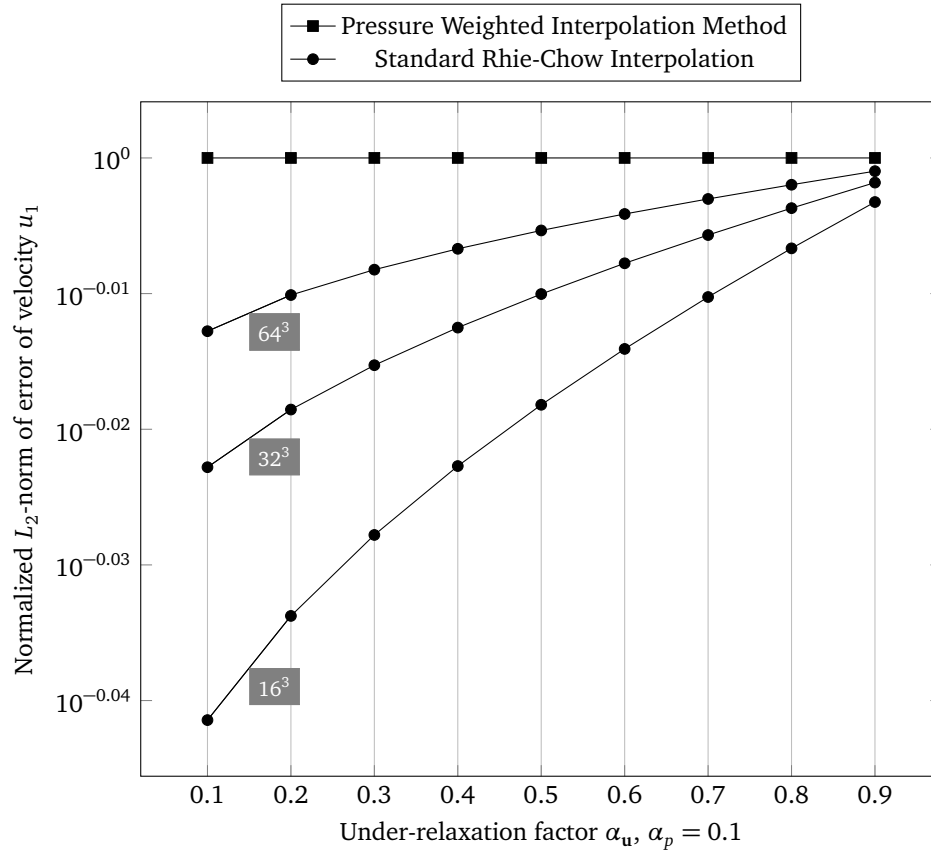


Figure 12: Comparison of calculated error for different under-relaxation factors α_u on grids with $16 \times 16 \times 16$, $32 \times 32 \times 32$ and $64 \times 64 \times 64$ cells

8 Comparison of Solver Concepts

This section conducts the performance analysis of both developed solvers. In the first part of this section, the used hardware and software is described. Furthermore, important performance measures are introduced. After performing a preliminary benchmark to measure the sustainable bandwidth, the Speed-Up is measured for both solver programs. The second part of this section tests the performance of the algorithm on a test case with complex geometry and on a cavity benchmark for flows with natural convection. In the end the results of a study on the effect of non-orthogonal corrections on the convergence behaviour of the coupled solution algorithm is presented.

8.1 Parallel Performance

In many cases, scientific code is used to solve complex problems regarding memory requirements or to make calculation results available within short time. In both scenarios, code that is able to run in parallel can address the mentioned challenges. Code that runs in parallel can allocate more memory resources which makes the calculation of complex problems feasible. If the code is scalable, the program execution can be shortened by involving more processes to solve a problem of constant size on the compatible hardware.

The solver framework that has been developed in the course of the present thesis, has been parallelized using the PETSc library. After introducing the used hardware and software, the central measures of parallel performance are presented. Then preliminary test results using low-level benchmarks are performed, which establish upper performance bounds on the parallel efficiency and the scalability of the developed solver framework. The results of the efficiency evaluation of the solver framework are presented in the last subsections.

8.1.1 Employed Hardware and Software – The Lichtenberg-High Performance Computer

All performance analyses that are presented in this thesis were conducted on the Lichtenberg-High Performance Computer, also known as *HHLR* (*Hessischer Hochleistungsrechner*). The cluster consist of different sections according to the used hardware. Throughout the thesis, tests were performed using the first and the second MPI section of the cluster. The first section consists of 705 nodes of which each runs two Intel®Xeon®E5-2670 processors and offers 32GB of memory. The second section consists of 356 nodes of which each runs two Intel Xeon E5-2680 v3 processors and offers 64GB of memory. As interconnect for both sections, FDR-14 InfiniBand is used.

All tests programs were compiled using the Intel compiler suite version 15.0.0 and the compiler options

—O3 —xHost.

As MPI implementation, Open MPI version 1.8.2 was chosen. Furthermore the PETSc version 3.5.3 was configured using the options

```
---with-blas-lapack-dir=/shared/apps/intel/2015/composer_xe_2015/mkl/lib/intel64/ \
---with-mpi-dir=/shared/apps/openmpi/1.8.2_intel \
COPTFLAGS="--O3 -xHost" \
FOPTFLAGS="--O3 -xHost" \
CXXOPTFLAGS="--O3 -xHost" \
---with-debugging=0 \
---download-hypre \
---download-ml
```

It should be noted, that, as the configurations options show, to maximize the efficiency of PETSc, a math kernel library should be used, that has been optimized for the underlying hardware architecture. Is in the case of the present thesis the Intel *MKL* (*Math Kernel Library*) has been used. The Open MPI library was compiled using Intel compilers as well.

8.1.2 Measures of Performance

This section establishes the needed set of measures to evaluate the performance of a solver program, which will be used in the following sections. The first measure is the plain measure of runtime T_p taken by a computer to solve a given problem, where $P \in \mathbb{N}$ denotes the number of involved processes. This so called *wall-clock* time can be measured directly, by calling subroutines of the underlying operating system and corresponds to the human perception of the time, that has passed. It must be noted, that this time does not correspond to the often mentioned *CPU* time. In fact, CPU time is only one contributor to wall-clock time. In one side, wall-clock time contains the time needed for communication and I/O and hence considers idle states of the processor. On the other side, CPU time only considers the time in which the processor is actively working. This makes wall-clock time not only a more complete time measure, but also more accurate time measure, when dealing with parallel processors, since processor idle times due to communication are actively considered, while neglected in CPU time.

While wall-clock time is an absolute measure that can be used to compare different solver programs, further relative measures are needed to evaluate the efficiency of one program regarding the parallelisation implementation. The main purpose of these measures is to assess the different causes of degrading efficiency due to parallelisation of the algorithm or the parallel execution of the algorithm. A simple model [19, 55] considers three contributions, that form the total efficiency

$$E_p^{tot} = E_p^{num} \cdot E_p^{par} \cdot E_p^{load}.$$

The *numerical efficiency*

$$E_p^{num} := \frac{\text{FLOPS}(1)}{P \cdot \text{FLOPS}(P)}$$

considers the degradation of the efficiency of the underlying algorithm due to the parallelisation. Many efficient algorithms owe their efficiency to recursions inside the algorithm. In the process of resolving this recursions, the efficiency of the algorithm degrades. It follows that this efficiency is completely independent of the underlying hardware.

The *parallel efficiency*

$$E_p^{par} := \frac{\text{TIME}(\text{parallel algorithm on one processor})}{P \cdot \text{TIME}(\text{parallel algorithm on } P \text{ processors})}$$

describes the effect of the need for inter process communication, if more than one processor is involved in the solution process. It should be noted, that this form of efficiency does explicitly exclude any algorithm related degradation, since the time measured corresponds to the exact same algorithms. It follows that the parallel efficiency only depends on the implementation of the communication and the hardware related latencies.

The *load balancing efficiency*

$$E_p^{load} := \frac{\text{TIME}(\text{calculation on complete domain})}{P \cdot \text{TIME}(\text{calculation on biggest subdomain})}$$

is formed by the quotient of the wall-clock times needed for the complete problem domain and partial solves on subdomains. This measure does neither depend on hardware nor on the used implementation. Instead it directly relates to the size and partition of the grid.

It is not possible to calculate all three efficiencies at the same time using only plain wall-clock time measurements of a given application. Different solver configurations have to be used to calculate them separately. Since the focus of investigation of the present thesis does not lie on load balancing, for the remainder of the thesis $E_p^{load} = 100\%$ is assumed. This does not present a considerable drawback, since an ideal load balancing is easily obtainable nowadays by the use of grid partitioning algorithms [19]. Using identical algorithms for different numbers of involved processes, implicitly achieves $E_p^{num} = 100\%$. In this case the parallel efficiency of an application can be measured through the quotient of the needed wall-clock time. To measure the numerical efficiency of an algorithm the respective hardware counters have to be evaluated. This can be done using the built-in log file functionality of PETSc.

Another common performance measure is the *Speed-Up*

$$S_p = \frac{T_1}{T_p} = P \cdot E_p^{tot}.$$

Speed-Up and parallel efficiency characterize the parallel scalability of an application and determine the regimes of efficient use of hardware resources.

8.1.3 Preliminary Upper Bounds on Performance – The STREAM Benchmark

Scientific applications, which solve partial differential equations, rely on sparse matrix computations, which usually exhibit the sustainable memory bandwidth as bottleneck with respect to the runtime performance of the program [26]. The purpose of this section is to establish a frame in terms of an upper bound on performance in which the efficiency of the developed solver framework can be evaluated critically. As common measure for the maximum sustainable bandwidth, low-level benchmarks can be used, which focus on evaluating specific properties of the deployed hardware architecture. In this case, the STREAM benchmark suite [39, 40] provides apt tests, which are designed to work with data sets that exceed the cache size of the involved processor architecture. This forces the processors to stream the needed data directly from memory instead of reusing the data residing in their caches. These types of tests can be used to calculate an upper bound on the memory bandwidth for the CAFFA framework.




Figure 13: Sustainable memory bandwidth as determined by the STREAM benchmark (Triad) for different process binding options on one node of the MPI1 section

In terms of parallel scalability, the STREAM benchmark can also be used as an upper performance bound. According to [5] the parallel performance of memory bandwidth limited codes correlates with the parallel performance of the STREAM benchmark, i.e. a scalable increase in memory bandwidth is necessary for scalable application performance. The intermediate results of the benchmark can then be used to test different configurations that bind hardware resources to the involved processes. Figures 13 and 14 show the test results of the STREAM benchmark on the two different MPI sections of the HHRL. It should be noted that it is central for the results to be reproducible, to always run the tests exclusively on each node, i.e. that no other process is running on the same node.

As can be seen from figure 13 and 14, the scaling of the sustainable bandwidth behaves rather erratic, such that for process counts up to eight for the MPI1 section and process counts up to six for the MPI2 section, no reliable results can be obtained from the STREAM benchmark. It is assumed that this kind of behaviour is due to the automatic turbo-boost the deployed processors apply, which cannot be controlled other than by turning it off on the nodes directly from the BIOS. For the following performance analyses with respect to parallel performance, it is desirable to minimize this kind of unpredictable side effects. Since all performance measures are relative but the plain measurement of wall-clock time, the reference value for the following relative performance measurements will be taken from the program execution for the maximum number of processes that can be bound without overlap to one deployed socket.

It should be taken into account that the STREAM benchmark used to test the present architecture was provided by the PETSc framework. Special effort has been made to adapt this benchmark to handle the described problem, by always using all available cores per socket, but only running the benchmark on n processes while the other $24 - n$ or $16 - n$ processes run calculations which only use data from the L1-cache of their corresponding core. The intention of this approach was to regulate the turbo-boost effects by keeping a constant load on the processors for different numbers of STREAM benchmark processes. However this approach didn't generate results different from figures 13 and 14.




Figure 14: Sustainable memory bandwidth as determined by the STREAM benchmark (Triad) for different process binding options on one node of the MPI2 section

8.1.4 Speed-Up Measurement and Impact of Coupling Algorithm for Analytic Test Case

This section presents the results from the Speed-Up measurements conducted on the supercomputer HHLR, which has been presented in section 8.1.1. Furthermore, the impact of the scaling algorithm on the running time of the corresponding solver program will be shown. This property of applications is often termed *strong* scalability [26]. All Speed-Up tests are conducted on the MPI1 section. The performance analyses will show the absolute time needed to solve for the flow of the analytic solution presented in section 7.2. In order to evaluate the parallel performance the Speed-Up will be calculated based on the presented results. Since the performance of segregated algorithms is linked to the amount of under-relaxation applied to the variables in each outer iteration, a preliminary test is conducted to support choosing the optimal amount of under-relaxation. As pointed out in [19, 55], it is recommended to choose the amount of under-relaxation for the pressure depending on the amount of under-relaxation for the velocity as

$$\alpha_p = 1 - \alpha_u.$$

Figure 15 shows the results of the preliminary runtime measurements for different choices of the under-relaxation factor α_u . The problem that was solved during this test, corresponds to the manufactured solution for the velocities and the pressure presented in section 7.2. It can be observed, that the optimal under-relaxation factor can be found in a neighborhood of 0.95. Based on the results presented in figure 15 the under-relaxation was chosen to $\alpha_u = 0.9$.




Figure 15: Wall-clock time comparison for the SIMPLE-algorithm, solving for the analytic solution presented in section 7.2 on a structured grid with 64^3 cells using different under-relaxation factors up to $\alpha_u = 0.999$

The performance of all solvers significantly depends on the performance of the linear solvers and the corresponding preconditioners used in the solution process. For all conducted tests, the linear system for the discretized momentum balances will be solved with a GMRES (*Generalized Minimal Residual*) [53] Krylov subspace solver and an incomplete LU-factorization as preconditioner. The pressure correction equation is solved by a CG (*Conjugated Gradient*) solver [50] with the GAMG (*Geometric Algebraic Multi Grid*) preconditioner from PETSc. The coupled systems are solved by a GMRES Krylov subspace solver with GAMG preconditioning. The GAMG preconditioner was chosen so that a high numerical efficiency is obtained, constituting a basis for further parallel performance measurements based on plain wall-clock time measurements. All further solver parameters are listed in table 4.

From the comparison in figure 16 it is evident, that both algorithms are scalable and the coupled solution algorithm accelerates the solution process almost up to one order of magnitude. For higher amounts of under-relaxation for the velocities the difference in performance is expected to increase as figure 15 suggests. It has to be noted, that the presented results for the coupled solution algorithm were obtained using black-box solvers from the PETSc framework. It is assumed that the performance of the coupled solution algorithm with respect to the needed wall-clock time would benefit from multigrid solvers especially designed for the coupled system of linear equations. Examples for such special solvers can be found in [12, 30, 36]. Another approach is represented by *physics-based* solvers which take advantage of the matrix structure as shown in figure 10. References [8, 41] describe how PETSc can be used to construct physics based preconditioners in the context of multiphysics simulations.

Table 4: Characteristic problem properties used in the performance measurements, solving for the analytic solution presented in section 7.2

Property	Value	Unit
Density	1E-0	kg/m^3
Viscosity	1E-0	Ns/m^2
Under-relaxation u	0.9	
Under-relaxation p	0.1	
Relative tolerance	1E-8	




Figure 16: Wall-clock time and Speed-Up comparison for segregated and fully coupled solution algorithm, solving for the analytic solution presented in section 7.2 on a structured grid with 128^3 cells

8.2 Weak Scaling Comparison of Coupled and Segregated Solution Algorithm

It has been previously reported in [12, 61], that the coupled solution algorithm also scales well with the size of the problem domain with respect to the involved number of unknowns. It is known that this is not true for the SIMPLE-algorithm as representative of the set of segregated solution algorithms. An algorithm that scales well with the problem size is apt to be tested on a benchmark that analyzes the so called *weak* scaling of the implementation [26]. Starting from a moderate number of unknowns for a tree-dimensional test case solving for the analytic solution presented in section 7.2, this benchmark will investigate how the running-time and the number of needed inner and outer iterations behaves when the computational load per process is held constant, while the number of involved processes is continuously increased. *Weak* scaling benchmarks show, if it is possible to solve a more complex problem in the same time by involving more processes running the solver application. Each process will get assigned a block with $32 \times 32 \times 32$ unknowns. Thus, the number of involved unknowns in each test run can be calculated as $N = 32 \times 32 \times 32 \times 32$. Figure 17 compares the needed number of outer iterations and the corresponding wall-clock time for an implementation of the SIMPLE-algorithm and the fully coupled solution algorithm.

The comparison in figure 17 shows, that the fully coupled algorithm scales ideally with respect to the number of outer iterations, which remains almost constant throughout the conducted study. The number of iterations at most differed by 2 outer iterations. These differences are suspected to be caused by the loose tolerance criterion of residual reduction in each outer iteration. As an effect of accumulating rounding errors, some test cases did not accomplish the total tolerance criterion and thus needed to perform another additional outer iteration. From the wall-clock time measurements presented in figure 17 it is evident that, despite the approximately constant number of outer iterations, the fully coupled solution algorithm did not scale ideally. Figure 18 shows that weak scaling is prevented because the linear solver does not scale ideally. This shows that weak scaling cannot be obtained using black-box linear solvers and




Figure 17: Comparison of the number of outer iterations and the corresponding wall-clock time, needed to achieve a reduction of $1E-8$ of the initial residual, for different methods to resolve pressure-velocity coupling

that further research regarding the development of special linear solvers for the systems resulting of a fully coupled discretization of the Navier-Stokes equations is needed.




Figure 18: Average number of inner iterations for different process counts solving for the analytic solution presented in section 7.2 with the fully coupled solution algorithm

8.3 Realistic Testing Scenario – Complex Geometry

Fluid flow inside closed applications is a common situation in mechanical engineering. The flow through a channel with rectangular cross section can be seen as a simple test case, that is a part of complex applications. This section compares the single process performance of the segregated and the fully coupled solution algorithm for a flow problem within a complex geometry. The geometry of the domain is based on a channel flow problem with square cross section, with the special property, that inside the channel reside two obstacles with a square cross section of which one has been twisted against the other. Figure 19 shows a sketch of the problem domain.




Figure 19: Sketch of the channel flow problem domain

This case exercises all of the previously introduced boundary conditions for flow problems, including the treatment of non-matching block boundaries. For the velocities at the inflow boundary, the parabolic distribution

$$\mathbf{u}(x_1, x_2, x_3) = \begin{bmatrix} u_1 \\ u_2 \\ u_3 \end{bmatrix} = \begin{bmatrix} \frac{16 * 0.45 * x_2 * x_3 * (0.41 - x_2) * (0.41 - x_3)}{0.41^4} \\ 0 \\ 0 \end{bmatrix}$$

was chosen. At the outlet a Dirichlet pressure boundary condition was used. All other boundaries used a no-slip solid non-moving wall boundary condition. All problem parameters were chosen to assure the flow problem resides in the regime of a non-turbulent stationary flow, for which the presented solver framework has been developed. Table 5 lists the remaining material and geometrical characteristics of the test case. A study similar to section 8.1.4 was conducted to optimize the used amount of under-relaxation.

Table 5: Characteristic problem properties used in the channel flow test case

Property	Value	Unit
Density	1E-0	kg/m^3
Viscosity	1E-3	Ns/m^2
Height	0.41	m
Length	2.5	m
Side length cube	0.1	m
Under-relaxation u	0.9	
Under-relaxation p	0.1	
Relative tolerance	1E-8	

Using the data from table 5 and the maximal inflow velocity of $0.45m/s$ the Reynolds number Re can be calculated as

$$Re = \frac{\rho \text{mean}(u_1)l}{\mu} = 20,$$

where $\text{mean}(u_1) = 0.45 * \frac{4}{9} m/s = 0.2m/s$. This shows, that the flow resides in a laminar regime.

The present test case shows one advantage of the treatment of block boundaries, which has been introduced in section 4.7.3. Since no assumptions on the geometry of a neighboring block are necessary, the mesh for each block can be constructed independently which increases the flexibility of the meshing of geometries. Furthermore, because of the fully implicit handling of block boundaries, the number of used blocks does not negatively affect the convergence of the deployed linear solvers. Figure 20 shows the mesh at the left and right bounding walls. It is evident that this mesh leads to non-trivial transitions between the blocks. Figure 21 shows the domain decomposition into structured grid blocks

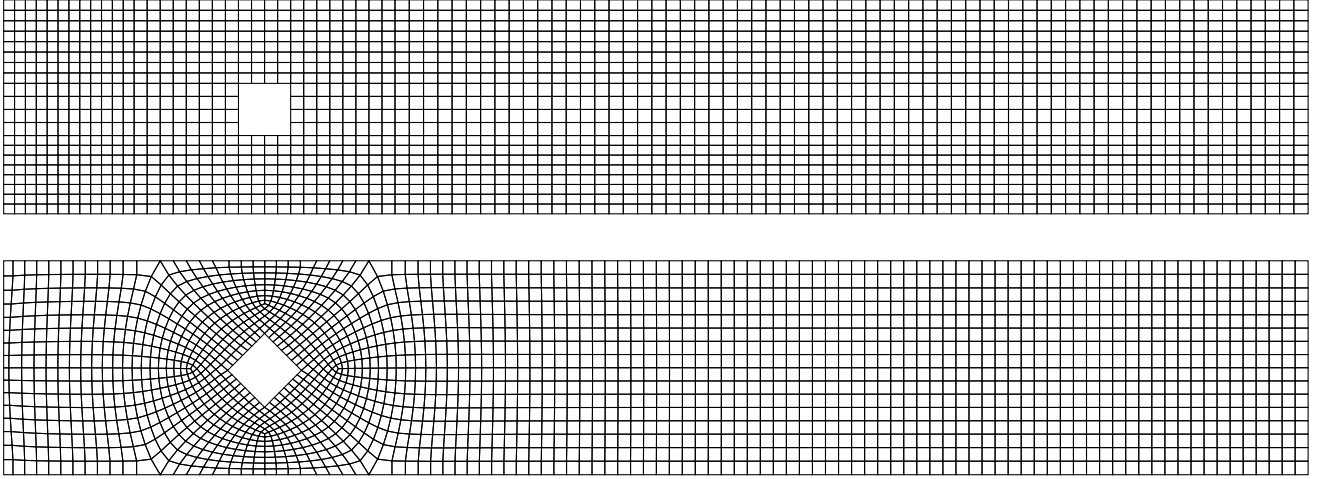


Figure 20: West and east boundary of the numerical grid for the channel flow problem




Figure 21: Blocking for the two different obstacles within the problem domain of the channel flow

around the two obstacles within the problem domain and emphasizes the need for accurate handling of non-matching block boundaries.

The solution of the linear systems resulting from the discretization of the problem, takes up more time for the coupled solution algorithm than in the segregated coupled algorithm. For small problem sizes the additional overhead for the solution methods for linear systems and the fact that the segregated solution algorithm for small numbers of unknowns does not need many outer iterations to converge, leads to the conclusion that moderate to big problems with respect to the number of involved unknowns are necessary for the coupled solution algorithm show significant advantages in performance.

In contrast to the segregated algorithm, the fully coupled solution algorithm achieves an approximately constant amount of needed outer iterations, independent of the number of involved unknowns. The tests regarding the weak scalability of the coupled solution algorithm, presented in section 8.2 emphasize this property. Table 6 compares the measured wall-clock time for different numbers of unknowns. The mesh shown in figure 20 was generated for the first number of unknowns and successively refined to achieve higher mesh resolutions. The two other numbers of unknowns result from up to two times bisectioning the mesh in each direction, every time scaling the number of unknowns by a factor of approximately eight. The tests were conducted on the formerly presented HHLR cluster, using the MPI2 section.

The timing results show, that already after the first grid refinement step, the fully coupled solution algorithm performs better with respect to the needed wall-clock time for computation. This effect is clearly visible for higher mesh resolutions.

8.4 Classical Benchmarking Case – Temperature-Driven Cavity Flow

This section deals with the evaluation and comparison of the velocity-to-temperature coupling, through the Boussinesq approximation, and the temperature-velocity/pressure coupling, through the Newton-Raphson linearization of the convective term of the temperature equation. For this, the standard temperature-driven cavity flow [11, 14] is adapted for a three dimensional domain and the material and geometric parameters are chosen in a way that a non-turbulent station-

Table 6: Performance analysis results of the channel flow problem for different numbers of unknowns comparing the segregated (SEG) to the fully coupled (CPLD) solution algorithm using one process on the MPI2 section of the HHLR supercomputer.

No. of Unknowns	SEG - time s	CPLD - time s	SEG - its	CPLD - its
75768	0.2226E+02	0.2674E+02	151	67
408040	0.4053E+03	0.1499E+03	355	42
2611080	1.1352E+05	0.3105E+04	1592	39

ary flow exists, which means that the solution lies within the regime of the approximations made by the solver. Figure 22 shows an example for a temperature and velocity field obtained with the developed solver framework, solving for the temperature-driven cavity flow problem.




Figure 22: Temperature field and velocity field in the coordinate direction of the gravitational force of the temperature-driven cavity flow problem for a cross section in the middle of the problem domain

Essential for this benchmarking case is the nature of the flow. The fluid motion is a consequence of the effect of volume forces caused by temperature differences in the solution domain, hence the mathematical problem exhibits a strong coupling between the involved variables velocity, pressure and temperature. This relation is represented by the Rayleigh and Prandtl number. It is assumed that for this kind of flow the fully implicit treatment of the temperature coupling will yield further benefits with respect to wall-clock time, compared with solution approaches that solve for the temperature separately. Table 7 lists the geometrical and solver parameters used for the performance analysis of this section.

Using the data from table 7 and given a Prandtl number for the flow problem $Pr = 0.71$, the Rayleigh number Ra , a dimensionless number used to characterize buoyancy driven flows, can be calculated. From the definition of the Prandtl number, the thermal diffusivity α can be calculated as

$$\alpha = \frac{\mu}{\rho Pr} = 2.1304 * 10^{-5}.$$

It follows, that the Rayleigh number can be determined to

$$Ra = \frac{\rho g \beta \Delta T h^3}{\mu \alpha} \approx 10^4,$$

which implies, according to [11], that the flow is still stationary and non-turbulent.

The tests were conducted on the formerly presented HHLR cluster, using the MPI2 section. Table 8 compares the wall-clock times and numbers of needed outer iterations for different solver and coupling approaches. The presented results are in good agreement with [61], which show monotonic decrease of the number of iterations with increased

Table 7: Characteristic problem properties used in the heated cavity flow test case

Property	Value	Unit
Density	1.19	kg/m^3
Viscosity	1.8E-5	Ns/m^2
Height	0.021277	m
Temperature difference	10	K
Coefficient of thermal expansion	0.00341	1/K
Under-relaxation u	0.8	
Under-relaxation p	0.2	
Under-relaxation T	1.0	
Relative tolerance	1E-8	

implicit coupling for a different test case. It is notable, that the pressure velocity coupling is responsible for the biggest decrease in the number of non-linear iterations. Different to the results presented in section 8.3, this does not yield significant performance benefits with respect to wall-clock time. In order to achieve the benefits of a fully coupled solution algorithm, the coupling has to be increased to also involve semi-implicit temperature-to-velocity/pressure and implicit velocity-to-temperature coupling. In this context the abbreviations *SEG* for the segregated solution algorithm and *CPLD* for the fully coupled solution algorithm without implicit temperature coupling are used. For the fully coupled solver configurations, *TCPLD* includes only implicit velocity-to-temperature coupling, whereas *NRCPLD* additionally includes implicit temperature-to-velocity/pressure coupling via the Newton-Raphson linearization.

Table 8: Performance analysis results of the temperature-driven cavity flow problem comparing the SIMPLE-algorithm with segregated temperature solve (SEG), the fully coupled solution algorithm with segregated temperature solve (CPLD), the fully coupled solution algorithm with an implicit Boussinesq approximation (TCPLD) and the fully coupled solution algorithm using an implicit Boussinesq approximation and a semi-implicit Newton-Raphson linearization of the convective part of the temperature equation (NRCPLD).

Resolution	Solver configuration	Time s	No. Non-linear its.
32x32x32	SEG	0.3719E+02	203
	CPLD	0.6861E+02	62
	TCPLD	0.1012E+03	31
	NRCPLD	0.2153E+02	22
64x64x64	SEG	0.1997E+04	804
	CPLD	0.7687E+03	63
	TCPLD	0.1278E+04	59
	NRCPLD	0.4240E+03	17
128x128x128	SEG	0.5197E+05	3060
	CPLD	0.1860E+05	74
	TCPLD	0.1950E+05	50
	NRCPLD	0.6155E+04	18

It can be noticed from table 8, that the solely use of velocity-to-temperature coupling does not result in benefits compared to the coupled solution for velocities and pressure combined with a decoupled solve for the temperature equation. This is accredited to the treatment of the non-linearity of the temperature equation in the TCPLD solver configuration. Even though the momentum balances implicitly use the temperature from the next iteration, the temperature equation does not use the velocities of the next iteration. Instead the convective fluxes are calculated with the velocities from the previous iteration. Even though the implicit velocity-to-temperature coupling reduces the number of needed non-linear iterations an increase in the needed wall-clock time for computation can be noticed. This is attributed to the augmented costs during the application of the linear solver algorithm as a result of the additional degree of freedom that the linear

system embraces. The necessary number of outer iterations is higher than in the NRCPLD configuration, since the convective fluxes in the temperature equation are not coupled to the momentum balances, which degrades the convergence of the non-linear iteration process.

It can be concluded, that in order to profit from the benefits of implicit coupling, the key lies in the semi-implicit temperature-to-velocity/pressure coupling. However, it should be noted, that the implicit consideration of the corresponding terms furthermore increases the amount of memory needed for computation and hence does not come without deficiencies.

8.5 Effect of Different Non-Orthogonal Correctors on Solver Convergence

Section 3.3 introduced different ways to address the degradation of the grid quality due to non-orthogonality of the grid. This section compares the three different correctors, the orthogonal correction, the minimum correction and the over-relaxed approach, which are used in the discretization process of the gradients at cell boundary faces. For this purpose the grid generator program of section 6.2 was extended to generate skewed grids. A random number generator was used to move each inner grid point within a specified neighborhood of the original location and to skew the grid. Figure 23 illustrates the concept for a two-dimensional 8×8 grid. It is evident that the previous orthogonality of the grid is lost and the need for non-orthogonal correctors becomes visible. It should be noted that even in the case of orthogonal grids, non-orthogonal correction may become necessary to treat fluxes across block boundaries if neighboring blocks have been locally refined. To maintain the grid's integrity after the movement of the grid points the neighborhoods were limited by half the distance Δx of two neighboring grid points, as indicated for one point in figure 23.




Figure 23: Skewing of an initially equidistant, structured and orthogonal grid via random movement of inner grid points within a small neighborhood of their original location. The neighborhood with maximal diameter is indicated by a circle around one inner grid point

To measure the effect of different non-orthogonal corrections on the solution process, tests for different skewed grids were performed, measuring the number of needed outer iterations to solve for the analytic solution presented in section 7.2. For this, all correctors addressing grid non-orthogonality from section 3.3 have been implemented in the coupled solver program. The grid skewness was parametrized with the relative diameter α of the maximal neighborhood, which constrains the movement of grid points. $\alpha \in [0, 1)$, where $\alpha = 0$ corresponds to no movement of grid points at all and $\alpha = 1$ would move grid points up to $\frac{\Delta x}{2}$ away from their original location. The choice of $\alpha = 1$ is not permitted to maintain the grid's integrity. Such a choice would permit the grid generator to place two grid points at the same location. The calculations were terminated after a reduction of 10^{-14} of the relative initial residual was obtained.

The results presented in figure 24 affirm the results presented in [28]. It can be observed, that for small movements of grid points the choice of corrector does not affect convergence. However as movements of grid points increase, and thus higher non-orthogonality increases, the over-relaxed approach outperforms the other correctors with respect to the number of needed outer iterations. This does directly translate into performance with respect to wall-clock time, since the computational effort of the application of each non-orthogonal corrector is approximately the same.




Figure 24: Number of needed outer iterations for different relative diameters of the neighborhood in which the movement of grid points takes place, parametrized by the orthogonal corrector

9 TESTING ONLY

9.1 Achieving Ideal Load Balancing by Automatic Matrix Partitioning

The performance analyses conducted in subsection 8.1.4 relied on perfect load balancing and thus a load balancing efficiency of 100%. However, in practice achieving ideal load balancing may be difficult, because of local refinement strategies or an unfavorable grid blocking, i.e. the partition of the solution domain into grid blocks, that lead to an imbalance of load that is distributed among the involved processes.

One possibility partially to address load imbalances is motivated by the fact, that solution algorithms which solve partial differential equations spend a significant amount of their time solving linear algebraic systems. Figures ?? and ?? show, that approximately 50% of the wall-clock time of the segregated SIMPLE-solution algorithm, and up-to 89% of wall clock time of the coupled solution algorithm are spent solving linear systems. The goal of this subsection is to present a simple method by which ideal load balancing with respect to the linear equation solver can be obtained by automatic matrix partitioning.

The formula PETSc uses to evenly distribute vector or matrix objects containing N rows across m involved processes reads for the p th process

$$n(p) = \begin{cases} \text{floor}(N/m) + 1, & \text{if } N \bmod m > p \\ \text{floor}(N/m), & \text{else} \end{cases},$$

where $p = 0, \dots, m-1$.

Unknowns A random number generator was used to create different grid resolutions for each one of the eight block of a block structured grid consisting of 826891 cells. Figure ?? shows the initial distribution of unknowns and the balanced number of unknowns with respect to the matrix.

Figure ?? now shows the rest.




Figure 25: Comparison of the number of unknowns assigned to each process and the assigned rows of the matrix using automatic repartitioning for load balancing of the 826891 unknowns.




Figure 26: Distribution of wall-clock time for one and eight processes, solving for an analytic solution, using the default data partitioning for matrix assembly and automatic partitioning to balance the load. MAIN refers to the flux calculation and matrix assembly, MOM to the solution of the linear systems for the momentum balance and PRESS to the solution of the linear systems for the pressure correction.




Figure 27: Distribution of wall-clock time spent to solve for an analytic solution with the coupled solution algorithm on one processor

10 Conclusion and Outlook

In the present thesis, a fully coupled solution algorithm for the Navier-Stokes equations with a finite-volume method was implemented on co-located, locally refined, block-structured grids. The main difference of this algorithm, compared to the commonly used segregated solution algorithms, was the implicit pressure-velocity coupling. The pressure-velocity coupling is a characterizing aspect of solution algorithms for Navier-Stokes equations with the potential to significantly reduce the time for computation. Furthermore, the solution algorithm was extended to solve for buoyancy driven flows, by using an implicit Boussinesq approximation in the momentum balances to achieve velocity-to-temperature coupling and a semi-implicit Newton-Raphson linearization to realize implicit temperature-to-velocity/pressure coupling.

The CAFFA solver framework, which was extended during this thesis was verified using a grid convergence study and a three-dimensional manufactured solution for the velocities, the pressure and the temperature. Further studies on the segregated solution algorithm showed, that a special modification of the widely used Rhie-Chow momentum interpolation scheme is necessary to calculate results that are independent of the under-relaxation factor and thus comparable with the results obtained with a fully coupled solution algorithm. The implementation of the algorithm considered the use of high performance computers by using the PETSc library for parallelisation of the involved data structures and solving the resulting linear systems with solvers and preconditioners provided by the same toolkit.

A comparison study was performed, dealing with the effect of implicit pressure-velocity coupling in the numerical solution process of the Navier-Stokes equations. Parallel performance measurements showed, that, for the implementation of the segregated and the coupled solution algorithm, scalable program behaviour is possible for high numbers of involved unknowns. Performance studies on flows through complex geometries demonstrated, that the fully coupled solution algorithm outperforms the segregated solution algorithm.

The analyses regarding the different degrees of velocity-to-temperature and temperature-to-velocity/pressure coupling showed that the implicit temperature-to-velocity/pressure coupling is the key aspect of maintaining high efficiency regarding the computation time for flow problems involving temperature transport and strong coupling between the velocities, the pressure and the temperature.

The present thesis revealed different starting points to further examine the effects of coupled solution algorithms in the context of finite-volume flow solvers. In the present thesis only stationary problems were solved. Due to the moderate Rayleigh number for the heated cavity test case a stationary solution existed. By increasing the Rayleigh number the coupling between the velocities, the pressure and the temperature gets stronger and the flow will get instationary. A further investigation could examine if coupled solution methods maintain the higher efficiency compared to segregated solution algorithms when applied to instationary flows.

In this thesis a temperature equation was successfully coupled to the modeling equations of a fluid. One straightforward application of the solver would deal with conjugate heat transfer problems. In practice there exist other similar transport equations that exhibit strong coupling to the velocities. Examples are the scalar transport of volume fractions, when multiphase flow problems are solved with Volume-of-Fluid methods, or turbulent flows, which, depending on the turbulence model, involve further scalar quantities that are strongly coupled to the velocities. Subsequent studies could examine, if for the corresponding transport equations an implicit coupling method could be efficiently implemented.

All linear systems surging from the discretization process of the coupled system of partial differential equations were solved using black-box implementations of the available solvers in the PETSc library. As studies on finite-volume [12, 30, 36] and finite-element [8, 15, 16, 41, 57, 60] discretizations show, a lot of information has not been used in the solution process of the linear system yet, which imposes a barrier on performance. In order to exploit the coupling and the structure of the linear systems two different approaches can be considered. One approach features so called *physics-based* preconditioners as the SIMPLE preconditioner or other Schur-type preconditioners introduced in [15, 30]. These kind of preconditioners accelerate the solution of the linear systems by taking into account the matrix structure, as it is presented in figure 10. At the moment further active research is needed to develop preconditioners for fully coupled systems surging from finite-volume discretizations. Another approach that additionally maintains high scalability of the solution algorithm of the resulting linear systems, have been presented in [12, 36]. In these approaches, algebraic multigrid methods are constructed, that, different to the GAMG preconditioner in PETSc, take into account the matrix structure.

As the studies in the present thesis demonstrated, segregated solution algorithms perform well compared to fully coupled solution algorithms for smaller numbers of unknowns. Furthermore due to the lower requirements on system memory, the problems that can be solved with segregated solution algorithms may for very high numbers of unknowns, not be computable with fully coupled solution algorithms. Especially for flows involving temperature transport many other implementations of segregated solution algorithms exist, whose performance has not been compared with a coupled solution method. Examples can be found in [34, 45].

References

- [1] ACHARYA, S., BALIGA, B. R., KARKI, K., MURTHY, J. Y., PRAKASH, C., AND VANKA, S. P. Pressure-based finite-volume methods in computational fluid dynamics. *Journal of Heat Transfer* 129, 4 (Jan 2007), 407–424.
- [2] ANDERSON, D. A., TANNEHILL, J. C., AND PLETCHER, R. H. *Computational Fluid Mechanics and Heat Transfer*. Hemisphere Publishing Corporation, Wahington, 1984.
- [3] ARIS, R. *Vectors, Tensors and the Basic Equations of Fluid Mechanics*. Prentice-Hall, Inc., Englewood Cliffs, New Jersey, 1962.
- [4] BALAY, S., ABHYANKAR, S., ADAMS, M. F., BROWN, J., BRUNE, P., BUSCHELMAN, K., EIJKHOUT, V., GROPE, W. D., KAUSHIK, D., KNEPLEY, M. G., MCINNES, L. C., RUPP, K., SMITH, B. F., AND ZHANG, H. PETSc users manual. Tech. Rep. ANL-95/11 - Revision 3.5, Argonne National Laboratory, 2014.
- [5] BALAY, S., ABHYANKAR, S., ADAMS, M. F., BROWN, J., BRUNE, P., BUSCHELMAN, K., EIJKHOUT, V., GROPE, W. D., KAUSHIK, D., KNEPLEY, M. G., MCINNES, L. C., RUPP, K., SMITH, B. F., AND ZHANG, H. PETSc Web page. <http://www.mcs.anl.gov/petsc>, 2014.
- [6] BALAY, S., GROPE, W. D., MCINNES, L. C., AND SMITH, B. F. Efficient management of parallelism in object oriented numerical software libraries. In *Modern Software Tools in Scientific Computing* (1997), E. Arge, A. M. Bruaset, and H. P. Langtangen, Eds., Birkhäuser Press, pp. 163–202.
- [7] BONFIGLIOLI, A., CAMPOBASSO, S., CARPENTIERI, B., AND BOLLHÖFER, M. A parallel 3d unstructured implicit rans solver for compressible and incompressible cfd simulations. In *Parallel Processing and Applied Mathematics*, R. Wyrzykowski, J. Dongarra, K. Karczewski, and J. Waśniewski, Eds., vol. 7204 of *Lecture Notes in Computer Science*. Springer Berlin Heidelberg, 2012, pp. 313–322.
- [8] BROWN, J., KNEPLEY, M. G., MAY, D. A., MCINNES, L. C., AND SMITH, B. Composable linear solvers for multiphysics. In *Proceedings of the 2012 11th International Symposium on Parallel and Distributed Computing* (Washington, DC, USA, 2012), ISPD'12, IEEE Computer Society, pp. 55–62.
- [9] CHEN, Z., AND PRZEKAS, A. A coupled pressure-based computational method for incompressible/compressible flows. *Journal of Computational Physics* 229, 24 (2010), 9150 – 9165.
- [10] CHOI, S. K. Note on the use of momentum interpolation method for unsteady flows. *Numerical Heat Transfer, Part A: Applications* 36, 5 (1999), 545–550.
- [11] CHRISTON, M. A., GRESHO, P. M., AND SUTTON, S. B. Computational predictability of time-dependent natural convection flows in enclosures (including a benchmark solution). *International Journal for Numerical Methods in Fluids* 40, 8 (2002), 953–980.
- [12] DARWISH, M., SRAJ, I., AND MOUKALLED, F. A coupled finite volume solver for the solution of incompressible flows on unstructured grids. *Journal of Computational Physics* 228, 1 (2009), 180 – 201.
- [13] DARWISH, F. MOUKALLED, M. A unified formulation of the segregated class of algorithms for fluid flow at all speeds. *Numerical Heat Transfer, Part B: Fundamentals* 37, 1 (2000), 103–139.
- [14] DE VAHL DAVIS, G. Natural convection of air in a square cavity: A bench mark numerical solution. *International Journal for Numerical Methods in Fluids* 3, 3 (1983), 249–264.
- [15] ELMAN, H., HOWLE, V. E., SHADID, J., SHUTTLEWORTH, R., AND TUMINARO, R. A taxonomy and comparison of parallel block multi-level preconditioners for the incompressible navier-stokes equations. *Journal of Computational Physics* 227, 3 (Jan. 2008), 1790–1808.
- [16] ELMAN, H., HOWLE, V. E., SHADID, J., AND TUMINARO, R. A parallel block multi-level preconditioner for the 3d incompressible navier-stokes equations. *Journal of Computational Physics* 187 (2003), 504–523.
- [17] FALGOUT, R., AND YANG, U. hypre: A library of high performance preconditioners. In *Computational Science — ICCS 2002*, P. Sloot, A. Hoekstra, C. Tan, and J. Dongarra, Eds., vol. 2331 of *Lecture Notes in Computer Science*. Springer Berlin Heidelberg, 2002, pp. 632–641.

-
- [18] FALK, U., AND SCHÄFER, M. A fully coupled finite volume solver for the solution of incompressible flows on locally refined non-matching block-structured grids. In *Adaptive Modeling and Simulation 2013* (Barcelona, Spain, June 2013), J. P. M. de Almeida, P. Diez, C. Tiago, and N. Parež, Eds., pp. 235–246.
- [19] FERZIGER, J. H., AND PERIĆ, M. *Numerische Strömungsmechanik*. Springer Verlag, Berlin, 2002.
- [20] GALPIN, P. F., AND RAITHEY, G. D. Numerical solution of problems in incompressible fluid flow: Treatment of the temperature-velocity coupling. *Numerical Heat Transfer* 10, 2 (1986), 105–129.
- [21] GEE, M., SIEFERT, C., HU, J., TUMINARO, R., AND SALA, M. ML 5.0 smoothed aggregation user’s guide. Tech. Rep. SAND2006-2649, Sandia National Laboratories, 2006.
- [22] GRAY, D. D., AND GIORGINI, A. The validity of the boussinesq approximation for liquids and gases. *International Journal of Heat and Mass Transfer* 19, 5 (1976), 545 – 551.
- [23] GROPP, W., LUSK, E., AND SKJELLUM, A. *Using MPI: portable parallel programming with the message-passing interface*, 2. ed. The MIT Press, Cambridge, Massachusetts, 1999.
- [24] GROPP, W. D., KAUSHIK, D. K., KEYES, D. E., AND SMITH, B. F. High performance parallel implicit cfd. *Parallel Computing* 27 (2000), 337–362.
- [25] HACKBUSCH, W. *Theorie und Numerik elliptischer Differentialgleichungen mit Beispielen und Übungsaufgaben*. Teubner-Studienbücher Mathematik. Teubner, Stuttgart, 1996.
- [26] HAGER, G., AND WELLEIN, G. *Introduction to High Performance Computing for Scientists and Engineers*. CRC, Boca Raton, 2011.
- [27] HENDERSON, A. Paraview guide, a parallel visualization application, 2007.
- [28] JASAK, H. *Error Analysis and Estimation for the Finite Volume Method with Applications to Fluid Flows*. PhD thesis, Imperial College of Science, Technology and Medicine, Jun 1996.
- [29] KARIMIAN, S. M., AND STRAATMAN, A. G. Benchmarking of a 3d, unstructured, finite volume code of incompressible navier-stokes equation on a cluster of distributed-memory computers. *High Performance Computing Systems and Applications, Annual International Symposium on O* (2005), 11–16.
- [30] KLAIJ, C. M., AND VUIK, C. Simple-type preconditioners for cell-centered, colocated finite volume discretization of incompressible reynolds-averaged navier-stokes equations. *International Journal for Numerical Methods in Fluids* 71, 7 (2013), 830–849.
- [31] KUNDU, P. K., COHEN, I. M., AND DOWNLING, D. R. *Fluid Mechanics*, 5 ed. Elsevier, 2012.
- [32] LANGE, C. F., SCHÄFER, M., AND DURST, F. Local block refinement with a multigrid flow solver. *International Journal for Numerical Methods in Fluids* 38, 1 (2002), 21–41.
- [33] LILEK, U., MUZAFERIJA, S., PERIĆ, M., AND SEIDL, V. An implicit finite-volume method using nonmatching blocks of structured grid. *Numerical Heat Transfer, Part B: Fundamentals* 32, 4 (1997), 385–401.
- [34] LIU, X., TAO, W., AND HE, Y. A simple method for improving the simpler algorithm for numerical simulations of incompressible fluid flow and heat transfer problems. *Engineering Computations* 22, 8 (2005), 921–939.
- [35] MAJUMDAR, S. Role of underrelaxation in momentum interpolation for calculation of flow with nonstaggered grids. *Numerical Heat Transfer* 13, 1 (1988), 125–132.
- [36] MANGANI, L., BUCHMAYR, M., AND DARWISH, M. Development of a novel fully coupled solver in openfoam: Steady-state incompressible turbulent flows in rotational reference frames. *Numerical Heat Transfer, Part B: Fundamentals* 66, 6 (2014), 526–543.
- [37] MAPLE. *version 18*. Waterloo Maple Inc. (Maplesoft), Waterloo, Ontario, 2014.
- [38] MATLAB. *version 8.4.150421 (R2014b)*. The MathWorks Inc., Natick, Massachusetts, 2014.
- [39] MCCALPIN, J. D. Stream: Sustainable memory bandwidth in high performance computers. Tech. rep., University of Virginia, Charlottesville, Virginia, 1991-2007. A continually updated technical report. <http://www.cs.virginia.edu/stream/>.

-
- [40] McCALPIN, J. D. Memory bandwidth and machine balance in current high performance computers. *IEEE Computer Society Technical Committee on Computer Architecture (TCCA) Newsletter* (Dec. 1995), 19–25.
- [41] MCINNES, L. C., SMITH, B., ZHANG, H., AND MILLS, R. T. Hierarchical krylov and nested krylov methods for extreme-scale computing. *Parallel Computing* 40, 1 (Jan. 2014), 17–31.
- [42] MILLER, T. F., AND SCHMIDT, F. W. Use of a pressure-weighted interpolation method for the solution of the incompressible navier-stokes equations on a nonstaggered grid system. *Numerical Heat Transfer* 14, 2 (1988), 213–233.
- [43] MUZAFERIJA, S. *Adaptive finite volume method for flow predictions using unstructured meshes and multigrid approach*. PhD thesis, University of London, 1994.
- [44] OBERKAMPE, W. L., AND TRUCANO, T. G. Verification and validation in computational fluid dynamics. *Progress in Aerospace Sciences* 38, 3 (2002), 209 – 272.
- [45] OLIVEIRA, P. J., AND ISSA, R. I. An improved piso algorithm for the computation of buoyancy-driven flows. *Numerical Heat Transfer, Part B: Fundamentals* 40, 6 (2001), 473–493.
- [46] PATANKAR, S., AND SPALDING, D. A calculation procedure for heat, mass and momentum transfer in three-dimensional parabolic flows. *International Journal of Heat and Mass Transfer* 15, 10 (1972), 1787 – 1806.
- [47] PATANKAR, S. V. *Numerical heat transfer and fluid flow*. McGraw - Hill, New York, 1980.
- [48] PERIĆ, M. Analysis of pressure-velocity coupling on nonorthogonal grids. *Numerical Heat Transfer* 17 (Jan. 1990), 63–82.
- [49] POPE, S. B. *Turbulent Flows*. Cambridge University Press, New York, 2000.
- [50] R., H. M., AND E., S. Methods of conjugate gradients for solving linear systems. *Journal of Research of the National Bureau of Standards* 49 (1952), 409.
- [51] RHIE, C. M., AND CHOW, W. L. Numerical study of the turbulent flow past an airfoil with trailing edge separation. *AIAA Journal* 21 (Nov. 1983), 1525–1532.
- [52] SAAD, Y. *Iterative methods for sparse linear systems*. SIAM, 2003.
- [53] SAAD, Y., AND SCHULTZ, M. H. Gmres: A generalized minimal residual algorithm for solving nonsymmetric linear systems. *SIAM Journal on Scientific and Statistical Computing* 7, 3 (July 1986), 856–869.
- [54] SALARI, K., AND KNUPP, P. Code verification by the method of manufactured solutions. Tech. Rep. SAND2000-1444, Sandia National Labs., Albuquerque, NM (US); Sandia National Labs., Livermore, CA (US), Jun 2000.
- [55] SCHÄFER, M. *Numerik im Maschinenbau*. Springer Verlag, Berlin, 1999.
- [56] SHEU, T. W. H., AND LIN, R. K. Newton linearization of the incompressible navier–stokes equations. *International Journal for Numerical Methods in Fluids* 44, 3 (2004), 297–312.
- [57] SILVESTER, D., ELMAN, H., KAY, D., AND WATHEN, A. Efficient preconditioning of the linearized navier–stokes equations for incompressible flow. *Journal of Computational and Applied Mathematics* 128, 1–2 (2001), 261 – 279. Numerical Analysis 2000. Vol. VII: Partial Differential Equations.
- [58] SPURK, J. H., AND AKSEL, N. *Strömungslehre: Einführung in die Theorie der Strömungen*, 8. ed. Springer Verlag, Berlin, 2010.
- [59] TAYLOR, G. I., AND GREEN, A. E. Mechanism of the production of small eddies from large ones. *Proceedings of the Royal Society of London. Series A, Mathematical and Physical Sciences* 158, 895 (1937), pp. 499–521.
- [60] TUREK, S., AND SCHMACHTEL, R. Fully coupled and operator-splitting approaches for natural convection flows in enclosures. *International Journal for Numerical Methods in Fluids* 40, 8 (2002), 1109–1119.
- [61] VAKILIPOUR, S., AND ORMISTON, S. J. A coupled pressure-based co-located finite-volume solution method for natural-convection flows. *Numerical Heat Transfer, Part B: Fundamentals* 61, 2 (2012), 91–115.
- [62] VAN DOORMAAL, J. P., AND RAITHBY, G. D. Enhancements of the simple method for predicting incompressible fluid flows. *Numerical Heat Transfer* 7, 2 (1984), 147–163.

-
- [63] ZHANG, S., ZHAO, X., AND BAYYUK, S. Generalized formulations for the rhie–chow interpolation. *Journal of Computational Physics* 258, 0 (2014), 880 – 914.